

Assignment 7

Generative Models

In [0]:

```
import tensorflow as tf
tf.enable_eager_execution()

import functools
import matplotlib.pyplot as plt
import numpy as np
import pdb
```

Downloading code of utils

In [3]:

```
# Download the class repository
! git clone https://github.com/aamini/introtodeeplearning_labs.git > /dev/null 2>&1
% cd introtodeeplearning_labs
! git pull
% cd ..

!cp -r introtodeeplearning_labs/lab1 lab1
!cp -r /content/lab1/util.py util.py

# Import the necessary class-specific utility files for this lab
import introtodeeplearning_labs as util

/content/introtodeeplearning_labs
Already up to date.
/content
```

Download Datasets

In [4]:

```
# Get the training data: both images from CelebA and ImageNet
path_to_training_data = tf.keras.utils.get_file('train_face.h5', 'https://www.dropbox.com/s/l5iqduhe0gwxumq/train_face.h5?dl=1')

Downloading data from https://www.dropbox.com/s/l5iqduhe0gwxumq/train_face.h5?dl=1
1263894528/1263889489 [=====] - 26s 0us/step
1263902720/1263889489 [=====] - 26s 0us/step
```

In [5]:

```
# Instantiate a TrainingDatasetLoader using the downloaded dataset
loader = util.TrainingDatasetLoader(path_to_training_data)
number_of_training_examples = loader.get_train_size()
(images, labels) = loader.get_batch(100)
```

```
Opening /root/.keras/datasets/train_face.h5
Loading data into memory...
```

Preview dataset

In [6]:

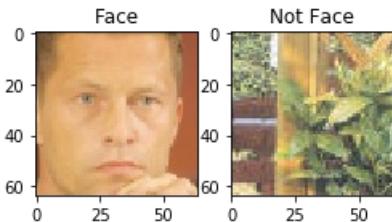
```
#@title Change the sliders to look at positive and negative training examples! { run: "auto" }

face_images = images[np.where(labels==1)[0]]
not_face_images = images[np.where(labels==0)[0]]

idx_face = 39 #@param {type:"slider", min:0, max:50, step:1}
idx_not_face = 39 #@param {type:"slider", min:0, max:50, step:1}

plt.figure(figsize=(4,2))
plt.subplot(1, 2, 1)
plt.imshow(face_images[idx_face])
plt.title("Face")
plt.grid(False)

plt.subplot(1, 2, 2)
plt.imshow(not_face_images[idx_not_face])
plt.title("Not Face")
plt.grid(False)
```



In [7]:

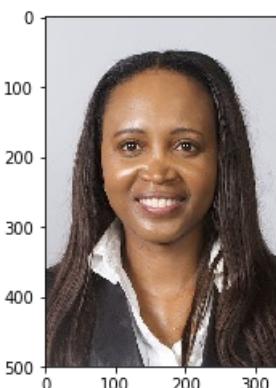
```
#@title Change the sliders to look at positive and negative training examples! { run: "auto" }

ppb = util.PPBFaceEvaluator(skip=4) # create the dataset handler

gender = "female" #@param ["male", "female"]
skin_color = "darker" #@param ["lighter", "darker"]

img = ppb.get_sample_faces_from_demographic(gender, skin_color)
plt.imshow(img)
plt.grid(False)
```

Downloading data from <https://www.dropbox.com/s/l0lp6qxeplumouf/PPB.tar?dl=1>
86245376/86241280 [=====] - 2s 0us/step
86253568/86241280 [=====] - 2s 0us/step



Build Baseline Model

In [0]:

```
n_outputs = 1 # number of outputs (i.e., face or not face)
n_filters = 12 # base number of convolutional filters

'''Function to define a standard CNN model'''
def make_standard_classifier():
    Conv2D = functools.partial(tf.keras.layers.Conv2D, padding='same', activation='relu')
    BatchNormalization = tf.keras.layers.BatchNormalization
    Flatten = tf.keras.layers.Flatten
    Dense = functools.partial(tf.keras.layers.Dense, activation='relu')

    model = tf.keras.Sequential([
        # TODO: define a convolutional layer with n_filters 5x5 filters and 2x2 stride
        Conv2D(n_filters, 5, strides=(2,2)),
        BatchNormalization(),

        # TODO: define a convolutional layer with 2*n_filters 5x5 filters and 2x2 stride
        Conv2D(2 * n_filters, 5, strides=(2,2)),
        BatchNormalization(),

        # TODO: define a convolutional layer with 4*n_filters 3x3 filters and 2x2 stride
        Conv2D(4 * n_filters, 3, strides=(2,2)),
        BatchNormalization(),

        # TODO: define a convolutional layer with 6*n_filters 3x3 filters and 1x1 stride
        Conv2D(6 * n_filters, 3, strides=(1,1)),
        BatchNormalization(),

        Flatten(),
        Dense(1, activation=None),
        tf.keras.layers.Dropout(0.5)
    ])
    return model
```

Training Baseline

In [57]:

```
standard_classifier = make_standard_classifier()

batch_size = 36
num_epochs = 10 # keep small to run faster
learning_rate = 1e-3

optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate) # define our optimizer
loss_history = util.LossHistory(smoothing_factor=0.99) # to record the evolution of the loss
plotter = util.PeriodicPlotter(sec=2, scale='semilogy')

# The training loop!
for epoch in range(num_epochs):

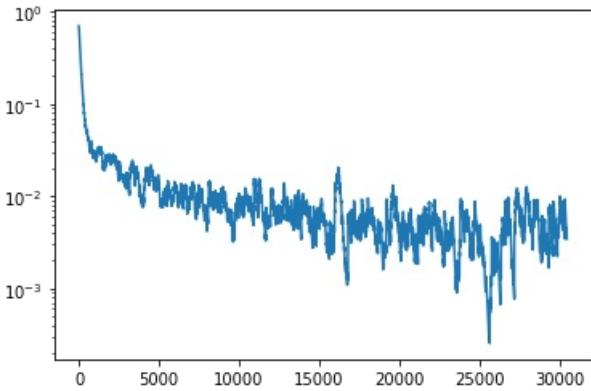
    custom_msg = util.custom_progress_text("Epoch: %(epoch).0f Loss: %(loss)2.2f")
    bar = util.create_progress_bar(custom_msg)

    for idx in bar(range(loader.get_train_size()//batch_size)):
        # First grab a batch of training data and convert the input images to tensors
        x, y = loader.get_batch(batch_size)
        x = tf.convert_to_tensor(x, dtype=tf.float32)
        y = tf.convert_to_tensor(y, dtype=tf.float32)

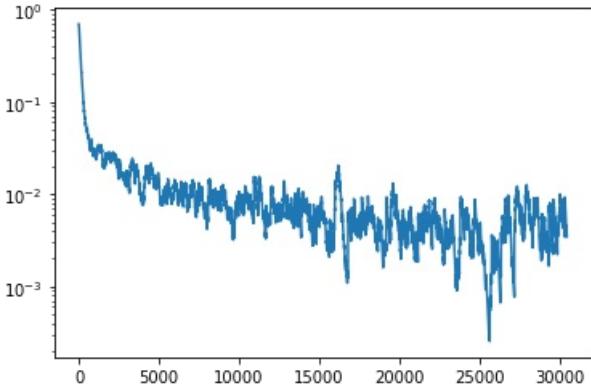
        # GradientTape to record differentiation operations
        with tf.GradientTape() as tape:
            logits = standard_classifier(x) # feed the images into the model
            loss_value = tf.nn.sigmoid_cross_entropy_with_logits(labels=y, logits=logits) # compute the loss

        custom_msg.update_mapping(epoch=epoch, loss=loss_value.numpy().mean())
        # Backpropagation
        grads = tape.gradient(loss_value, standard_classifier.variables)
        optimizer.apply_gradients(zip(grads, standard_classifier.variables), global_step=tf.train.get_or_create_global_step())

    loss_history.append(loss_value.numpy().mean())
    plotter.plot(loss_history.get())
```



100% #####| Time: 0:01:40 Epoch: 9 Loss: 0.00



Evaluate Baseline CNN

In [58]:

```
# Evaluate on a subset of CelebA+Imagenet
(batch_x, batch_y) = loader.get_batch(5000)
y_pred_standard = tf.round(tf.nn.sigmoid(standard_classifier.predict(batch_x)))
acc_standard = tf.reduce_mean(tf.cast(tf.equal(batch_y, y_pred_standard), tf.float32))
print "Standard CNN accuracy on (potentially biased) training set: {:.4f}".format(acc_standard.numpy())

# Evaluate on PPB dataset (takes ~3 minutes)
standard_cnn_accuracy = []
for skin_color in ['lighter', 'darker']:
    for gender in ['male', 'female']:
        standard_cnn_accuracy.append(ppb.evaluate([standard_classifier], gender, skin_color, from_logit=True)[0])
    print
    print "{} {}: {}".format(gender, skin_color, standard_cnn_accuracy[-1])

plt.bar(range(4), standard_cnn_accuracy)
plt.xticks(range(4), ('LM', 'LF', 'DM', 'DF'))
plt.ylim(np.min(standard_cnn_accuracy)-0.1,np.max(standard_cnn_accuracy)+0.1)
plt.ylabel('Accuracy')
```

WARNING:tensorflow:From /usr/local/lib/python2.7/dist-packages/tensorflow/python/keras/layers/core.py:143: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

N/A% (0 of 97) | Elapsed Time: 0:00:00 ETA: ---::--

Standard CNN accuracy on (potentially biased) training set: 0.9984

100% (97 of 97) |#####| Elapsed Time: 0:01:09 Time: 0:01:09
N/A% (0 of 72) | Elapsed Time: 0:00:00 ETA: ---::--

male lighter: 0.824742268041

100% (72 of 72) |#####| Elapsed Time: 0:00:49 Time: 0:00:49
N/A% (0 of 78) | Elapsed Time: 0:00:00 ETA: ---::--

female lighter: 0.861111111111

100% (78 of 78) |#####| Elapsed Time: 0:00:51 Time: 0:00:51
N/A% (0 of 71) | Elapsed Time: 0:00:00 ETA: ---::--

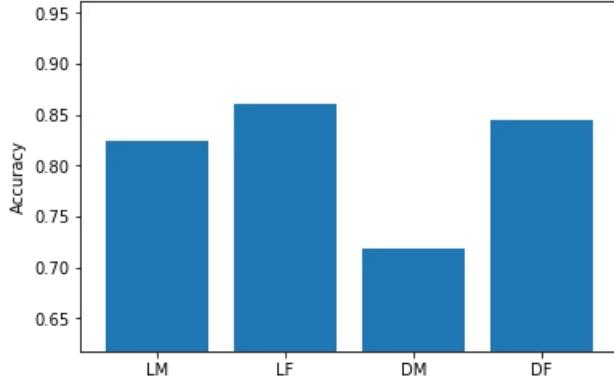
male darker: 0.717948717949

100% (71 of 71) |#####| Elapsed Time: 0:00:45 Time: 0:00:45

female darker: 0.845070422535

Out[58]:

Text(0,0.5,'Accuracy')



We can notice a bias since for lighter skin tone, the accuracy is considerably higher. There might also be a slight bias for females over males.

Vae Loss Function

In [0]:

```
# Function to calculate VAE loss given an input x, reconstructed output x_pred,
# encoded means mu, encoded log of standard deviation logsigma, and weight parameter for the latent loss
def vae_loss_function(x, x_pred, mu, logsigma, kl_weight=0.0005):
    '''TODO: Define the latent loss'''
    latent_loss = 0.5 * tf.reduce_sum(mu**2 - logsigma + tf.exp(logsigma) - 1)
    '''TODO: Define the reconstruction loss. Hint: you'll need to use tf.reduce_mean'''
    reconstruction_loss = tf.reduce_mean((x - x_pred) ** 2)
    '''TODO: Define the VAE loss'''
    vae_loss = kl_weight * latent_loss + reconstruction_loss
    return vae_loss
```

Reparameterization and Sampling

In [0]:

```
"""Reparameterization trick by sampling from an isotropic unit Gaussian.
# Arguments
args (tensor): mean and log of standard deviation of latent distribution (Q(z|X))
# Returns
z (tensor): sampled latent vector
"""

def sampling(args):
    z_mean, z_logsigtma = args
    batch = z_mean.shape[0]
    dim = z_mean.shape[1]

    # by default, random_normal has mean=0 and std=1.0
    epsilon = tf.random_normal(tf.shape(z_mean))
    '''TODO: Define the reparameterization computation!'''
    z = z_mean + tf.exp(0.5 * z_logsigtma) * epsilon
    return z
```

Debiasing Vae Loss Function

In [0]:

```
# Loss function for DB-VAE
def debiasing_loss_function(x, x_pred, y, y_logit, mu, logsigma):

    '''TODO: call the relevant function to obtain VAE loss'''
    vae_loss = vae_loss_function(x, x_pred, mu, logsigma, kl_weight=0.0005)
    '''TODO: define the classification loss'''
    classification_loss = tf.nn.sigmoid_cross_entropy_with_logits(labels=y, logits=y_logit)
    # Use the training data labels to create variable face_mask
    face_mask = tf.cast(tf.equal(y, 1), tf.float32)

    '''TODO: define the DB-VAE total loss! Hint: think about the dimensionality of your output.'''
    total_loss = tf.reduce_mean(classification_loss + face_mask * vae_loss)
    return total_loss, classification_loss
```

Vae Model Creation

In [0]:

```
latent_dim = 100
```

Encoder

In [0]:

```
'''Define the encoder network for the DB-VAE'''
def make_face_encoder_network():
    Conv2D = functools.partial(tf.keras.layers.Conv2D, padding='same', activation='relu')
    BatchNormalization = tf.keras.layers.BatchNormalization
    Flatten = tf.keras.layers.Flatten
    Dense = functools.partial(tf.keras.layers.Dense, activation='relu')

    inputs = tf.keras.layers.Input(shape=(64,64,3))

    hidden = Conv2D(filters=1*n_filters, kernel_size=[5,5], strides=[2,2])(inputs)
    hidden = BatchNormalization()(hidden)
    hidden = Conv2D(filters=2*n_filters, kernel_size=[5,5], strides=[2,2])(hidden)
    hidden = BatchNormalization()(hidden)
    hidden = Conv2D(filters=4*n_filters, kernel_size=[3,3], strides=[2,2])(hidden)
    hidden = BatchNormalization()(hidden)
    hidden = Conv2D(filters=6*n_filters, kernel_size=[3,3], strides=[1,1])(hidden)
    hidden = BatchNormalization()(hidden)

    hidden = Flatten(name='flatten')(hidden)

    '''Encoder outputs:
        y_logit: supervised class prediction
        z_mean: means in the latent space
        z_logsigma: standard deviations in the latent space'''
    y_logit = Dense(1, activation=None, name='y_logit')(hidden)
    z_mean = Dense(latent_dim, name='z_mean')(hidden)
    z_logsigma = Dense(latent_dim, name='z_logsigma')(hidden)

    # use reparameterization trick to sample from the latent space
    z = tf.keras.layers.Lambda(sampling, output_shape=(latent_dim,))([z_mean, z_logsigma])

    # define the outputs that the encoder model should return
    outputs = [y_logit, z_mean, z_logsigma, z]
    # finalize the encoder model
    encoder = tf.keras.Model(inputs=inputs, outputs=outputs, name='encoder')

    # get the shape of the final convolutional output (right before the flatten)
    flatten_layer_idx = encoder.layers.index(encoder.get_layer('flatten'))
    pre_flatten_shape = encoder.layers[flatten_layer_idx-1].get_output_at(0).shape[1:]

    return encoder, inputs, outputs, pre_flatten_shape
```

Decoder

In [0]:

```
'''Define the decoder network for the DB-VAE'''
def make_face_decoder_network(pre_flatten_shape):
    Conv2DTranspose = functools.partial(tf.keras.layers.Conv2DTranspose, padding='same', activation='relu')
    BatchNormalization = tf.keras.layers.BatchNormalization
    Flatten = tf.keras.layers.Flatten
    Dense = functools.partial(tf.keras.layers.Dense, activation='relu')

    latent_inputs = tf.keras.layers.Input(shape=(latent_dim,))

    hidden = Dense(tf.reduce_prod(pre_flatten_shape))(latent_inputs)
    hidden = tf.keras.layers.Reshape(pre_flatten_shape)(hidden)

    # series of deconvolutional layers with batch normalization
    hidden = Conv2DTranspose(filters=4*n_filters, kernel_size=[3,3], strides=[1,1])(hidden)
    hidden = BatchNormalization()(hidden)
    hidden = Conv2DTranspose(filters=2*n_filters, kernel_size=[3,3], strides=[2,2])(hidden)
    hidden = BatchNormalization()(hidden)
    hidden = Conv2DTranspose(filters=1*n_filters, kernel_size=[5,5], strides=[2,2])(hidden)
    hidden = BatchNormalization()(hidden)

    x_hat = Conv2DTranspose(filters=3, kernel_size=[5,5], strides=[2,2])(hidden)

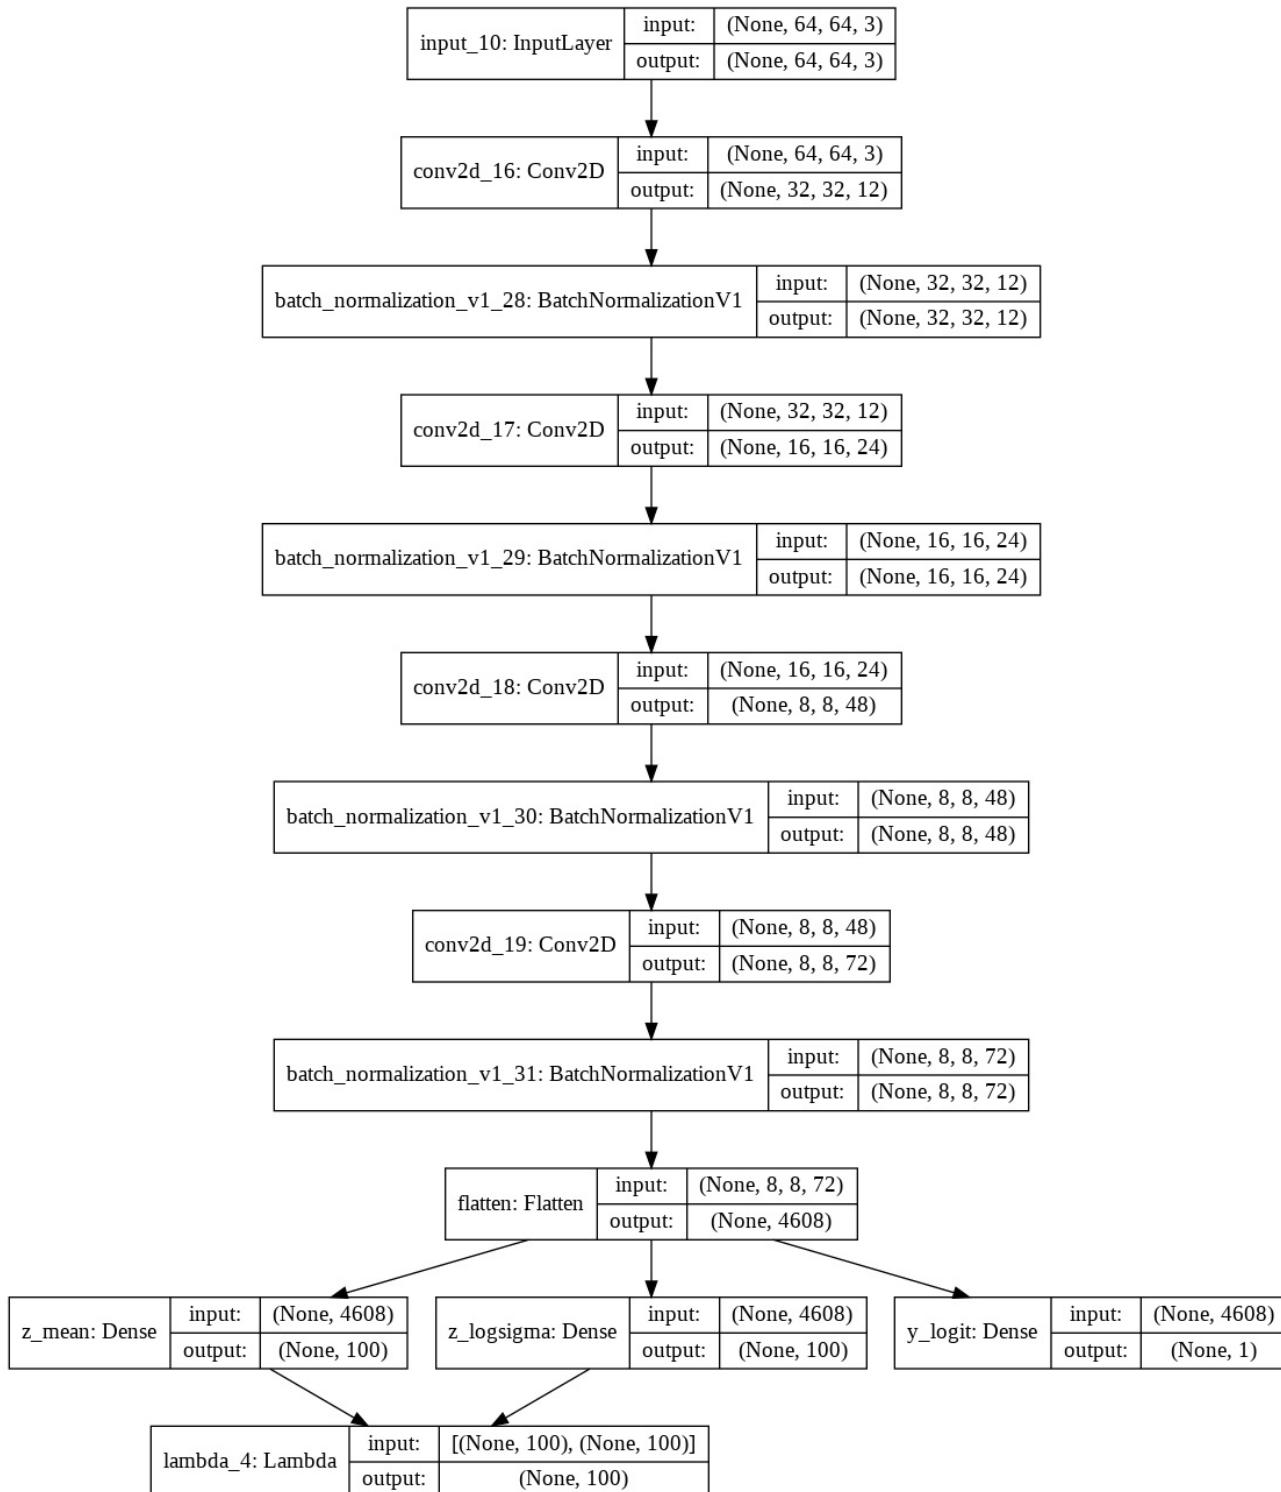
    # instantiate decoder model
    decoder = tf.keras.Model(inputs=latent_inputs, outputs=x_hat, name='decoder')
    return decoder
```

In [50]:

```
'''TODO: create the encoder and decoder networks'''
encoder, inputs, outputs, pre_flatten_shape = make_face_encoder_network()
decoder = make_face_decoder_network(pre_flatten_shape)
# initialize the models
encoder_output = encoder(inputs)
y_logit, z_mean, z_logsigma, z = encoder_output
reconstructed_inputs = decoder(z)

vae = tf.keras.Model(inputs, reconstructed_inputs)
util.display_model(encoder)
```

Out[50]:



VAE Utils

In [0]:

```
# Function to return the means for an input image batch
def get_latent_mu(images, encoder, batch_size=1024):
    N = images.shape[0]
    mu = np.zeros((N, latent_dim))
    for start_ind in xrange(0, N, batch_size):
        end_ind = min(start_ind+batch_size, N+1)
        batch = images[start_ind:end_ind]
        batch = tf.convert_to_tensor(batch, dtype=tf.float32)/255.
        _, batch_mu, _, _ = encoder(batch)
        mu[start_ind:end_ind] = batch_mu
    return mu

'''Function that recomputes the sampling probabilities for images within a batch
based on how they distribute across the '''
def get_training_sample_probabilities(images, encoder, bins=10, smoothing_fac=0.0):
    print "Recomputing the sampling probabilities"

    mu = get_latent_mu(images, encoder)
    # sampling probabilities for the images
    training_sample_p = np.zeros(mu.shape[0])

    # consider the distribution for each latent variable
    for i in range(latent_dim):

        latent_distribution = mu[:,i]
        # generate a histogram of the latent distribution
        hist_density, bin_edges = np.histogram(latent_distribution, density=True, bins=bins)

        # find which latent bin every data sample falls in
        # https://docs.scipy.org/doc/numpy-1.13.0/reference/generated/numpy.digitize.html
        bin_edges[0] = -float('inf')
        bin_edges[-1] = float('inf')
        '''TODO: call the digitize function to find which bins in the latent distribution
        every data sample falls in to'''
        bin_index = np.digitize(latent_distribution, bin_edges)
        # smooth the density function [Eq. #]
        hist_smoothed_density = hist_density + smoothing_fac
        hist_smoothed_density = hist_smoothed_density / np.sum(hist_smoothed_density)

        '''TODO: invert the density function to compute the sampling probability!
        HINT: think carefully about the indexing of the bins! What is the length of bin_edges?'''
        p = 1 / hist_smoothed_density[bin_index - 1]
        # normalize all probabilities
        p = p / np.sum(p)

        # update sampling probabilities
        training_sample_p = np.maximum(p, training_sample_p)

    # final normalization
    training_sample_p /= np.sum(training_sample_p)

    return training_sample_p
```

VAE Training and unbiased Model

In [55]:

```
loss_history = []

batch_size = 36
num_epochs = 10 # keep small to run faster
learning_rate = 1e-3

optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate)

enable_debiasing = True
all_faces = loader.get_all_train_faces() # parameter from data loader

for epoch in range(num_epochs):

    # progress message and bar
    custom_msg = util.custom_progress_text("Epoch: %(epoch).0f Iter: %(idx).0f Class Loss: %(class_loss)2.
2f Loss: %(loss)2.2f")
    bar = util.create_progress_bar(custom_msg)

    p_faces = None
    if enable_debiasing:
        # Recompute data sampling probabilities if debiasing is enabled
        '''TODO: write the function call to recompute the sampling probabilities
        when debiasing is enabled'''
        p_faces = get_training_sample_probabilities(all_faces, encoder, bins=10, smoothing_fac=0.0)
    for idx in bar(range(loader.get_train_size()//batch_size)):
        # load a batch of data
        (x, y) = loader.get_batch(batch_size, p_pos=p_faces)
        x = tf.convert_to_tensor(x, dtype=tf.float32)
        y = tf.convert_to_tensor(y, dtype=tf.float32)

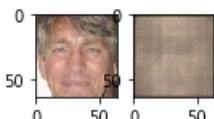
        # define GradientTape for automatic differentiation
        with tf.GradientTape() as tape:
            y_logit, mu, logsigma, z = encoder(x)
            x_hat = decoder(z)
            '''TODO: call the relevant loss function to compute the loss'''
            loss, class_loss = debiasing_loss_function(x, x_hat, y, y_logit, mu, logsigma)
            '''TODO: use the GradientTape.gradient method to compute the gradients'''
            grads = tape.gradient(loss, vae.variables)
        # apply gradients to variables
        optimizer.apply_gradients(zip(grads, vae.variables),
                                global_step=tf.train.get_or_create_global_step())

        # track the losses
        class_loss_value = class_loss.numpy().mean()
        loss_value = loss.numpy().mean()
        loss_history.append((class_loss_value, loss_value))
        custom_msg.update_mapping(epoch=epoch, idx=idx, loss=loss_value, class_loss=class_loss_value)

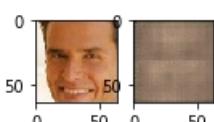
        # plot the progress every 100 steps
        if idx%100 == 0:
            util.plot_sample(x,y,vae)
```

Recomputing the sampling probabilities

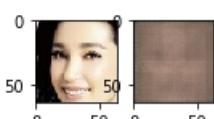
N/A%| |ETA: --:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



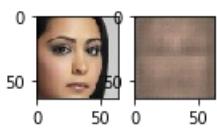
3%| |ETA: 0:02:44 Epoch: 0 Iter: 99 Class Loss: 0.02 Loss: 0.05



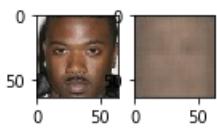
6%| |ETA: 0:02:40 Epoch: 0 Iter: 199 Class Loss: 0.01 Loss: 0.04



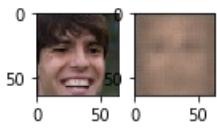
9%| |ETA: 0:02:32 Epoch: 0 Iter: 298 Class Loss: 0.11 Loss: 0.15



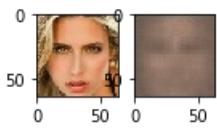
13% | ETA: 0:02:31 Epoch: 0 Iter: 398 Class Loss: 0.07 Loss: 0.11



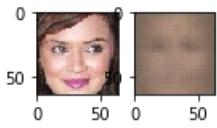
16% | ETA: 0:02:21 Epoch: 0 Iter: 499 Class Loss: 0.01 Loss: 0.04



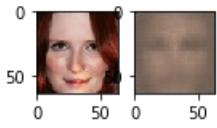
19% | ETA: 0:02:17 Epoch: 0 Iter: 599 Class Loss: 0.01 Loss: 0.04



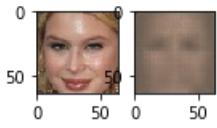
22% | # ETA: 0:02:09 Epoch: 0 Iter: 699 Class Loss: 0.01 Loss: 0.03



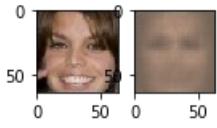
26% | # ETA: 0:02:05 Epoch: 0 Iter: 798 Class Loss: 0.07 Loss: 0.11



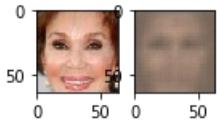
29% | # ETA: 0:01:57 Epoch: 0 Iter: 898 Class Loss: 0.02 Loss: 0.06



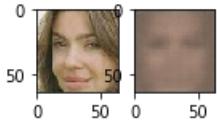
32% | # ETA: 0:01:54 Epoch: 0 Iter: 999 Class Loss: 0.00 Loss: 0.03



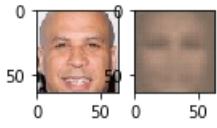
35% | # ETA: 0:01:48 Epoch: 0 Iter: 1098 Class Loss: 0.13 Loss: 0.16



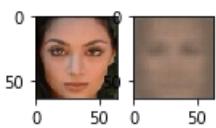
39% | # ETA: 0:01:57 Epoch: 0 Iter: 1198 Class Loss: 0.06 Loss: 0.09



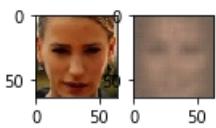
42% | # ETA: 0:01:52 Epoch: 0 Iter: 1299 Class Loss: 0.03 Loss: 0.07



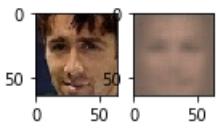
45% | # ETA: 0:01:27 Epoch: 0 Iter: 1398 Class Loss: 0.00 Loss: 0.04



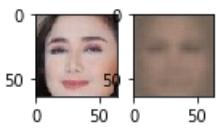
49%|#| ETA: 0:01:23 Epoch: 0 Iter: 1499 Class Loss: 0.00 Loss: 0.03



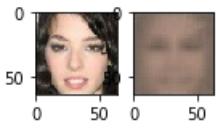
52%|##| ETA: 0:01:18 Epoch: 0 Iter: 1599 Class Loss: 0.00 Loss: 0.04



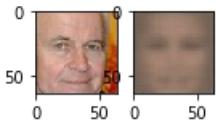
55%|##| ETA: 0:01:12 Epoch: 0 Iter: 1699 Class Loss: 0.04 Loss: 0.07



58%|##| ETA: 0:01:07 Epoch: 0 Iter: 1798 Class Loss: 0.01 Loss: 0.05



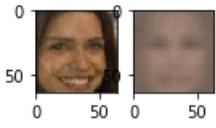
62%|##| ETA: 0:01:00 Epoch: 0 Iter: 1899 Class Loss: 0.02 Loss: 0.05



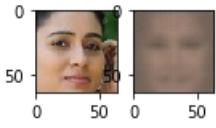
65%|##| ETA: 0:00:54 Epoch: 0 Iter: 1998 Class Loss: 0.01 Loss: 0.04



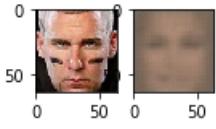
68%|##| ETA: 0:01:00 Epoch: 0 Iter: 2098 Class Loss: 0.01 Loss: 0.04



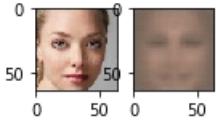
72%|##| ETA: 0:00:54 Epoch: 0 Iter: 2198 Class Loss: 0.00 Loss: 0.03



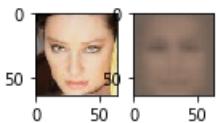
75%|###| ETA: 0:00:40 Epoch: 0 Iter: 2299 Class Loss: 0.01 Loss: 0.04



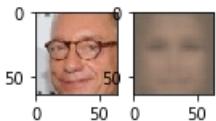
78%|###| ETA: 0:00:34 Epoch: 0 Iter: 2399 Class Loss: 0.00 Loss: 0.03



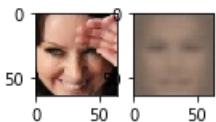
81%|###| ETA: 0:00:30 Epoch: 0 Iter: 2498 Class Loss: 0.00 Loss: 0.04



85%|###| ETA: 0:00:28 Epoch: 0 Iter: 2599 Class Loss: 0.00 Loss: 0.04



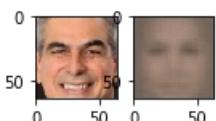
88%|###| ETA: 0:00:20 Epoch: 0 Iter: 2699 Class Loss: 0.00 Loss: 0.03



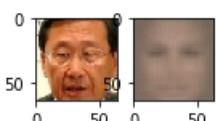
91%|###| ETA: 0:00:13 Epoch: 0 Iter: 2799 Class Loss: 0.00 Loss: 0.04



94%|###| ETA: 0:00:08 Epoch: 0 Iter: 2898 Class Loss: 0.00 Loss: 0.04



98%|###| ETA: 0:00:02 Epoch: 0 Iter: 2999 Class Loss: 0.03 Loss: 0.06



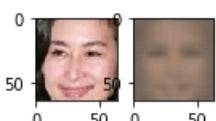
100%|####| Time: 0:02:56 Epoch: 0 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

N/A%| | ETA: ---- Epoch: nan Iter: nan Class Loss: nan Loss: nan



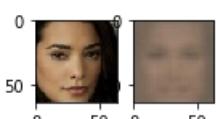
3%| | ETA: 0:02:16 Epoch: 1 Iter: 97 Class Loss: 0.34 Loss: 0.37



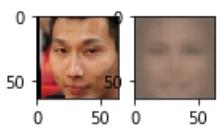
6%| | ETA: 0:02:14 Epoch: 1 Iter: 198 Class Loss: 0.00 Loss: 0.04



9%| | ETA: 0:02:08 Epoch: 1 Iter: 298 Class Loss: 0.01 Loss: 0.04



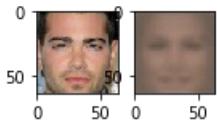
13%| | ETA: 0:02:02 Epoch: 1 Iter: 398 Class Loss: 0.00 Loss: 0.04



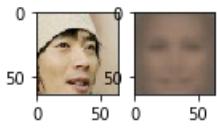
16% | ETA: 0:02:01 Epoch: 1 Iter: 499 Class Loss: 0.00 Loss: 0.03



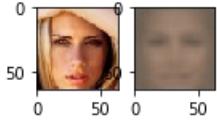
19% | ETA: 0:01:54 Epoch: 1 Iter: 597 Class Loss: 0.01 Loss: 0.04



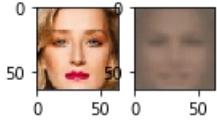
22% |# ETA: 0:01:49 Epoch: 1 Iter: 698 Class Loss: 0.00 Loss: 0.03



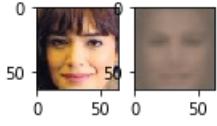
26% |# ETA: 0:01:45 Epoch: 1 Iter: 799 Class Loss: 0.09 Loss: 0.12



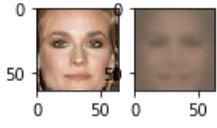
29% |# ETA: 0:01:41 Epoch: 1 Iter: 898 Class Loss: 0.00 Loss: 0.03



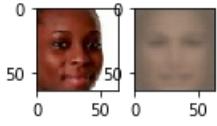
32% |# ETA: 0:01:52 Epoch: 1 Iter: 998 Class Loss: 0.05 Loss: 0.08



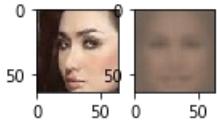
35% |# ETA: 0:01:47 Epoch: 1 Iter: 1098 Class Loss: 0.02 Loss: 0.05



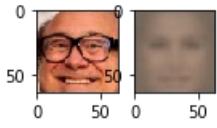
39% |# ETA: 0:01:26 Epoch: 1 Iter: 1198 Class Loss: 0.02 Loss: 0.05



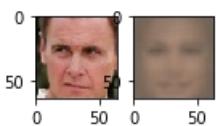
42% |# ETA: 0:01:23 Epoch: 1 Iter: 1297 Class Loss: 0.00 Loss: 0.03



45% |# ETA: 0:01:18 Epoch: 1 Iter: 1397 Class Loss: 0.01 Loss: 0.04



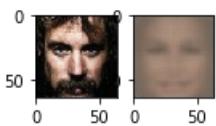
49% |# ETA: 0:01:12 Epoch: 1 Iter: 1499 Class Loss: 0.00 Loss: 0.03



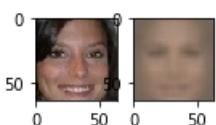
52%|##| ETA: 0:01:07 Epoch: 1 Iter: 1599 Class Loss: 0.01 Loss: 0.04



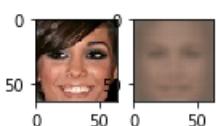
55%|##| ETA: 0:01:04 Epoch: 1 Iter: 1697 Class Loss: 0.00 Loss: 0.03



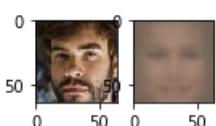
58%|##| ETA: 0:00:59 Epoch: 1 Iter: 1798 Class Loss: 0.01 Loss: 0.04



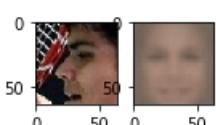
62%|##| ETA: 0:00:54 Epoch: 1 Iter: 1899 Class Loss: 0.00 Loss: 0.04



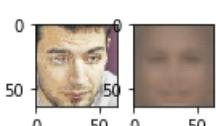
65%|##| ETA: 0:00:48 Epoch: 1 Iter: 1999 Class Loss: 0.00 Loss: 0.04



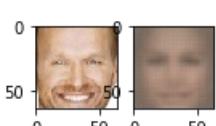
68%|##| ETA: 0:00:44 Epoch: 1 Iter: 2097 Class Loss: 0.00 Loss: 0.04



72%|##| ETA: 0:00:40 Epoch: 1 Iter: 2199 Class Loss: 0.00 Loss: 0.03



75%|###| ETA: 0:00:35 Epoch: 1 Iter: 2297 Class Loss: 0.00 Loss: 0.03



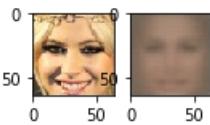
78%|###| ETA: 0:00:30 Epoch: 1 Iter: 2398 Class Loss: 0.00 Loss: 0.03



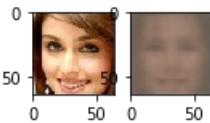
81%|###| ETA: 0:00:26 Epoch: 1 Iter: 2499 Class Loss: 0.00 Loss: 0.03



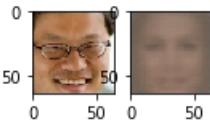
85%|###| ETA: 0:00:21 Epoch: 1 Iter: 2599 Class Loss: 0.00 Loss: 0.03



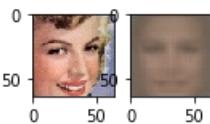
88%|###|ETA: 0:00:18 Epoch: 1 Iter: 2698 Class Loss: 0.02 Loss: 0.05



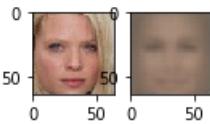
91%|###|ETA: 0:00:12 Epoch: 1 Iter: 2799 Class Loss: 0.00 Loss: 0.04



94%|###|ETA: 0:00:07 Epoch: 1 Iter: 2898 Class Loss: 0.00 Loss: 0.03



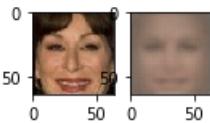
98%|###|ETA: 0:00:02 Epoch: 1 Iter: 2997 Class Loss: 0.00 Loss: 0.03



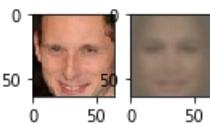
100%|####|Time: 0:02:30 Epoch: 1 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

N/A%| |ETA: --:--- Epoch: nan Iter: nan Class Loss: nan Loss: nan



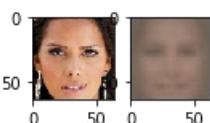
3%| |ETA: 0:02:24 Epoch: 2 Iter: 98 Class Loss: 0.00 Loss: 0.04



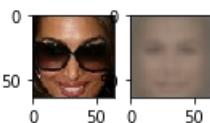
6%| |ETA: 0:02:18 Epoch: 2 Iter: 198 Class Loss: 0.00 Loss: 0.03



9%| |ETA: 0:02:17 Epoch: 2 Iter: 297 Class Loss: 0.00 Loss: 0.03



13%| |ETA: 0:02:11 Epoch: 2 Iter: 398 Class Loss: 0.00 Loss: 0.03



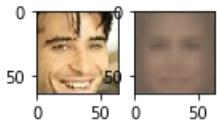
16%| |ETA: 0:02:04 Epoch: 2 Iter: 499 Class Loss: 0.01 Loss: 0.05



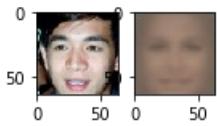
19% | ETA: 0:01:59 Epoch: 2 Iter: 599 Class Loss: 0.00 Loss: 0.03



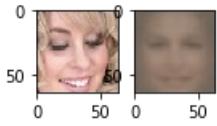
22% |# ETA: 0:01:55 Epoch: 2 Iter: 698 Class Loss: 0.00 Loss: 0.03



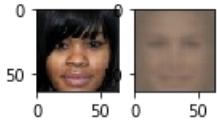
26% |# ETA: 0:01:51 Epoch: 2 Iter: 797 Class Loss: 0.08 Loss: 0.11



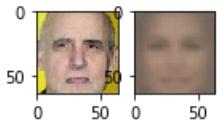
29% |# ETA: 0:01:45 Epoch: 2 Iter: 897 Class Loss: 0.00 Loss: 0.04



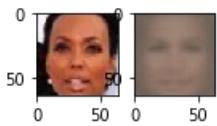
32% |# ETA: 0:02:02 Epoch: 2 Iter: 998 Class Loss: 0.00 Loss: 0.03



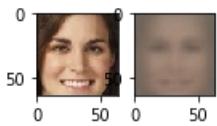
35% |# ETA: 0:01:38 Epoch: 2 Iter: 1098 Class Loss: 0.00 Loss: 0.04



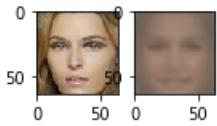
39% |# ETA: 0:01:46 Epoch: 2 Iter: 1198 Class Loss: 0.00 Loss: 0.03



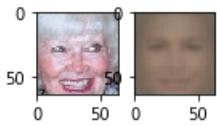
42% |# ETA: 0:01:26 Epoch: 2 Iter: 1299 Class Loss: 0.00 Loss: 0.03



45% |# ETA: 0:01:23 Epoch: 2 Iter: 1397 Class Loss: 0.00 Loss: 0.03



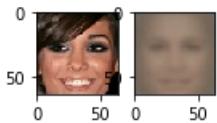
49% |# ETA: 0:01:14 Epoch: 2 Iter: 1499 Class Loss: 0.01 Loss: 0.05



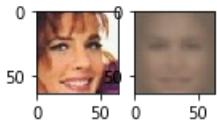
52% |## ETA: 0:01:12 Epoch: 2 Iter: 1599 Class Loss: 0.01 Loss: 0.03



55%|##| ETA: 0:01:07 Epoch: 2 Iter: 1697 Class Loss: 0.00 Loss: 0.03



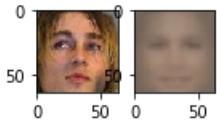
58%|##| ETA: 0:01:03 Epoch: 2 Iter: 1797 Class Loss: 0.01 Loss: 0.04



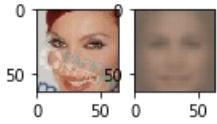
62%|##| ETA: 0:01:06 Epoch: 2 Iter: 1899 Class Loss: 0.00 Loss: 0.03



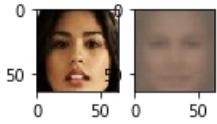
65%|##| ETA: 0:01:05 Epoch: 2 Iter: 1998 Class Loss: 0.00 Loss: 0.03



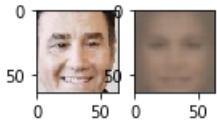
68%|##| ETA: 0:00:47 Epoch: 2 Iter: 2099 Class Loss: 0.03 Loss: 0.06



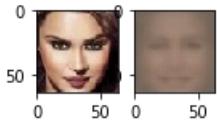
72%|##| ETA: 0:00:41 Epoch: 2 Iter: 2199 Class Loss: 0.00 Loss: 0.03



75%|###| ETA: 0:00:37 Epoch: 2 Iter: 2297 Class Loss: 0.00 Loss: 0.03



78%|###| ETA: 0:00:32 Epoch: 2 Iter: 2399 Class Loss: 0.00 Loss: 0.03



81%|###| ETA: 0:00:27 Epoch: 2 Iter: 2498 Class Loss: 0.00 Loss: 0.03



85%|###| ETA: 0:00:22 Epoch: 2 Iter: 2597 Class Loss: 0.00 Loss: 0.04



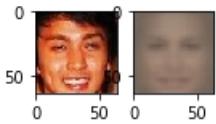
88%|###| ETA: 0:00:20 Epoch: 2 Iter: 2698 Class Loss: 0.00 Loss: 0.03



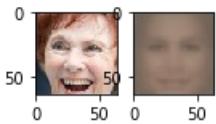
91%|###| ETA: 0:00:14 Epoch: 2 Iter: 2799 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:07 Epoch: 2 Iter: 2898 Class Loss: 0.00 Loss: 0.04



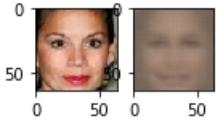
98%|###| ETA: 0:00:02 Epoch: 2 Iter: 2999 Class Loss: 0.01 Loss: 0.04



100%|####| Time: 0:02:37 Epoch: 2 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

N/A%| | ETA: --:--- Epoch: nan Iter: nan Class Loss: nan Loss: nan



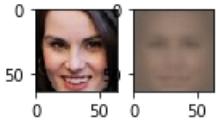
3%| | ETA: 0:02:26 Epoch: 3 Iter: 99 Class Loss: 0.00 Loss: 0.03



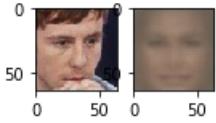
6%| | ETA: 0:02:21 Epoch: 3 Iter: 199 Class Loss: 0.00 Loss: 0.04



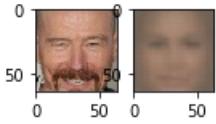
9%| | ETA: 0:02:18 Epoch: 3 Iter: 297 Class Loss: 0.00 Loss: 0.03



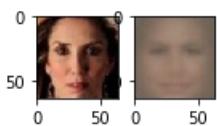
13%| | ETA: 0:02:14 Epoch: 3 Iter: 397 Class Loss: 0.00 Loss: 0.04



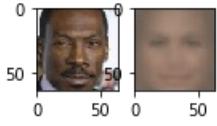
16%| | ETA: 0:02:06 Epoch: 3 Iter: 498 Class Loss: 0.00 Loss: 0.03



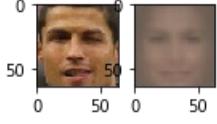
19%| | ETA: 0:02:02 Epoch: 3 Iter: 599 Class Loss: 0.01 Loss: 0.04



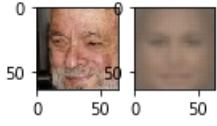
22%|# |ETA: 0:01:56 Epoch: 3 Iter: 698 Class Loss: 0.00 Loss: 0.03



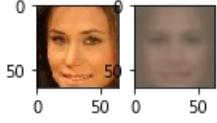
26%|# |ETA: 0:01:53 Epoch: 3 Iter: 797 Class Loss: 0.00 Loss: 0.04



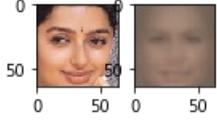
29%|# |ETA: 0:01:48 Epoch: 3 Iter: 897 Class Loss: 0.00 Loss: 0.03



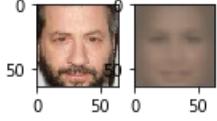
32%|# |ETA: 0:01:44 Epoch: 3 Iter: 997 Class Loss: 0.01 Loss: 0.04



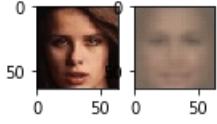
35%|# |ETA: 0:01:41 Epoch: 3 Iter: 1098 Class Loss: 0.00 Loss: 0.03



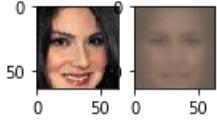
39%|# |ETA: 0:01:48 Epoch: 3 Iter: 1198 Class Loss: 0.00 Loss: 0.04



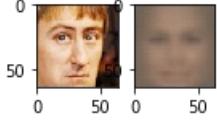
42%|# |ETA: 0:01:29 Epoch: 3 Iter: 1297 Class Loss: 0.00 Loss: 0.03



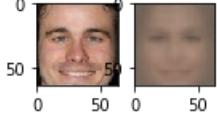
45%|# |ETA: 0:01:25 Epoch: 3 Iter: 1397 Class Loss: 0.00 Loss: 0.03



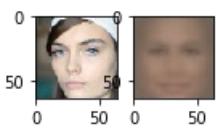
49%|# |ETA: 0:01:16 Epoch: 3 Iter: 1499 Class Loss: 0.00 Loss: 0.03



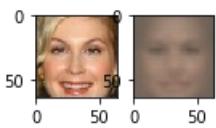
52%|## |ETA: 0:01:12 Epoch: 3 Iter: 1598 Class Loss: 0.00 Loss: 0.04



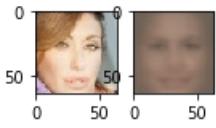
55%|## |ETA: 0:01:08 Epoch: 3 Iter: 1697 Class Loss: 0.09 Loss: 0.12



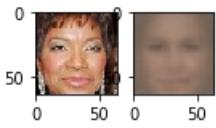
58%|##| ETA: 0:01:03 Epoch: 3 Iter: 1799 Class Loss: 0.00 Loss: 0.03



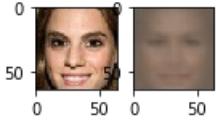
62%|##| ETA: 0:00:57 Epoch: 3 Iter: 1898 Class Loss: 0.00 Loss: 0.03



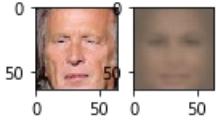
65%|##| ETA: 0:00:54 Epoch: 3 Iter: 1998 Class Loss: 0.14 Loss: 0.18



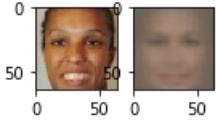
68%|##| ETA: 0:00:48 Epoch: 3 Iter: 2099 Class Loss: 0.00 Loss: 0.03



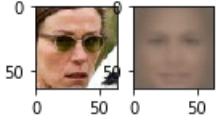
71%|##| ETA: 0:00:41 Epoch: 3 Iter: 2197 Class Loss: 0.00 Loss: 0.03



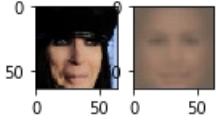
75%|###| ETA: 0:00:37 Epoch: 3 Iter: 2298 Class Loss: 0.00 Loss: 0.03



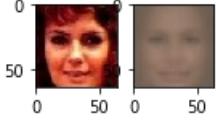
78%|###| ETA: 0:00:32 Epoch: 3 Iter: 2399 Class Loss: 0.00 Loss: 0.03



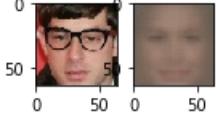
81%|###| ETA: 0:00:29 Epoch: 3 Iter: 2498 Class Loss: 0.00 Loss: 0.03



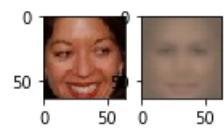
85%|###| ETA: 0:00:23 Epoch: 3 Iter: 2599 Class Loss: 0.01 Loss: 0.05



88%|###| ETA: 0:00:20 Epoch: 3 Iter: 2698 Class Loss: 0.00 Loss: 0.03



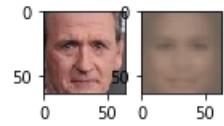
91%|###| ETA: 0:00:15 Epoch: 3 Iter: 2798 Class Loss: 0.00 Loss: 0.04



94%|###| ETA: 0:00:07 Epoch: 3 Iter: 2898 Class Loss: 0.00 Loss: 0.03



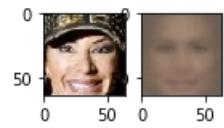
98%|###| ETA: 0:00:02 Epoch: 3 Iter: 2997 Class Loss: 0.00 Loss: 0.04



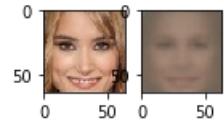
100%|####| Time: 0:02:37 Epoch: 3 Iter: 3052 Class Loss: 0.02 Loss: 0.05

Recomputing the sampling probabilities

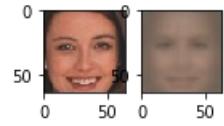
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



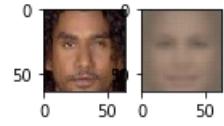
3%| | ETA: 0:02:26 Epoch: 4 Iter: 99 Class Loss: 0.00 Loss: 0.04



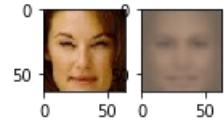
6%| | ETA: 0:02:20 Epoch: 4 Iter: 199 Class Loss: 0.02 Loss: 0.05



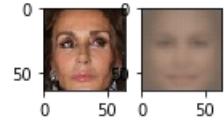
9%| | ETA: 0:02:17 Epoch: 4 Iter: 297 Class Loss: 0.00 Loss: 0.03



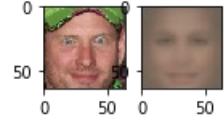
13%| | ETA: 0:02:11 Epoch: 4 Iter: 397 Class Loss: 0.01 Loss: 0.04



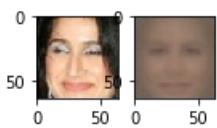
16%| | ETA: 0:02:07 Epoch: 4 Iter: 497 Class Loss: 0.03 Loss: 0.06



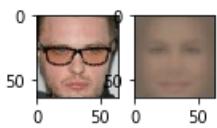
19%| | ETA: 0:02:05 Epoch: 4 Iter: 597 Class Loss: 0.00 Loss: 0.04



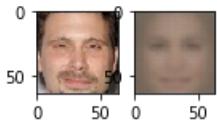
22%|#| ETA: 0:02:19 Epoch: 4 Iter: 699 Class Loss: 0.00 Loss: 0.03



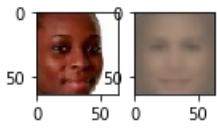
26%|# |ETA: 0:01:53 Epoch: 4 Iter: 798 Class Loss: 0.00 Loss: 0.03



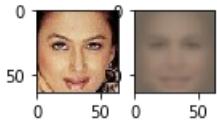
29%|# |ETA: 0:01:51 Epoch: 4 Iter: 897 Class Loss: 0.00 Loss: 0.03



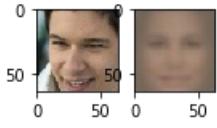
32%|# |ETA: 0:01:44 Epoch: 4 Iter: 998 Class Loss: 0.00 Loss: 0.03



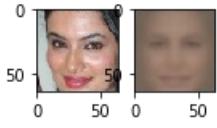
35%|# |ETA: 0:01:42 Epoch: 4 Iter: 1098 Class Loss: 0.00 Loss: 0.03



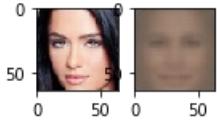
39%|# |ETA: 0:01:46 Epoch: 4 Iter: 1198 Class Loss: 0.05 Loss: 0.08



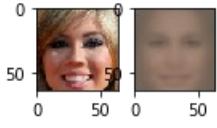
42%|# |ETA: 0:01:27 Epoch: 4 Iter: 1297 Class Loss: 0.00 Loss: 0.03



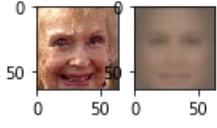
45%|# |ETA: 0:01:22 Epoch: 4 Iter: 1397 Class Loss: 0.00 Loss: 0.03



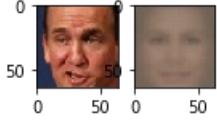
49%|# |ETA: 0:01:16 Epoch: 4 Iter: 1499 Class Loss: 0.00 Loss: 0.04



52%## |ETA: 0:01:12 Epoch: 4 Iter: 1598 Class Loss: 0.00 Loss: 0.03



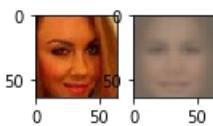
55%## |ETA: 0:01:22 Epoch: 4 Iter: 1699 Class Loss: 0.01 Loss: 0.04



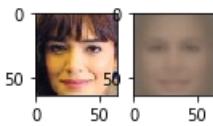
58%## |ETA: 0:01:16 Epoch: 4 Iter: 1799 Class Loss: 0.00 Loss: 0.03



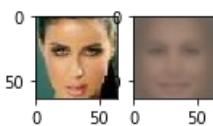
62%|##| ETA: 0:00:58 Epoch: 4 Iter: 1899 Class Loss: 0.00 Loss: 0.03



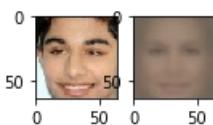
65%|##| ETA: 0:00:52 Epoch: 4 Iter: 1999 Class Loss: 0.00 Loss: 0.04



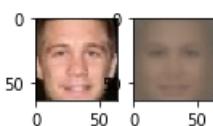
68%|##| ETA: 0:00:48 Epoch: 4 Iter: 2097 Class Loss: 0.03 Loss: 0.07



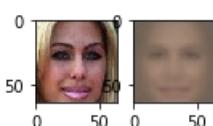
71%|##| ETA: 0:00:42 Epoch: 4 Iter: 2197 Class Loss: 0.00 Loss: 0.03



75%|###| ETA: 0:00:37 Epoch: 4 Iter: 2297 Class Loss: 0.00 Loss: 0.03



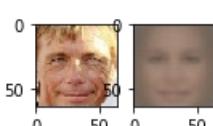
78%|###| ETA: 0:00:32 Epoch: 4 Iter: 2398 Class Loss: 0.00 Loss: 0.03



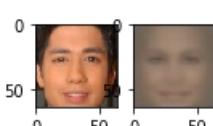
81%|###| ETA: 0:00:27 Epoch: 4 Iter: 2499 Class Loss: 0.00 Loss: 0.03



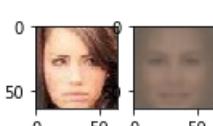
85%|###| ETA: 0:00:22 Epoch: 4 Iter: 2597 Class Loss: 0.00 Loss: 0.03



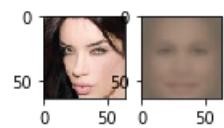
88%|###| ETA: 0:00:20 Epoch: 4 Iter: 2698 Class Loss: 0.00 Loss: 0.03



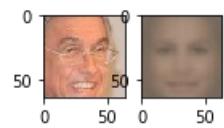
91%|###| ETA: 0:00:14 Epoch: 4 Iter: 2799 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:07 Epoch: 4 Iter: 2897 Class Loss: 0.00 Loss: 0.03



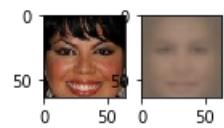
98%|###| ETA: 0:00:02 Epoch: 4 Iter: 2997 Class Loss: 0.00 Loss: 0.03



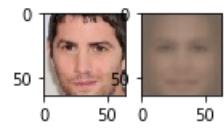
100%|####| Time: 0:02:39 Epoch: 4 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

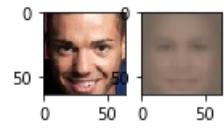
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



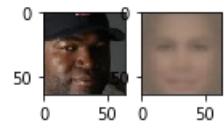
3%| | ETA: 0:02:29 Epoch: 5 Iter: 98 Class Loss: 0.00 Loss: 0.03



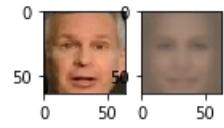
6%| | ETA: 0:02:24 Epoch: 5 Iter: 198 Class Loss: 0.00 Loss: 0.03



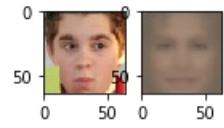
9%| | ETA: 0:02:17 Epoch: 5 Iter: 299 Class Loss: 0.00 Loss: 0.04



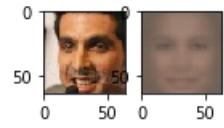
13%| | ETA: 0:02:17 Epoch: 5 Iter: 399 Class Loss: 0.00 Loss: 0.03



16%| | ETA: 0:02:06 Epoch: 5 Iter: 498 Class Loss: 0.05 Loss: 0.08



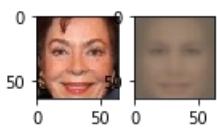
19%| | ETA: 0:02:03 Epoch: 5 Iter: 599 Class Loss: 0.00 Loss: 0.03



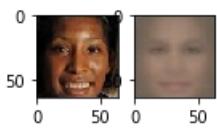
22%|#| ETA: 0:01:54 Epoch: 5 Iter: 697 Class Loss: 0.00 Loss: 0.03



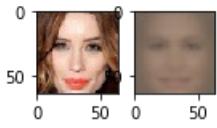
26%|#| ETA: 0:01:53 Epoch: 5 Iter: 798 Class Loss: 0.01 Loss: 0.04



29%|# |ETA: 0:01:47 Epoch: 5 Iter: 897 Class Loss: 0.00 Loss: 0.03



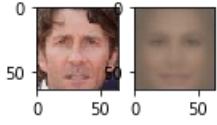
32%|# |ETA: 0:01:42 Epoch: 5 Iter: 998 Class Loss: 0.00 Loss: 0.04



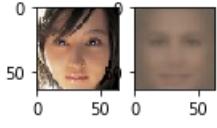
35%|# |ETA: 0:01:43 Epoch: 5 Iter: 1098 Class Loss: 0.00 Loss: 0.03



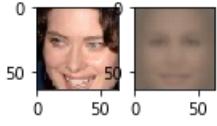
39%|# |ETA: 0:01:46 Epoch: 5 Iter: 1198 Class Loss: 0.00 Loss: 0.04



42%|# |ETA: 0:01:28 Epoch: 5 Iter: 1299 Class Loss: 0.00 Loss: 0.03



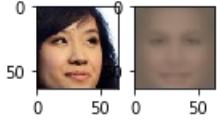
45%|# |ETA: 0:01:21 Epoch: 5 Iter: 1399 Class Loss: 0.00 Loss: 0.04



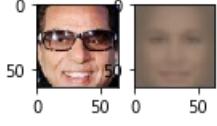
49%|# |ETA: 0:01:18 Epoch: 5 Iter: 1499 Class Loss: 0.00 Loss: 0.04



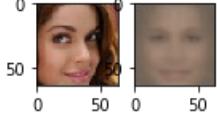
52%|## |ETA: 0:01:11 Epoch: 5 Iter: 1597 Class Loss: 0.00 Loss: 0.03



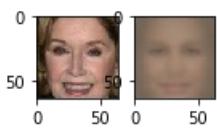
55%|## |ETA: 0:01:07 Epoch: 5 Iter: 1698 Class Loss: 0.00 Loss: 0.04



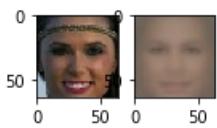
58%|## |ETA: 0:01:01 Epoch: 5 Iter: 1797 Class Loss: 0.00 Loss: 0.04



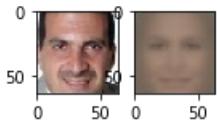
62%|## |ETA: 0:00:57 Epoch: 5 Iter: 1899 Class Loss: 0.00 Loss: 0.04



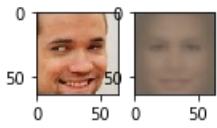
65%|##| ETA: 0:00:52 Epoch: 5 Iter: 1998 Class Loss: 0.00 Loss: 0.03



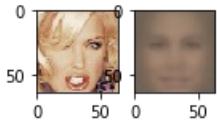
68%|##| ETA: 0:00:47 Epoch: 5 Iter: 2097 Class Loss: 0.03 Loss: 0.06



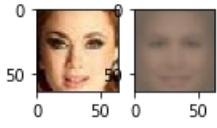
71%|##| ETA: 0:00:43 Epoch: 5 Iter: 2197 Class Loss: 0.00 Loss: 0.03



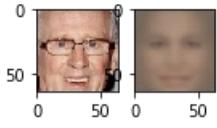
75%|###| ETA: 0:00:37 Epoch: 5 Iter: 2299 Class Loss: 0.00 Loss: 0.04



78%|###| ETA: 0:00:32 Epoch: 5 Iter: 2399 Class Loss: 0.00 Loss: 0.03



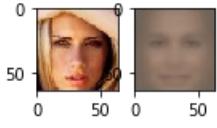
81%|###| ETA: 0:00:27 Epoch: 5 Iter: 2499 Class Loss: 0.00 Loss: 0.03



85%|###| ETA: 0:00:22 Epoch: 5 Iter: 2597 Class Loss: 0.00 Loss: 0.03



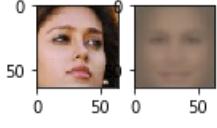
88%|###| ETA: 0:00:20 Epoch: 5 Iter: 2698 Class Loss: 0.00 Loss: 0.03



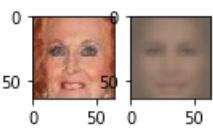
91%|###| ETA: 0:00:14 Epoch: 5 Iter: 2798 Class Loss: 0.00 Loss: 0.04



94%|###| ETA: 0:00:07 Epoch: 5 Iter: 2898 Class Loss: 0.00 Loss: 0.03



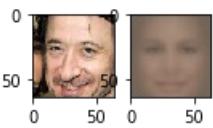
98%|###| ETA: 0:00:02 Epoch: 5 Iter: 2997 Class Loss: 0.00 Loss: 0.03



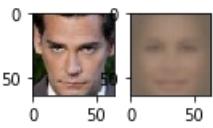
100%|####|Time: 0:02:36 Epoch: 5 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

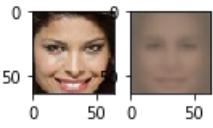
N/A%| |ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



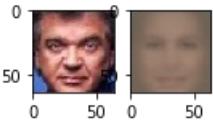
3%| |ETA: 0:02:28 Epoch: 6 Iter: 98 Class Loss: 0.00 Loss: 0.03



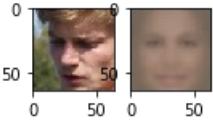
6%| |ETA: 0:02:20 Epoch: 6 Iter: 199 Class Loss: 0.00 Loss: 0.03



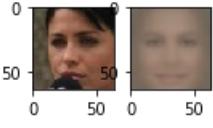
9%| |ETA: 0:02:20 Epoch: 6 Iter: 298 Class Loss: 0.00 Loss: 0.04



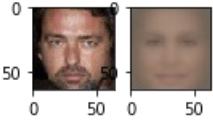
13%| |ETA: 0:02:38 Epoch: 6 Iter: 398 Class Loss: 0.00 Loss: 0.04



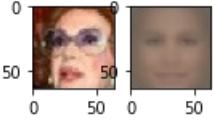
16%| |ETA: 0:02:04 Epoch: 6 Iter: 499 Class Loss: 0.00 Loss: 0.03



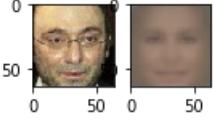
19%| |ETA: 0:02:03 Epoch: 6 Iter: 598 Class Loss: 0.00 Loss: 0.04



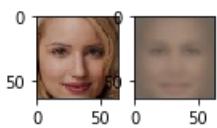
22%|# |ETA: 0:01:56 Epoch: 6 Iter: 698 Class Loss: 0.00 Loss: 0.03



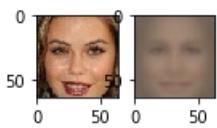
26%|# |ETA: 0:01:53 Epoch: 6 Iter: 798 Class Loss: 0.00 Loss: 0.04



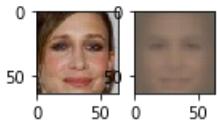
29%|# |ETA: 0:01:46 Epoch: 6 Iter: 897 Class Loss: 0.00 Loss: 0.03



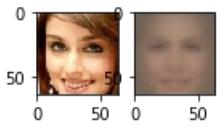
32%|# |ETA: 0:01:42 Epoch: 6 Iter: 999 Class Loss: 0.00 Loss: 0.03



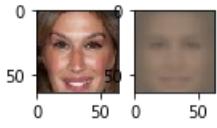
35%|# |ETA: 0:01:37 Epoch: 6 Iter: 1098 Class Loss: 0.00 Loss: 0.03



39%|# |ETA: 0:01:47 Epoch: 6 Iter: 1198 Class Loss: 0.00 Loss: 0.04



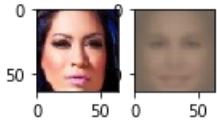
42%|# |ETA: 0:01:34 Epoch: 6 Iter: 1297 Class Loss: 0.00 Loss: 0.03



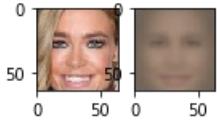
45%|# |ETA: 0:01:22 Epoch: 6 Iter: 1399 Class Loss: 0.00 Loss: 0.04



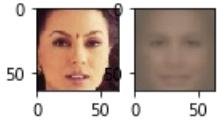
49%|# |ETA: 0:01:34 Epoch: 6 Iter: 1498 Class Loss: 0.00 Loss: 0.04



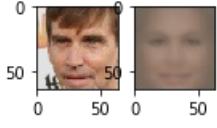
52%|## |ETA: 0:01:29 Epoch: 6 Iter: 1598 Class Loss: 0.00 Loss: 0.04



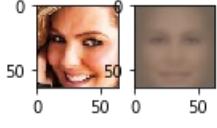
55%|## |ETA: 0:01:09 Epoch: 6 Iter: 1699 Class Loss: 0.00 Loss: 0.03



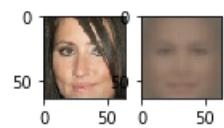
58%|## |ETA: 0:01:02 Epoch: 6 Iter: 1799 Class Loss: 0.00 Loss: 0.03



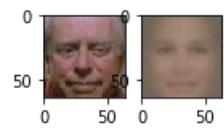
62%|## |ETA: 0:00:57 Epoch: 6 Iter: 1899 Class Loss: 0.00 Loss: 0.04



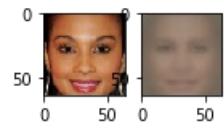
65%|## |ETA: 0:00:52 Epoch: 6 Iter: 1998 Class Loss: 0.00 Loss: 0.03



68%|##| ETA: 0:00:48 Epoch: 6 Iter: 2097 Class Loss: 0.00 Loss: 0.03



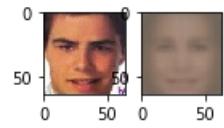
72%|##| ETA: 0:00:42 Epoch: 6 Iter: 2199 Class Loss: 0.00 Loss: 0.03



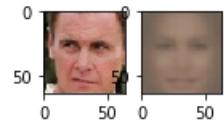
75%|###| ETA: 0:00:37 Epoch: 6 Iter: 2298 Class Loss: 0.00 Loss: 0.03



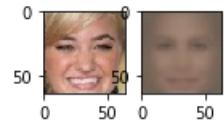
78%|###| ETA: 0:00:32 Epoch: 6 Iter: 2399 Class Loss: 0.00 Loss: 0.03



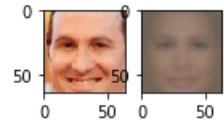
81%|###| ETA: 0:00:28 Epoch: 6 Iter: 2498 Class Loss: 0.00 Loss: 0.03



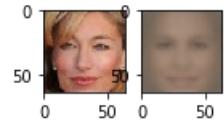
85%|###| ETA: 0:00:22 Epoch: 6 Iter: 2598 Class Loss: 0.00 Loss: 0.03



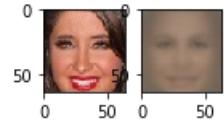
88%|###| ETA: 0:00:20 Epoch: 6 Iter: 2698 Class Loss: 0.00 Loss: 0.03



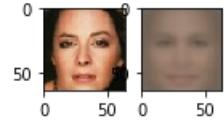
91%|###| ETA: 0:00:14 Epoch: 6 Iter: 2798 Class Loss: 0.05 Loss: 0.08



94%|###| ETA: 0:00:07 Epoch: 6 Iter: 2898 Class Loss: 0.00 Loss: 0.03



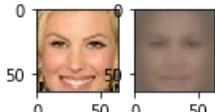
98%|###| ETA: 0:00:02 Epoch: 6 Iter: 2998 Class Loss: 0.00 Loss: 0.03



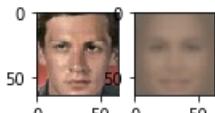
100%|####| Time: 0:02:39 Epoch: 6 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

N/A%| |ETA: --:--- Epoch: nan Iter: nan Class Loss: nan Loss: nan



3%| |ETA: 0:02:25 Epoch: 7 Iter: 98 Class Loss: 0.00 Loss: 0.03



6%| |ETA: 0:02:23 Epoch: 7 Iter: 199 Class Loss: 0.00 Loss: 0.03



9%| |ETA: 0:02:15 Epoch: 7 Iter: 299 Class Loss: 0.00 Loss: 0.04



13%| |ETA: 0:02:11 Epoch: 7 Iter: 398 Class Loss: 0.00 Loss: 0.03



16%| |ETA: 0:02:06 Epoch: 7 Iter: 498 Class Loss: 0.00 Loss: 0.03



19%| |ETA: 0:02:04 Epoch: 7 Iter: 599 Class Loss: 0.00 Loss: 0.04



22%|# |ETA: 0:02:00 Epoch: 7 Iter: 698 Class Loss: 0.00 Loss: 0.03



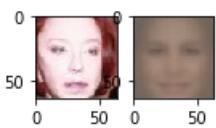
26%|# |ETA: 0:01:53 Epoch: 7 Iter: 798 Class Loss: 0.00 Loss: 0.03



29%|# |ETA: 0:01:45 Epoch: 7 Iter: 899 Class Loss: 0.00 Loss: 0.03



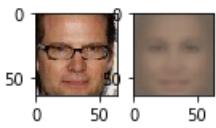
32%|# |ETA: 0:01:40 Epoch: 7 Iter: 998 Class Loss: 0.01 Loss: 0.04



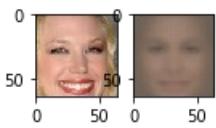
35%|# |ETA: 0:01:37 Epoch: 7 Iter: 1097 Class Loss: 0.00 Loss: 0.03



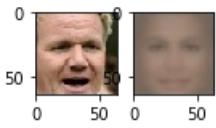
39%|# |ETA: 0:01:43 Epoch: 7 Iter: 1198 Class Loss: 0.00 Loss: 0.03



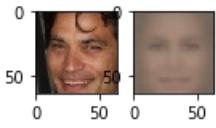
42%|# |ETA: 0:01:34 Epoch: 7 Iter: 1299 Class Loss: 0.00 Loss: 0.04



45%|# |ETA: 0:01:20 Epoch: 7 Iter: 1399 Class Loss: 0.00 Loss: 0.04



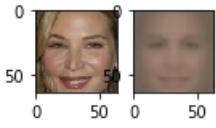
49%|# |ETA: 0:01:16 Epoch: 7 Iter: 1498 Class Loss: 0.00 Loss: 0.03



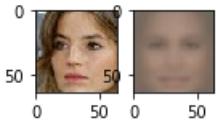
52%|## |ETA: 0:01:11 Epoch: 7 Iter: 1598 Class Loss: 0.00 Loss: 0.03



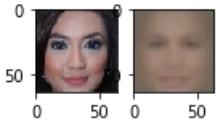
55%|## |ETA: 0:01:06 Epoch: 7 Iter: 1698 Class Loss: 0.00 Loss: 0.03



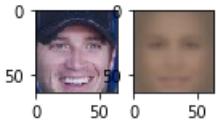
58%|## |ETA: 0:01:03 Epoch: 7 Iter: 1799 Class Loss: 0.00 Loss: 0.03



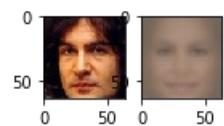
62%|## |ETA: 0:00:57 Epoch: 7 Iter: 1898 Class Loss: 0.00 Loss: 0.03



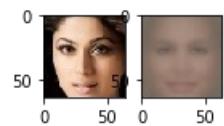
65%|## |ETA: 0:00:51 Epoch: 7 Iter: 1999 Class Loss: 0.02 Loss: 0.05



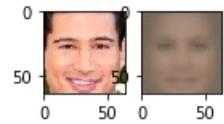
68%|## |ETA: 0:00:46 Epoch: 7 Iter: 2098 Class Loss: 0.00 Loss: 0.04



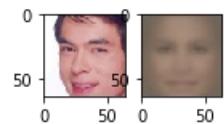
72%|##| ETA: 0:00:41 Epoch: 7 Iter: 2198 Class Loss: 0.00 Loss: 0.03



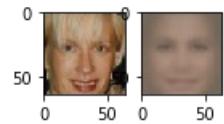
75%|###| ETA: 0:00:37 Epoch: 7 Iter: 2297 Class Loss: 0.00 Loss: 0.03



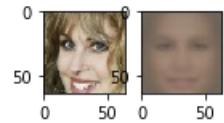
78%|###| ETA: 0:00:32 Epoch: 7 Iter: 2399 Class Loss: 0.00 Loss: 0.03



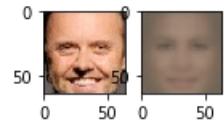
81%|###| ETA: 0:00:27 Epoch: 7 Iter: 2498 Class Loss: 0.00 Loss: 0.04



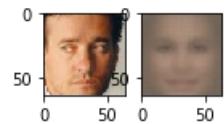
85%|###| ETA: 0:00:22 Epoch: 7 Iter: 2598 Class Loss: 0.00 Loss: 0.03



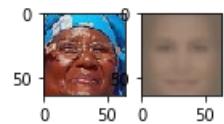
88%|###| ETA: 0:00:17 Epoch: 7 Iter: 2699 Class Loss: 0.01 Loss: 0.04



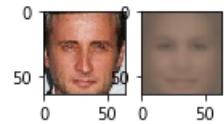
91%|###| ETA: 0:00:14 Epoch: 7 Iter: 2798 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:07 Epoch: 7 Iter: 2897 Class Loss: 0.00 Loss: 0.03



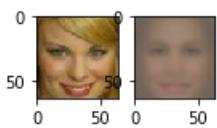
98%|###| ETA: 0:00:02 Epoch: 7 Iter: 2997 Class Loss: 0.00 Loss: 0.04



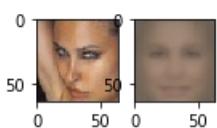
100%|####| Time: 0:02:35 Epoch: 7 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

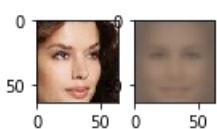
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



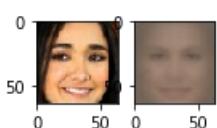
3% | ETA: 0:02:55 Epoch: 8 Iter: 99 Class Loss: 0.00 Loss: 0.03



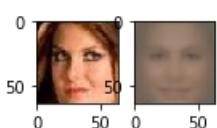
6% | ETA: 0:02:21 Epoch: 8 Iter: 199 Class Loss: 0.00 Loss: 0.04



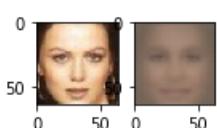
9% | ETA: 0:02:14 Epoch: 8 Iter: 298 Class Loss: 0.00 Loss: 0.03



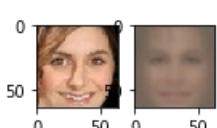
13% | ETA: 0:02:08 Epoch: 8 Iter: 398 Class Loss: 0.00 Loss: 0.03



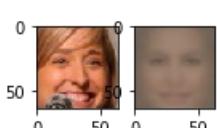
16% | ETA: 0:02:06 Epoch: 8 Iter: 498 Class Loss: 0.00 Loss: 0.03



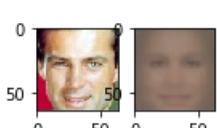
19% | ETA: 0:02:04 Epoch: 8 Iter: 597 Class Loss: 0.00 Loss: 0.03



22% | # ETA: 0:01:55 Epoch: 8 Iter: 697 Class Loss: 0.00 Loss: 0.03



26% | # ETA: 0:01:50 Epoch: 8 Iter: 798 Class Loss: 0.00 Loss: 0.04



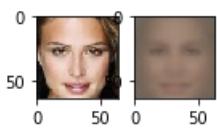
29% | # ETA: 0:01:47 Epoch: 8 Iter: 897 Class Loss: 0.00 Loss: 0.03



32% | # ETA: 0:01:40 Epoch: 8 Iter: 998 Class Loss: 0.00 Loss: 0.03



35% | # ETA: 0:01:37 Epoch: 8 Iter: 1097 Class Loss: 0.00 Loss: 0.03



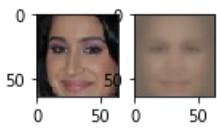
39%|# |ETA: 0:01:46 Epoch: 8 Iter: 1198 Class Loss: 0.00 Loss: 0.03



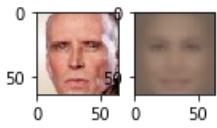
42%|# |ETA: 0:01:44 Epoch: 8 Iter: 1299 Class Loss: 0.00 Loss: 0.03



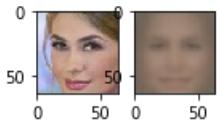
45%|# |ETA: 0:01:41 Epoch: 8 Iter: 1399 Class Loss: 0.00 Loss: 0.03



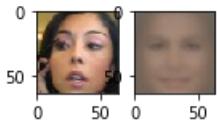
49%|# |ETA: 0:01:17 Epoch: 8 Iter: 1499 Class Loss: 0.00 Loss: 0.03



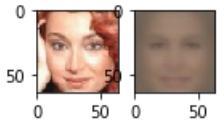
52%|## |ETA: 0:01:11 Epoch: 8 Iter: 1598 Class Loss: 0.00 Loss: 0.03



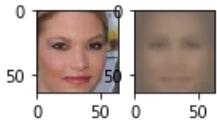
55%|## |ETA: 0:01:07 Epoch: 8 Iter: 1697 Class Loss: 0.00 Loss: 0.03



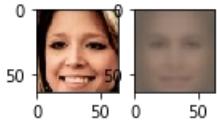
58%|## |ETA: 0:01:00 Epoch: 8 Iter: 1797 Class Loss: 0.00 Loss: 0.04



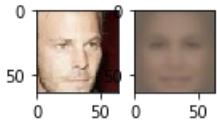
62%|## |ETA: 0:00:55 Epoch: 8 Iter: 1899 Class Loss: 0.00 Loss: 0.03



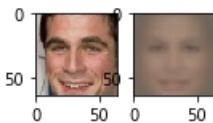
65%|## |ETA: 0:00:53 Epoch: 8 Iter: 1998 Class Loss: 0.00 Loss: 0.04



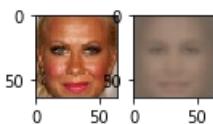
68%|## |ETA: 0:00:47 Epoch: 8 Iter: 2098 Class Loss: 0.00 Loss: 0.03



71%|## |ETA: 0:00:42 Epoch: 8 Iter: 2197 Class Loss: 0.00 Loss: 0.04



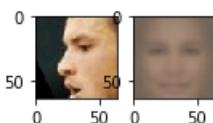
75%|###| ETA: 0:00:36 Epoch: 8 Iter: 2298 Class Loss: 0.00 Loss: 0.03



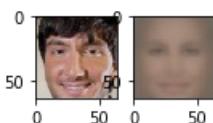
78%|###| ETA: 0:00:31 Epoch: 8 Iter: 2398 Class Loss: 0.00 Loss: 0.03



81%|###| ETA: 0:00:28 Epoch: 8 Iter: 2497 Class Loss: 0.00 Loss: 0.03



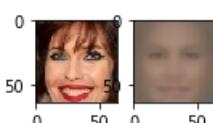
85%|###| ETA: 0:00:22 Epoch: 8 Iter: 2597 Class Loss: 0.00 Loss: 0.03



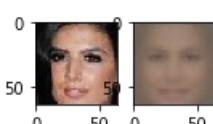
88%|###| ETA: 0:00:17 Epoch: 8 Iter: 2697 Class Loss: 0.00 Loss: 0.03



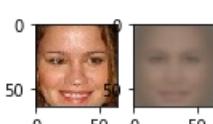
91%|###| ETA: 0:00:14 Epoch: 8 Iter: 2798 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:08 Epoch: 8 Iter: 2898 Class Loss: 0.00 Loss: 0.04



98%|###| ETA: 0:00:02 Epoch: 8 Iter: 2998 Class Loss: 0.00 Loss: 0.03



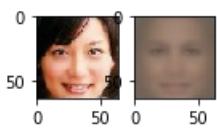
100%|####| Time: 0:02:37 Epoch: 8 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

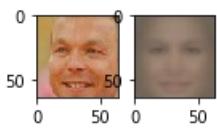
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



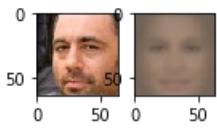
3%| | ETA: 0:02:25 Epoch: 9 Iter: 98 Class Loss: 0.00 Loss: 0.04



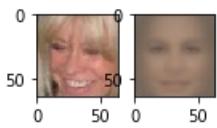
6% | ETA: 0:02:22 Epoch: 9 Iter: 198 Class Loss: 0.00 Loss: 0.03



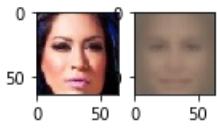
9% | ETA: 0:02:18 Epoch: 9 Iter: 299 Class Loss: 0.00 Loss: 0.03



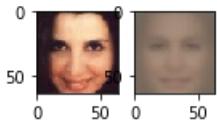
13% | ETA: 0:02:13 Epoch: 9 Iter: 397 Class Loss: 0.00 Loss: 0.03



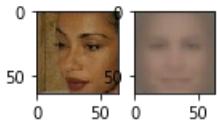
16% | ETA: 0:02:07 Epoch: 9 Iter: 499 Class Loss: 0.00 Loss: 0.03



19% | ETA: 0:02:02 Epoch: 9 Iter: 599 Class Loss: 0.00 Loss: 0.03



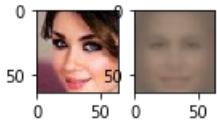
22% | # ETA: 0:01:59 Epoch: 9 Iter: 697 Class Loss: 0.00 Loss: 0.03



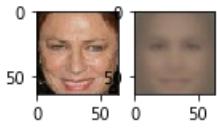
26% | # ETA: 0:01:53 Epoch: 9 Iter: 797 Class Loss: 0.00 Loss: 0.03



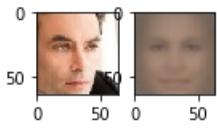
29% | # ETA: 0:01:46 Epoch: 9 Iter: 898 Class Loss: 0.00 Loss: 0.03



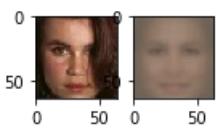
32% | # ETA: 0:01:40 Epoch: 9 Iter: 997 Class Loss: 0.00 Loss: 0.03



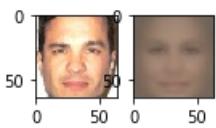
35% | # ETA: 0:01:37 Epoch: 9 Iter: 1098 Class Loss: 0.00 Loss: 0.03



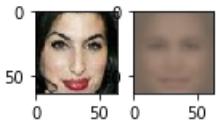
39% | # ETA: 0:01:36 Epoch: 9 Iter: 1198 Class Loss: 0.00 Loss: 0.03



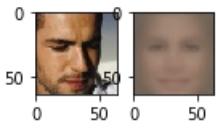
42%|# |ETA: 0:01:40 Epoch: 9 Iter: 1299 Class Loss: 0.00 Loss: 0.03



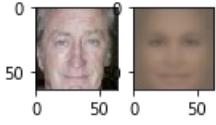
45%|# |ETA: 0:01:23 Epoch: 9 Iter: 1399 Class Loss: 0.00 Loss: 0.03



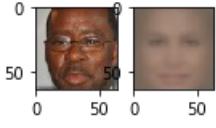
49%|# |ETA: 0:01:16 Epoch: 9 Iter: 1498 Class Loss: 0.00 Loss: 0.03



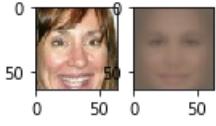
52%|## |ETA: 0:01:12 Epoch: 9 Iter: 1598 Class Loss: 0.00 Loss: 0.03



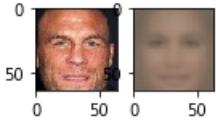
55%|## |ETA: 0:01:06 Epoch: 9 Iter: 1699 Class Loss: 0.00 Loss: 0.04



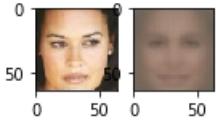
58%|## |ETA: 0:01:01 Epoch: 9 Iter: 1799 Class Loss: 0.00 Loss: 0.03



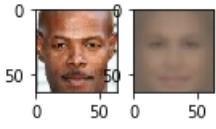
62%|## |ETA: 0:00:56 Epoch: 9 Iter: 1898 Class Loss: 0.00 Loss: 0.03



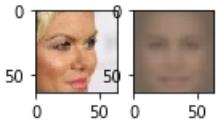
65%|## |ETA: 0:00:52 Epoch: 9 Iter: 1997 Class Loss: 0.00 Loss: 0.04



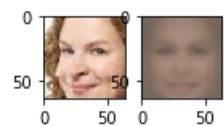
68%|## |ETA: 0:00:47 Epoch: 9 Iter: 2098 Class Loss: 0.00 Loss: 0.03



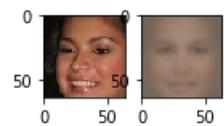
72%|## |ETA: 0:00:41 Epoch: 9 Iter: 2199 Class Loss: 0.00 Loss: 0.03



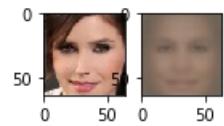
75%|### |ETA: 0:00:37 Epoch: 9 Iter: 2297 Class Loss: 0.00 Loss: 0.03



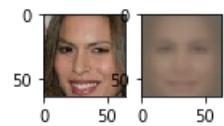
78%|###|ETA: 0:00:33 Epoch: 9 Iter: 2398 Class Loss: 0.00 Loss: 0.03



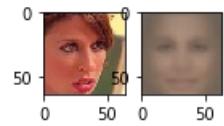
81%|###|ETA: 0:00:27 Epoch: 9 Iter: 2499 Class Loss: 0.00 Loss: 0.03



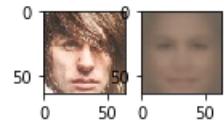
85%|###|ETA: 0:00:21 Epoch: 9 Iter: 2598 Class Loss: 0.00 Loss: 0.04



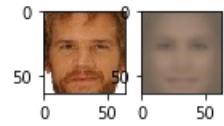
88%|###|ETA: 0:00:17 Epoch: 9 Iter: 2698 Class Loss: 0.00 Loss: 0.03



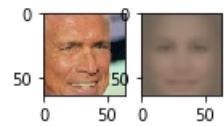
91%|###|ETA: 0:00:14 Epoch: 9 Iter: 2798 Class Loss: 0.00 Loss: 0.03



94%|###|ETA: 0:00:09 Epoch: 9 Iter: 2898 Class Loss: 0.00 Loss: 0.04



98%|###|ETA: 0:00:02 Epoch: 9 Iter: 2997 Class Loss: 0.00 Loss: 0.03



100%|####|Time: 0:02:36 Epoch: 9 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Debiased Evaluation

In [59]:

```
# Evaluate on PPB dataset (takes ~4 minutes)
accuracy_debiased = []
for skin_color in ['lighter', 'darker']:
    for gender in ['male', 'female']:
        accuracy_debiased.append( ppb.evaluate([encoder], gender, skin_color, output_idx=0, from_logit=True)[0] )
    print
print "{} {}: {}".format(gender, skin_color, accuracy_debiased[-1])

bar_width = 0.3
plt.bar(np.arange(4), standard_cnn_accuracy, width=bar_width)
plt.bar(np.arange(4)+bar_width, accuracy_debiased, width=bar_width)
plt.legend(['Standard Classifier', 'Debiased Classifier (DB-VAE)'])
plt.xticks(np.arange(4), ('LM', 'LF', 'DM', 'DF'))
plt.ylim(np.min([standard_cnn_accuracy,accuracy_debiased])-0.1,1)
plt.ylabel('Accuracy')
```

```
100% (97 of 97) |#####| Elapsed Time: 0:01:07 Time: 0:01:07
N/A% (0 of 72) | Elapsed Time: 0:00:00 ETA: ---:--
```

```
male lighter: 0.969072164948
```

```
100% (72 of 72) |#####| Elapsed Time: 0:00:49 Time: 0:00:49
N/A% (0 of 78) | Elapsed Time: 0:00:00 ETA: ---:--
```

```
female lighter: 1.0
```

```
100% (78 of 78) |#####| Elapsed Time: 0:00:52 Time: 0:00:52
N/A% (0 of 71) | Elapsed Time: 0:00:00 ETA: ---:--
```

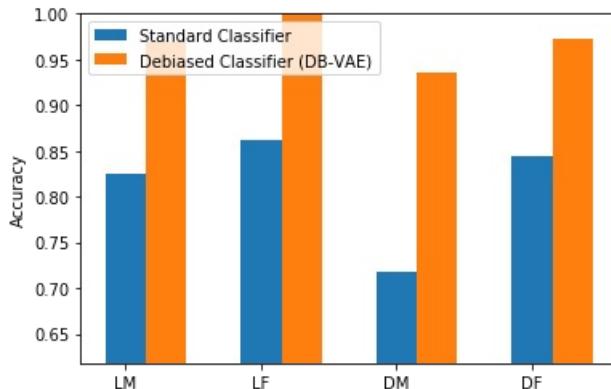
```
male darker: 0.935897435897
```

```
100% (71 of 71) |#####| Elapsed Time: 0:00:48 Time: 0:00:48
```

```
female darker: 0.971830985915
```

Out[59]:

```
Text(0,0.5,'Accuracy')
```



We can notice how much closer the different genders and skin tone have become, and how our accuracy increased for all.

Trying Different Smoothing Factors

In [0]:

```
def run_smooth_factor(factor_smooth):
    encoder, inputs, outputs, pre_flatten_shape = make_face_encoder_network()
    decoder = make_face_decoder_network(pre_flatten_shape)
    # initialize the models
    encoder_output = encoder(inputs)
    y_logit, z_mean, z_logsigma, z = encoder_output
    reconstructed_inputs = decoder(z)

    vae = tf.keras.Model(inputs, reconstructed_inputs)
    loss_history = []

    batch_size = 36
    num_epochs = 10 # keep small to run faster
    learning_rate = 1e-3

    optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate)
```

```

enable_debiasing = True
all_faces = loader.get_all_train_faces() # parameter from data loader

for epoch in range(num_epochs):

    # progress message and bar
    custom_msg = util.custom_progress_text("Epoch: %(epoch).0f Iter: %(idx).0f Class Loss: %(class_loss)2.2f Loss: %(loss)2.2f")
    bar = util.create_progress_bar(custom_msg)

    p_faces = None
    if enable_debiasing:
        # Recompute data sampling probabilities if debiasing is enabled
        '''TODO: write the function call to recompute the sampling probabilities when debiasing is enabled'''
        p_faces = get_training_sample_probabilities(all_faces, encoder, bins=10, smoothing_fac=factor_smooth)
    for idx in bar(range(loader.get_train_size()//batch_size)):
        # load a batch of data
        (x, y) = loader.get_batch(batch_size, p_pos=p_faces)
        x = tf.convert_to_tensor(x, dtype=tf.float32)
        y = tf.convert_to_tensor(y, dtype=tf.float32)

        # define GradientTape for automatic differentiation
        with tf.GradientTape() as tape:
            y_logit, mu, logsigma, z = encoder(x)
            x_hat = decoder(z)
            '''TODO: call the relevant loss function to compute the loss'''
            loss, class_loss = debiasing_loss_function(x, x_hat, y, y_logit, mu, logsigma)
            '''TODO: use the GradientTape.gradient method to compute the gradients'''
            grads = tape.gradient(loss, vae.variables)
        # apply gradients to variables
        optimizer.apply_gradients(zip(grads, vae.variables),
                                  global_step=tf.train.get_or_create_global_step())

        # track the losses
        class_loss_value = class_loss.numpy().mean()
        loss_value = loss.numpy().mean()
        loss_history.append((class_loss_value, loss_value))
        custom_msg.update_mapping(epoch=epoch, idx=idx, loss=loss_value, class_loss=class_loss_value)

    # plot the progress every 100 steps
    if idx%100 == 0:
        util.plot_sample(x,y,vae)
# Evaluate on PPB dataset (takes ~4 minutes)
accuracy_debiased = []
for skin_color in ['lighter', 'darker']:
    for gender in ['male', 'female']:
        accuracy_debiased.append(ppb.evaluate([encoder], gender, skin_color, output_idx=0, from_logit=True)[0])
print
print "{} {}: {}".format(gender, skin_color, accuracy_debiased[-1])

bar_width = 0.3
plt.bar(np.arange(4), standard_cnn_accuracy, width=bar_width)
plt.bar(np.arange(4)+bar_width, accuracy_debiased, width=bar_width)
plt.legend(['Standard Classifier', 'Debiased Classifier (DB-VAE)'])
plt.xticks(np.arange(4), ('LM', 'LF', 'DM', 'DF'))
plt.ylim(np.min([standard_cnn_accuracy,accuracy_debiased])-0.1,1)
plt.ylabel('Accuracy')

```

Smooth Factor of 0.2

In [61]:

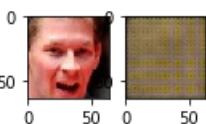
```
run_smooth_factor(0.2)
```

Recomputing the sampling probabilities

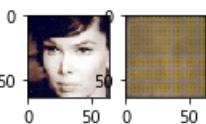
```
N/A%|      |ETA:  --- Epoch: nan Iter: nan Class Loss: nan Loss: nan
```



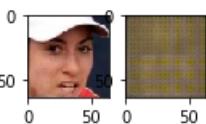
3%| ETA: 0:02:19 Epoch: 0 Iter: 99 Class Loss: 0.12 Loss: 0.16



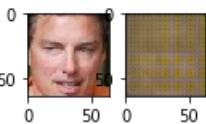
6%| ETA: 0:02:15 Epoch: 0 Iter: 199 Class Loss: 0.27 Loss: 0.31



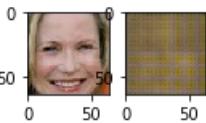
9%| ETA: 0:02:10 Epoch: 0 Iter: 299 Class Loss: 0.08 Loss: 0.12



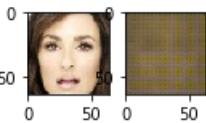
13%| ETA: 0:02:06 Epoch: 0 Iter: 399 Class Loss: 0.15 Loss: 0.19



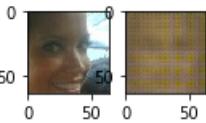
16%| ETA: 0:01:58 Epoch: 0 Iter: 498 Class Loss: 0.06 Loss: 0.10



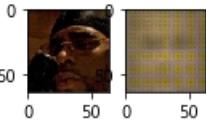
19%| ETA: 0:01:57 Epoch: 0 Iter: 598 Class Loss: 0.03 Loss: 0.07



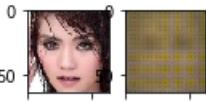
22%|# ETA: 0:02:06 Epoch: 0 Iter: 699 Class Loss: 0.01 Loss: 0.05



26%|# ETA: 0:01:51 Epoch: 0 Iter: 798 Class Loss: 0.01 Loss: 0.05



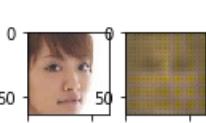
29%|# ETA: 0:01:41 Epoch: 0 Iter: 897 Class Loss: 0.00 Loss: 0.04



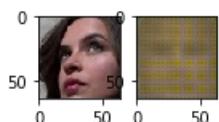
32%|# ETA: 0:01:36 Epoch: 0 Iter: 998 Class Loss: 0.04 Loss: 0.08



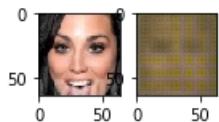
35%|# ETA: 0:01:32 Epoch: 0 Iter: 1097 Class Loss: 0.02 Loss: 0.06



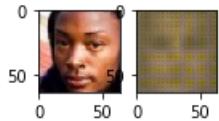
39%|# |ETA: 0:01:27 Epoch: 0 Iter: 1197 Class Loss: 0.04 Loss: 0.09



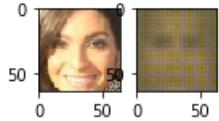
42%|# |ETA: 0:01:24 Epoch: 0 Iter: 1298 Class Loss: 0.07 Loss: 0.11



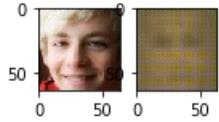
45%|# |ETA: 0:01:18 Epoch: 0 Iter: 1399 Class Loss: 0.33 Loss: 0.37



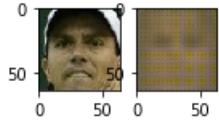
49%|# |ETA: 0:01:15 Epoch: 0 Iter: 1497 Class Loss: 0.02 Loss: 0.07



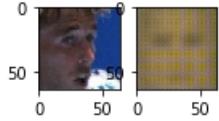
52%|## |ETA: 0:01:08 Epoch: 0 Iter: 1599 Class Loss: 0.00 Loss: 0.05



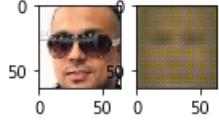
55%|## |ETA: 0:01:14 Epoch: 0 Iter: 1698 Class Loss: 0.00 Loss: 0.04



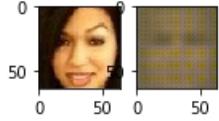
58%|## |ETA: 0:01:10 Epoch: 0 Iter: 1799 Class Loss: 0.01 Loss: 0.05



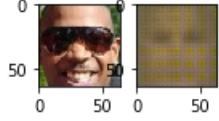
62%|## |ETA: 0:00:55 Epoch: 0 Iter: 1899 Class Loss: 0.05 Loss: 0.09



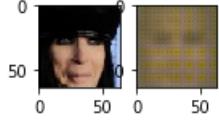
65%|## |ETA: 0:00:49 Epoch: 0 Iter: 1998 Class Loss: 0.01 Loss: 0.06



68%|## |ETA: 0:00:44 Epoch: 0 Iter: 2098 Class Loss: 0.01 Loss: 0.04



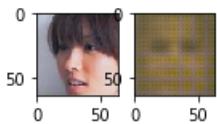
72%|## |ETA: 0:00:42 Epoch: 0 Iter: 2198 Class Loss: 0.00 Loss: 0.05



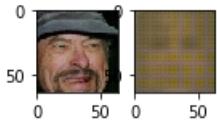
75%|### |ETA: 0:00:40 Epoch: 0 Iter: 2298 Class Loss: 0.00 Loss: 0.05



78%|### |ETA: 0:00:31 Epoch: 0 Iter: 2398 Class Loss: 0.13 Loss: 0.17



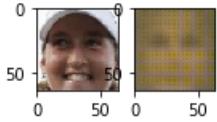
81%|### |ETA: 0:00:26 Epoch: 0 Iter: 2497 Class Loss: 0.00 Loss: 0.05



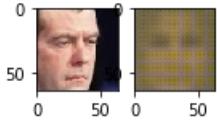
85%|### |ETA: 0:00:21 Epoch: 0 Iter: 2597 Class Loss: 0.00 Loss: 0.04



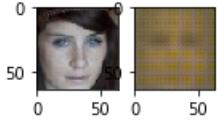
88%|### |ETA: 0:00:16 Epoch: 0 Iter: 2698 Class Loss: 0.02 Loss: 0.06



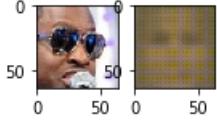
91%|### |ETA: 0:00:12 Epoch: 0 Iter: 2797 Class Loss: 0.00 Loss: 0.05



94%|### |ETA: 0:00:07 Epoch: 0 Iter: 2898 Class Loss: 0.00 Loss: 0.05



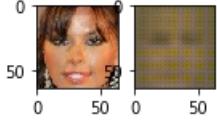
98%|### |ETA: 0:00:02 Epoch: 0 Iter: 2999 Class Loss: 0.01 Loss: 0.05



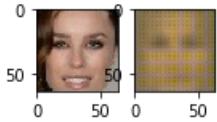
100%|####|Time: 0:02:32 Epoch: 0 Iter: 3052 Class Loss: 0.00 Loss: 0.05

Recomputing the sampling probabilities

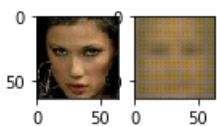
N/A%| |ETA: ---- Epoch: nan Iter: nan Class Loss: nan Loss: nan



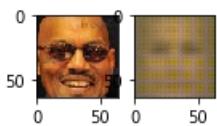
3%| |ETA: 0:02:17 Epoch: 1 Iter: 97 Class Loss: 0.01 Loss: 0.06



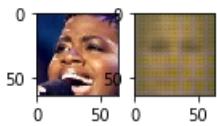
6%| |ETA: 0:02:14 Epoch: 1 Iter: 199 Class Loss: 0.03 Loss: 0.07



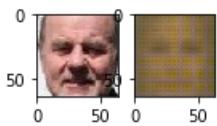
9% | ETA: 0:02:11 Epoch: 1 Iter: 297 Class Loss: 0.01 Loss: 0.05



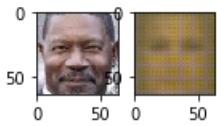
13% | ETA: 0:02:07 Epoch: 1 Iter: 397 Class Loss: 0.00 Loss: 0.04



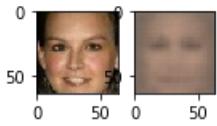
16% | ETA: 0:02:02 Epoch: 1 Iter: 497 Class Loss: 0.09 Loss: 0.12



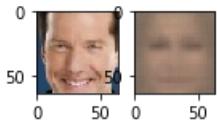
19% | ETA: 0:01:56 Epoch: 1 Iter: 599 Class Loss: 0.00 Loss: 0.04



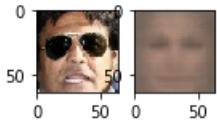
22% | # ETA: 0:01:54 Epoch: 1 Iter: 699 Class Loss: 0.00 Loss: 0.03



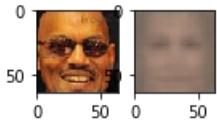
26% | # ETA: 0:02:03 Epoch: 1 Iter: 798 Class Loss: 0.00 Loss: 0.03



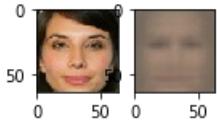
29% | # ETA: 0:01:46 Epoch: 1 Iter: 899 Class Loss: 0.00 Loss: 0.03



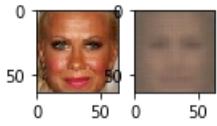
32% | # ETA: 0:01:36 Epoch: 1 Iter: 999 Class Loss: 0.01 Loss: 0.04



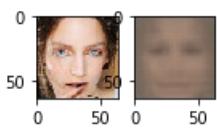
35% | # ETA: 0:01:33 Epoch: 1 Iter: 1097 Class Loss: 0.00 Loss: 0.03



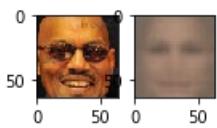
39% | # ETA: 0:01:27 Epoch: 1 Iter: 1198 Class Loss: 0.00 Loss: 0.04



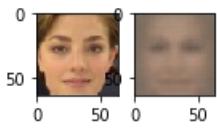
42% | # ETA: 0:01:23 Epoch: 1 Iter: 1297 Class Loss: 0.00 Loss: 0.03



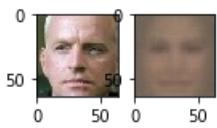
45%|# |ETA: 0:01:19 Epoch: 1 Iter: 1399 Class Loss: 0.09 Loss: 0.12



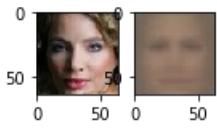
49%|# |ETA: 0:01:15 Epoch: 1 Iter: 1498 Class Loss: 0.00 Loss: 0.03



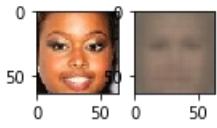
52%|## |ETA: 0:01:09 Epoch: 1 Iter: 1598 Class Loss: 0.00 Loss: 0.03



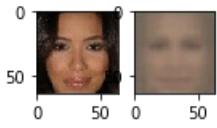
55%|## |ETA: 0:01:04 Epoch: 1 Iter: 1699 Class Loss: 0.01 Loss: 0.05



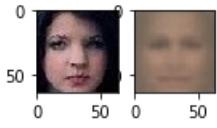
58%|## |ETA: 0:01:00 Epoch: 1 Iter: 1799 Class Loss: 0.00 Loss: 0.04



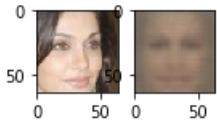
62%|## |ETA: 0:00:54 Epoch: 1 Iter: 1898 Class Loss: 0.00 Loss: 0.03



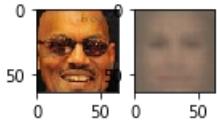
65%|## |ETA: 0:00:50 Epoch: 1 Iter: 1999 Class Loss: 0.01 Loss: 0.04



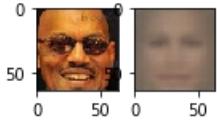
68%|## |ETA: 0:00:45 Epoch: 1 Iter: 2098 Class Loss: 0.00 Loss: 0.03



72%|## |ETA: 0:00:40 Epoch: 1 Iter: 2199 Class Loss: 0.00 Loss: 0.03



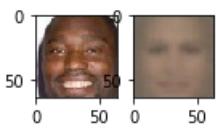
75%|### |ETA: 0:00:36 Epoch: 1 Iter: 2297 Class Loss: 0.07 Loss: 0.10



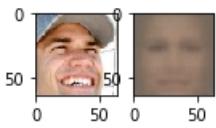
78%|### |ETA: 0:00:36 Epoch: 1 Iter: 2398 Class Loss: 0.00 Loss: 0.03



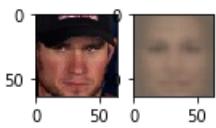
81%|###|ETA: 0:00:28 Epoch: 1 Iter: 2499 Class Loss: 0.00 Loss: 0.03



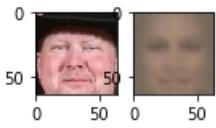
85%|###|ETA: 0:00:21 Epoch: 1 Iter: 2598 Class Loss: 0.00 Loss: 0.03



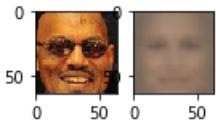
88%|###|ETA: 0:00:16 Epoch: 1 Iter: 2699 Class Loss: 0.04 Loss: 0.06



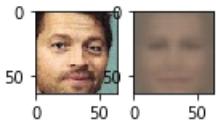
91%|###|ETA: 0:00:12 Epoch: 1 Iter: 2799 Class Loss: 0.00 Loss: 0.04



94%|###|ETA: 0:00:07 Epoch: 1 Iter: 2898 Class Loss: 0.00 Loss: 0.03



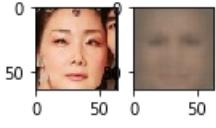
98%|###|ETA: 0:00:02 Epoch: 1 Iter: 2999 Class Loss: 0.00 Loss: 0.03



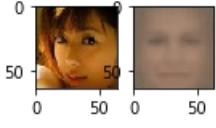
100%|####|Time: 0:02:32 Epoch: 1 Iter: 3052 Class Loss: 0.03 Loss: 0.06

Recomputing the sampling probabilities

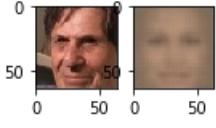
N/A%| |ETA: ---- Epoch: nan Iter: nan Class Loss: nan Loss: nan



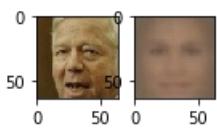
3%| |ETA: 0:02:19 Epoch: 2 Iter: 97 Class Loss: 0.00 Loss: 0.03



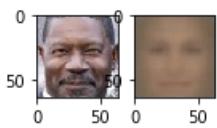
6%| |ETA: 0:02:16 Epoch: 2 Iter: 199 Class Loss: 0.00 Loss: 0.04



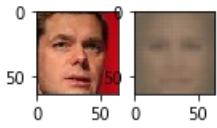
9%| |ETA: 0:02:11 Epoch: 2 Iter: 298 Class Loss: 0.00 Loss: 0.03



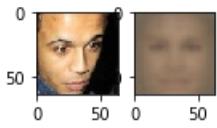
13% | ETA: 0:02:07 Epoch: 2 Iter: 399 Class Loss: 0.00 Loss: 0.03



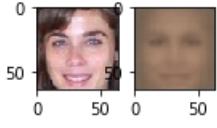
16% | ETA: 0:01:59 Epoch: 2 Iter: 499 Class Loss: 0.00 Loss: 0.03



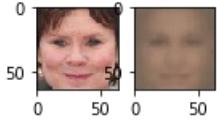
19% | ETA: 0:01:55 Epoch: 2 Iter: 599 Class Loss: 0.04 Loss: 0.07



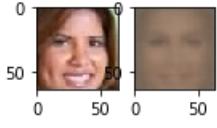
22% | # ETA: 0:01:50 Epoch: 2 Iter: 698 Class Loss: 0.00 Loss: 0.03



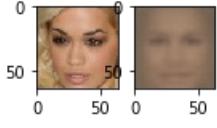
26% | # ETA: 0:01:47 Epoch: 2 Iter: 797 Class Loss: 0.00 Loss: 0.03



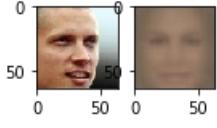
29% | # ETA: 0:01:57 Epoch: 2 Iter: 898 Class Loss: 0.00 Loss: 0.04



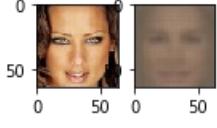
32% | # ETA: 0:01:49 Epoch: 2 Iter: 999 Class Loss: 0.00 Loss: 0.03



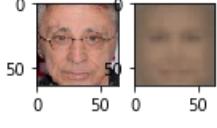
36% | # ETA: 0:01:32 Epoch: 2 Iter: 1099 Class Loss: 0.00 Loss: 0.03



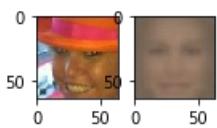
39% | # ETA: 0:01:27 Epoch: 2 Iter: 1198 Class Loss: 0.00 Loss: 0.04



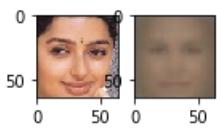
42% | # ETA: 0:01:22 Epoch: 2 Iter: 1299 Class Loss: 0.00 Loss: 0.03



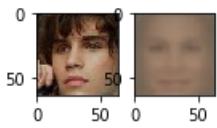
45% | # ETA: 0:01:18 Epoch: 2 Iter: 1397 Class Loss: 0.00 Loss: 0.04



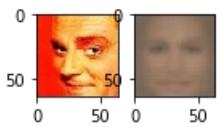
49%|#| ETA: 0:01:14 Epoch: 2 Iter: 1499 Class Loss: 0.00 Loss: 0.03



52%|##| ETA: 0:01:08 Epoch: 2 Iter: 1599 Class Loss: 0.00 Loss: 0.04



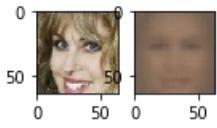
55%|##| ETA: 0:01:16 Epoch: 2 Iter: 1699 Class Loss: 0.01 Loss: 0.04



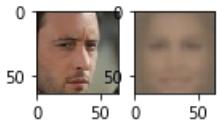
58%|##| ETA: 0:01:11 Epoch: 2 Iter: 1799 Class Loss: 0.07 Loss: 0.10



62%|##| ETA: 0:00:54 Epoch: 2 Iter: 1899 Class Loss: 0.00 Loss: 0.03



65%|##| ETA: 0:00:50 Epoch: 2 Iter: 1998 Class Loss: 0.00 Loss: 0.04



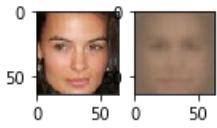
68%|##| ETA: 0:00:46 Epoch: 2 Iter: 2099 Class Loss: 0.00 Loss: 0.03



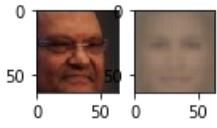
71%|##| ETA: 0:00:40 Epoch: 2 Iter: 2197 Class Loss: 0.00 Loss: 0.04



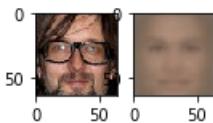
75%|###| ETA: 0:00:35 Epoch: 2 Iter: 2298 Class Loss: 0.00 Loss: 0.03



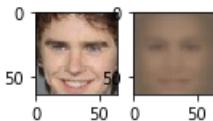
78%|###| ETA: 0:00:32 Epoch: 2 Iter: 2398 Class Loss: 0.00 Loss: 0.03



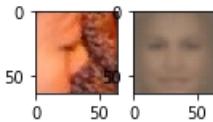
81%|###| ETA: 0:00:29 Epoch: 2 Iter: 2499 Class Loss: 0.00 Loss: 0.04



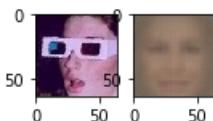
85%|###|ETA: 0:00:22 Epoch: 2 Iter: 2597 Class Loss: 0.00 Loss: 0.03



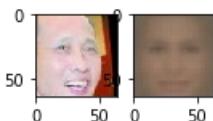
88%|###|ETA: 0:00:16 Epoch: 2 Iter: 2698 Class Loss: 0.00 Loss: 0.03



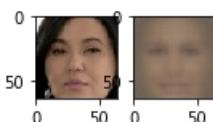
91%|###|ETA: 0:00:12 Epoch: 2 Iter: 2798 Class Loss: 0.01 Loss: 0.04



94%|###|ETA: 0:00:07 Epoch: 2 Iter: 2898 Class Loss: 0.00 Loss: 0.03



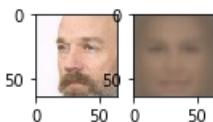
98%|###|ETA: 0:00:02 Epoch: 2 Iter: 2997 Class Loss: 0.00 Loss: 0.03



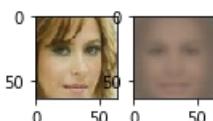
100%|####|Time: 0:02:34 Epoch: 2 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

N/A%| |ETA: ---- Epoch: nan Iter: nan Class Loss: nan Loss: nan



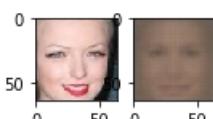
3%| |ETA: 0:02:18 Epoch: 3 Iter: 97 Class Loss: 0.04 Loss: 0.07



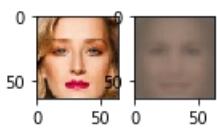
6%| |ETA: 0:02:15 Epoch: 3 Iter: 198 Class Loss: 0.00 Loss: 0.03



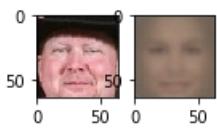
9%| |ETA: 0:02:09 Epoch: 3 Iter: 299 Class Loss: 0.00 Loss: 0.03



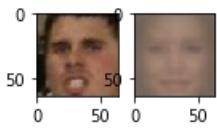
13%| |ETA: 0:02:05 Epoch: 3 Iter: 398 Class Loss: 0.00 Loss: 0.03



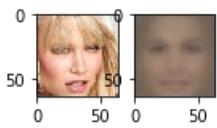
16% | ETA: 0:02:00 Epoch: 3 Iter: 497 Class Loss: 0.00 Loss: 0.03



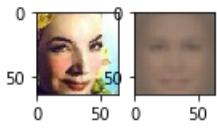
19% | ETA: 0:01:54 Epoch: 3 Iter: 599 Class Loss: 0.01 Loss: 0.03



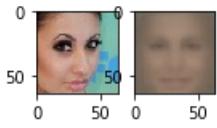
22% |# ETA: 0:01:51 Epoch: 3 Iter: 698 Class Loss: 0.00 Loss: 0.04



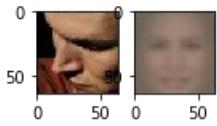
26% |# ETA: 0:01:46 Epoch: 3 Iter: 798 Class Loss: 0.00 Loss: 0.04



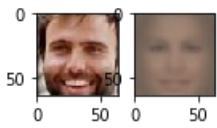
29% |# ETA: 0:01:41 Epoch: 3 Iter: 897 Class Loss: 0.00 Loss: 0.03



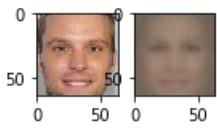
32% |# ETA: 0:01:49 Epoch: 3 Iter: 998 Class Loss: 0.00 Loss: 0.04



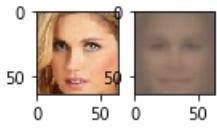
35% |# ETA: 0:01:42 Epoch: 3 Iter: 1097 Class Loss: 0.00 Loss: 0.04



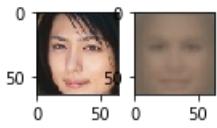
39% |# ETA: 0:01:28 Epoch: 3 Iter: 1198 Class Loss: 0.00 Loss: 0.03



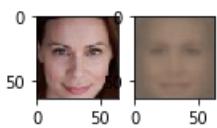
42% |# ETA: 0:01:23 Epoch: 3 Iter: 1298 Class Loss: 0.00 Loss: 0.03



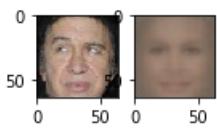
45% |# ETA: 0:01:17 Epoch: 3 Iter: 1399 Class Loss: 0.01 Loss: 0.04



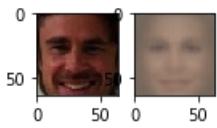
49% |# ETA: 0:01:12 Epoch: 3 Iter: 1497 Class Loss: 0.02 Loss: 0.06



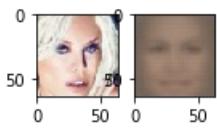
52%|##| ETA: 0:01:08 Epoch: 3 Iter: 1597 Class Loss: 0.02 Loss: 0.05



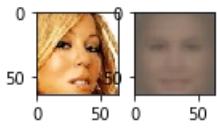
55%|##| ETA: 0:01:05 Epoch: 3 Iter: 1698 Class Loss: 0.00 Loss: 0.03



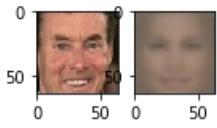
58%|##| ETA: 0:00:58 Epoch: 3 Iter: 1799 Class Loss: 0.00 Loss: 0.03



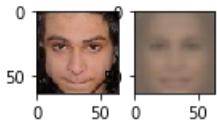
62%|##| ETA: 0:00:55 Epoch: 3 Iter: 1898 Class Loss: 0.00 Loss: 0.04



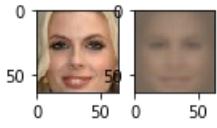
65%|##| ETA: 0:00:50 Epoch: 3 Iter: 1999 Class Loss: 0.00 Loss: 0.03



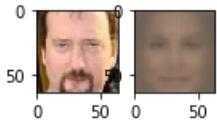
68%|##| ETA: 0:00:45 Epoch: 3 Iter: 2097 Class Loss: 0.03 Loss: 0.06



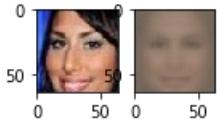
72%|##| ETA: 0:00:40 Epoch: 3 Iter: 2199 Class Loss: 0.00 Loss: 0.03



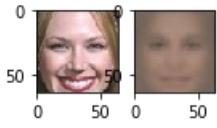
75%|###| ETA: 0:00:42 Epoch: 3 Iter: 2298 Class Loss: 0.01 Loss: 0.05



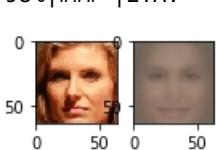
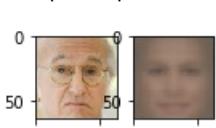
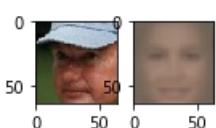
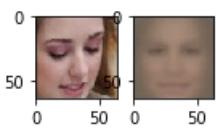
78%|###| ETA: 0:00:30 Epoch: 3 Iter: 2399 Class Loss: 0.00 Loss: 0.04



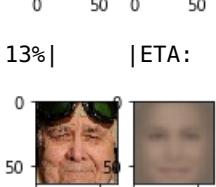
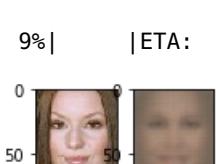
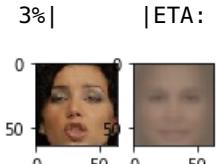
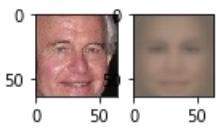
81%|###| ETA: 0:00:26 Epoch: 3 Iter: 2497 Class Loss: 0.00 Loss: 0.03

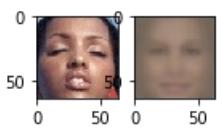


85%|###| ETA: 0:00:24 Epoch: 3 Iter: 2599 Class Loss: 0.00 Loss: 0.03

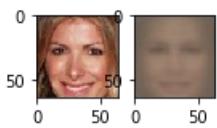


N/A%| | ETA: --:--- Epoch: nan Iter: nan Class Loss: nan Loss: nan





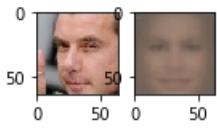
19% | ETA: 0:01:54 Epoch: 4 Iter: 597 Class Loss: 0.00 Loss: 0.04



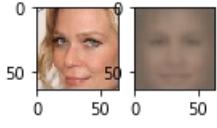
22% |# ETA: 0:01:50 Epoch: 4 Iter: 698 Class Loss: 0.00 Loss: 0.04



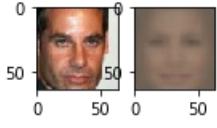
26% |# ETA: 0:01:45 Epoch: 4 Iter: 798 Class Loss: 0.01 Loss: 0.04



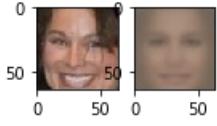
29% |# ETA: 0:01:43 Epoch: 4 Iter: 899 Class Loss: 0.08 Loss: 0.12



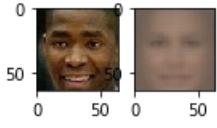
32% |# ETA: 0:01:37 Epoch: 4 Iter: 998 Class Loss: 0.00 Loss: 0.04



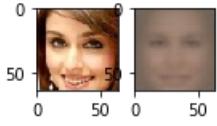
35% |# ETA: 0:01:46 Epoch: 4 Iter: 1098 Class Loss: 0.00 Loss: 0.03



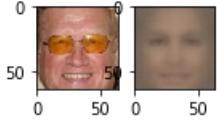
39% |# ETA: 0:01:39 Epoch: 4 Iter: 1198 Class Loss: 0.00 Loss: 0.03



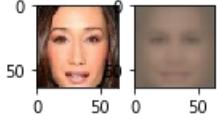
42% |# ETA: 0:01:25 Epoch: 4 Iter: 1298 Class Loss: 0.00 Loss: 0.04



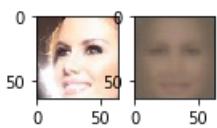
45% |# ETA: 0:01:18 Epoch: 4 Iter: 1397 Class Loss: 0.01 Loss: 0.05



49% |# ETA: 0:01:13 Epoch: 4 Iter: 1497 Class Loss: 0.02 Loss: 0.05



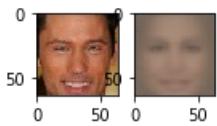
52% |## ETA: 0:01:10 Epoch: 4 Iter: 1598 Class Loss: 0.00 Loss: 0.03



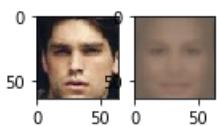
55%|##| ETA: 0:01:16 Epoch: 4 Iter: 1699 Class Loss: 0.00 Loss: 0.03



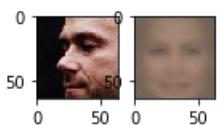
58%|##| ETA: 0:01:11 Epoch: 4 Iter: 1799 Class Loss: 0.01 Loss: 0.04



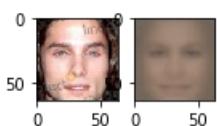
62%|##| ETA: 0:00:55 Epoch: 4 Iter: 1899 Class Loss: 0.00 Loss: 0.03



65%|##| ETA: 0:00:50 Epoch: 4 Iter: 1999 Class Loss: 0.00 Loss: 0.03



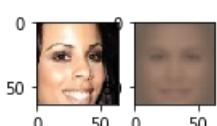
68%|##| ETA: 0:00:44 Epoch: 4 Iter: 2097 Class Loss: 0.06 Loss: 0.09



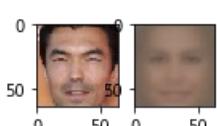
72%|##| ETA: 0:00:39 Epoch: 4 Iter: 2198 Class Loss: 0.00 Loss: 0.03



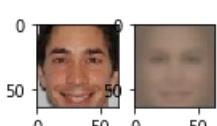
75%|###| ETA: 0:00:35 Epoch: 4 Iter: 2297 Class Loss: 0.00 Loss: 0.03



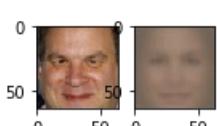
78%|###| ETA: 0:00:30 Epoch: 4 Iter: 2398 Class Loss: 0.00 Loss: 0.03



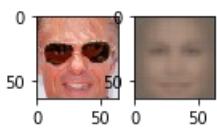
81%|###| ETA: 0:00:26 Epoch: 4 Iter: 2498 Class Loss: 0.00 Loss: 0.04



85%|###| ETA: 0:00:21 Epoch: 4 Iter: 2597 Class Loss: 0.00 Loss: 0.04



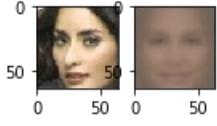
88%|###| ETA: 0:00:19 Epoch: 4 Iter: 2698 Class Loss: 0.00 Loss: 0.04



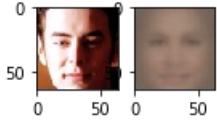
91%|###| ETA: 0:00:13 Epoch: 4 Iter: 2798 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:07 Epoch: 4 Iter: 2897 Class Loss: 0.01 Loss: 0.04



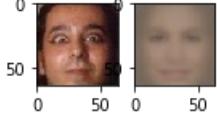
98%|###| ETA: 0:00:02 Epoch: 4 Iter: 2999 Class Loss: 0.00 Loss: 0.03



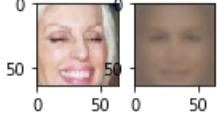
100%|####| Time: 0:02:32 Epoch: 4 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



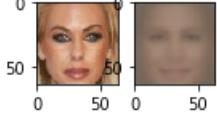
3%| | ETA: 0:02:21 Epoch: 5 Iter: 98 Class Loss: 0.00 Loss: 0.03



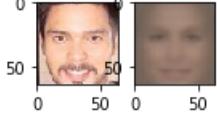
6%| | ETA: 0:02:14 Epoch: 5 Iter: 198 Class Loss: 0.00 Loss: 0.04



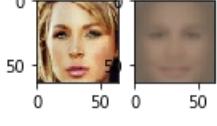
9%| | ETA: 0:02:12 Epoch: 5 Iter: 297 Class Loss: 0.00 Loss: 0.04



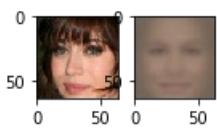
13%| | ETA: 0:02:04 Epoch: 5 Iter: 399 Class Loss: 0.00 Loss: 0.04



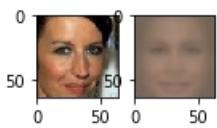
16%| | ETA: 0:01:59 Epoch: 5 Iter: 499 Class Loss: 0.00 Loss: 0.03



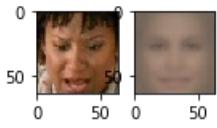
19%| | ETA: 0:01:54 Epoch: 5 Iter: 599 Class Loss: 0.01 Loss: 0.05



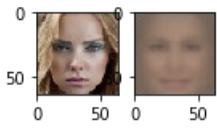
22%|# |ETA: 0:01:51 Epoch: 5 Iter: 698 Class Loss: 0.00 Loss: 0.03



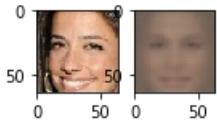
26%|# |ETA: 0:01:46 Epoch: 5 Iter: 798 Class Loss: 0.00 Loss: 0.04



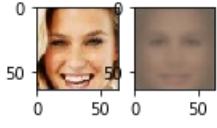
29%|# |ETA: 0:01:41 Epoch: 5 Iter: 897 Class Loss: 0.00 Loss: 0.03



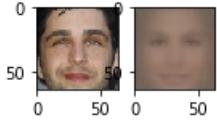
32%|# |ETA: 0:01:36 Epoch: 5 Iter: 997 Class Loss: 0.00 Loss: 0.03



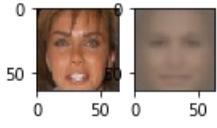
35%|# |ETA: 0:01:32 Epoch: 5 Iter: 1097 Class Loss: 0.00 Loss: 0.03



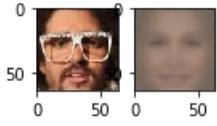
39%|# |ETA: 0:01:39 Epoch: 5 Iter: 1198 Class Loss: 0.00 Loss: 0.03



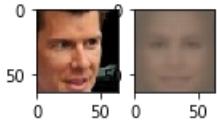
42%|# |ETA: 0:01:35 Epoch: 5 Iter: 1299 Class Loss: 0.00 Loss: 0.03



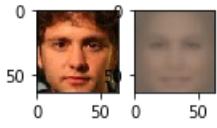
45%|# |ETA: 0:01:18 Epoch: 5 Iter: 1399 Class Loss: 0.00 Loss: 0.04



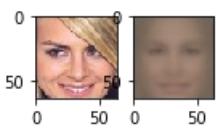
49%|# |ETA: 0:01:12 Epoch: 5 Iter: 1497 Class Loss: 0.00 Loss: 0.03



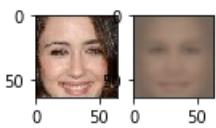
52%|## |ETA: 0:01:08 Epoch: 5 Iter: 1599 Class Loss: 0.00 Loss: 0.03



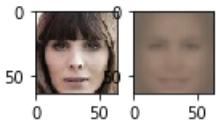
55%|## |ETA: 0:01:03 Epoch: 5 Iter: 1698 Class Loss: 0.00 Loss: 0.03



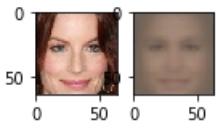
58%|##| ETA: 0:01:00 Epoch: 5 Iter: 1797 Class Loss: 0.04 Loss: 0.07



62%|##| ETA: 0:00:54 Epoch: 5 Iter: 1897 Class Loss: 0.00 Loss: 0.03



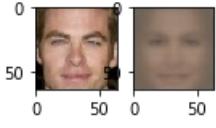
65%|##| ETA: 0:00:50 Epoch: 5 Iter: 1998 Class Loss: 0.00 Loss: 0.03



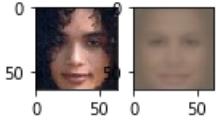
68%|##| ETA: 0:00:46 Epoch: 5 Iter: 2098 Class Loss: 0.00 Loss: 0.03



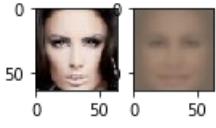
72%|##| ETA: 0:00:47 Epoch: 5 Iter: 2198 Class Loss: 0.00 Loss: 0.03



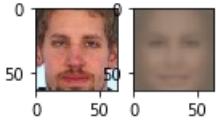
75%|###| ETA: 0:00:36 Epoch: 5 Iter: 2297 Class Loss: 0.00 Loss: 0.03



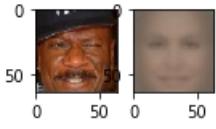
78%|###| ETA: 0:00:30 Epoch: 5 Iter: 2399 Class Loss: 0.00 Loss: 0.03



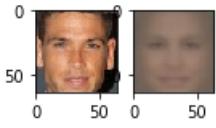
81%|###| ETA: 0:00:25 Epoch: 5 Iter: 2499 Class Loss: 0.00 Loss: 0.04



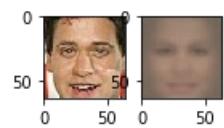
85%|###| ETA: 0:00:21 Epoch: 5 Iter: 2597 Class Loss: 0.00 Loss: 0.04



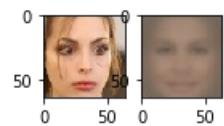
88%|###| ETA: 0:00:16 Epoch: 5 Iter: 2698 Class Loss: 0.00 Loss: 0.03



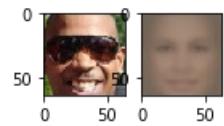
91%|###| ETA: 0:00:13 Epoch: 5 Iter: 2798 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:08 Epoch: 5 Iter: 2898 Class Loss: 0.00 Loss: 0.04



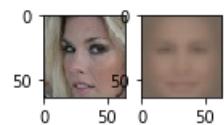
98%|###| ETA: 0:00:02 Epoch: 5 Iter: 2997 Class Loss: 0.00 Loss: 0.03



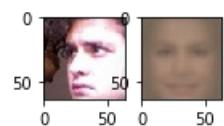
100%|####| Time: 0:02:32 Epoch: 5 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

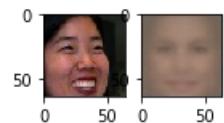
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



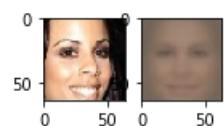
3%| | ETA: 0:02:17 Epoch: 6 Iter: 98 Class Loss: 0.00 Loss: 0.03



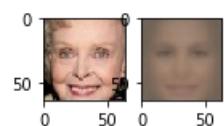
6%| | ETA: 0:02:13 Epoch: 6 Iter: 198 Class Loss: 0.00 Loss: 0.03



9%| | ETA: 0:02:10 Epoch: 6 Iter: 298 Class Loss: 0.27 Loss: 0.30



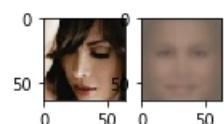
13%| | ETA: 0:02:04 Epoch: 6 Iter: 397 Class Loss: 0.00 Loss: 0.03



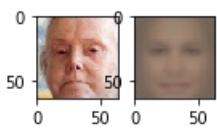
16%| | ETA: 0:01:59 Epoch: 6 Iter: 497 Class Loss: 0.00 Loss: 0.03



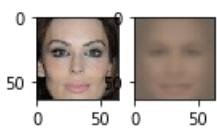
19%| | ETA: 0:01:54 Epoch: 6 Iter: 598 Class Loss: 0.00 Loss: 0.03



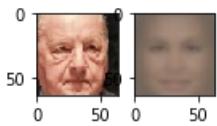
22%|#| ETA: 0:01:51 Epoch: 6 Iter: 698 Class Loss: 0.00 Loss: 0.03



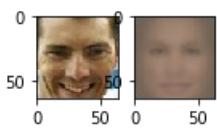
26%|# |ETA: 0:01:44 Epoch: 6 Iter: 799 Class Loss: 0.00 Loss: 0.04



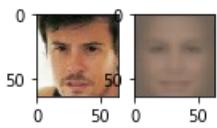
29%|# |ETA: 0:01:39 Epoch: 6 Iter: 899 Class Loss: 0.00 Loss: 0.03



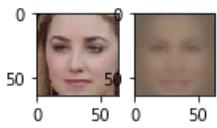
32%|# |ETA: 0:01:35 Epoch: 6 Iter: 998 Class Loss: 0.00 Loss: 0.03



36%|# |ETA: 0:01:31 Epoch: 6 Iter: 1099 Class Loss: 0.00 Loss: 0.03



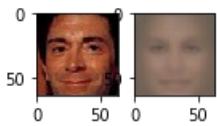
39%|# |ETA: 0:01:26 Epoch: 6 Iter: 1198 Class Loss: 0.00 Loss: 0.03



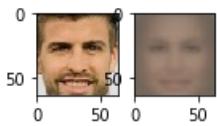
42%|# |ETA: 0:01:28 Epoch: 6 Iter: 1298 Class Loss: 0.00 Loss: 0.04



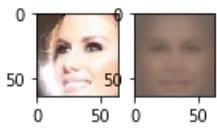
45%|# |ETA: 0:01:29 Epoch: 6 Iter: 1399 Class Loss: 0.00 Loss: 0.03



49%|# |ETA: 0:01:14 Epoch: 6 Iter: 1497 Class Loss: 0.00 Loss: 0.03



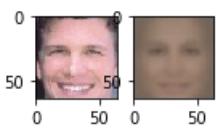
52%## |ETA: 0:01:07 Epoch: 6 Iter: 1598 Class Loss: 0.00 Loss: 0.04



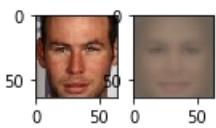
55%## |ETA: 0:01:17 Epoch: 6 Iter: 1699 Class Loss: 0.00 Loss: 0.04



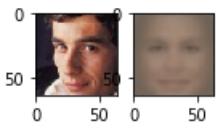
58%## |ETA: 0:01:11 Epoch: 6 Iter: 1799 Class Loss: 0.00 Loss: 0.04



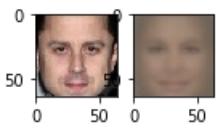
62%|##| ETA: 0:00:53 Epoch: 6 Iter: 1899 Class Loss: 0.00 Loss: 0.03



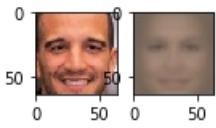
65%|##| ETA: 0:00:49 Epoch: 6 Iter: 1999 Class Loss: 0.00 Loss: 0.04



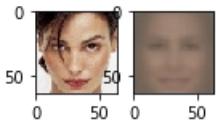
68%|##| ETA: 0:00:44 Epoch: 6 Iter: 2098 Class Loss: 0.00 Loss: 0.04



71%|##| ETA: 0:00:40 Epoch: 6 Iter: 2197 Class Loss: 0.00 Loss: 0.03



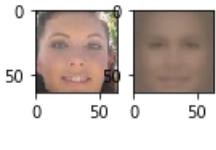
75%|###| ETA: 0:00:36 Epoch: 6 Iter: 2297 Class Loss: 0.00 Loss: 0.04



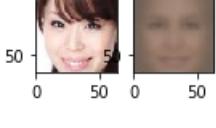
78%|###| ETA: 0:00:30 Epoch: 6 Iter: 2399 Class Loss: 0.00 Loss: 0.03



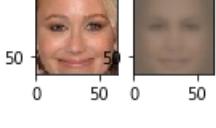
81%|###| ETA: 0:00:25 Epoch: 6 Iter: 2499 Class Loss: 0.00 Loss: 0.03



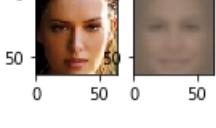
85%|###| ETA: 0:00:21 Epoch: 6 Iter: 2598 Class Loss: 0.00 Loss: 0.03



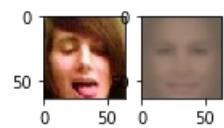
88%|###| ETA: 0:00:16 Epoch: 6 Iter: 2698 Class Loss: 0.00 Loss: 0.03



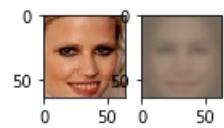
91%|###| ETA: 0:00:11 Epoch: 6 Iter: 2799 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:07 Epoch: 6 Iter: 2898 Class Loss: 0.00 Loss: 0.03



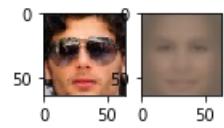
98%|###| ETA: 0:00:02 Epoch: 6 Iter: 2999 Class Loss: 0.00 Loss: 0.04



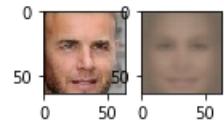
100%|####| Time: 0:02:32 Epoch: 6 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

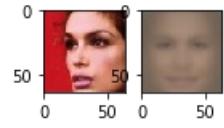
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



3%| | ETA: 0:02:17 Epoch: 7 Iter: 99 Class Loss: 0.00 Loss: 0.03



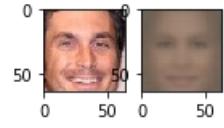
6%| | ETA: 0:02:12 Epoch: 7 Iter: 198 Class Loss: 0.00 Loss: 0.03



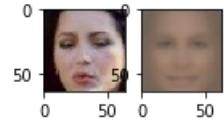
9%| | ETA: 0:02:08 Epoch: 7 Iter: 299 Class Loss: 0.00 Loss: 0.03



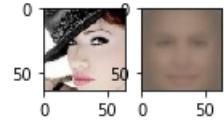
13%| | ETA: 0:02:03 Epoch: 7 Iter: 397 Class Loss: 0.00 Loss: 0.04



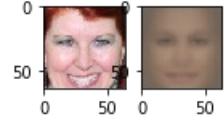
16%| | ETA: 0:01:58 Epoch: 7 Iter: 499 Class Loss: 0.00 Loss: 0.03



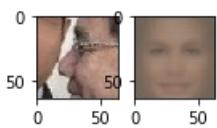
19%| | ETA: 0:01:54 Epoch: 7 Iter: 598 Class Loss: 0.00 Loss: 0.03



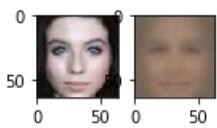
22%|#| ETA: 0:01:50 Epoch: 7 Iter: 697 Class Loss: 0.00 Loss: 0.04



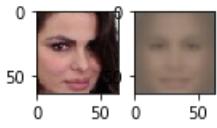
26%|#| ETA: 0:01:45 Epoch: 7 Iter: 798 Class Loss: 0.00 Loss: 0.03



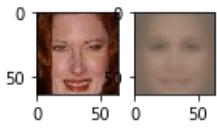
29%|# |ETA: 0:01:43 Epoch: 7 Iter: 897 Class Loss: 0.00 Loss: 0.03



32%|# |ETA: 0:01:38 Epoch: 7 Iter: 997 Class Loss: 0.25 Loss: 0.29



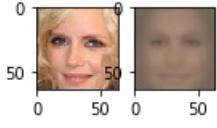
36%|# |ETA: 0:01:32 Epoch: 7 Iter: 1099 Class Loss: 0.00 Loss: 0.04



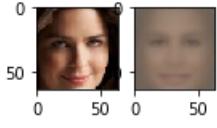
39%|# |ETA: 0:01:26 Epoch: 7 Iter: 1197 Class Loss: 0.00 Loss: 0.03



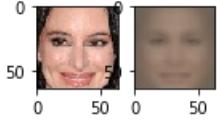
42%|# |ETA: 0:01:22 Epoch: 7 Iter: 1299 Class Loss: 0.00 Loss: 0.03



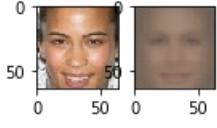
45%|# |ETA: 0:01:18 Epoch: 7 Iter: 1397 Class Loss: 0.00 Loss: 0.03



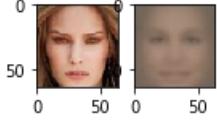
49%|# |ETA: 0:01:24 Epoch: 7 Iter: 1498 Class Loss: 0.00 Loss: 0.04



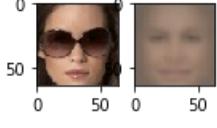
52%|## |ETA: 0:01:11 Epoch: 7 Iter: 1598 Class Loss: 0.00 Loss: 0.03



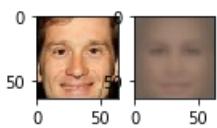
55%|## |ETA: 0:00:59 Epoch: 7 Iter: 1699 Class Loss: 0.00 Loss: 0.03



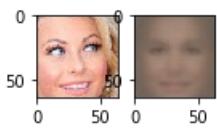
58%|## |ETA: 0:00:54 Epoch: 7 Iter: 1798 Class Loss: 0.00 Loss: 0.03



62%|## |ETA: 0:00:50 Epoch: 7 Iter: 1899 Class Loss: 0.00 Loss: 0.03



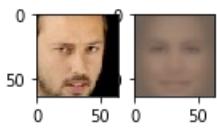
65%|##| ETA: 0:00:46 Epoch: 7 Iter: 1997 Class Loss: 0.00 Loss: 0.03



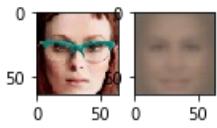
68%|##| ETA: 0:00:42 Epoch: 7 Iter: 2099 Class Loss: 0.00 Loss: 0.03



72%|##| ETA: 0:00:43 Epoch: 7 Iter: 2199 Class Loss: 0.13 Loss: 0.16



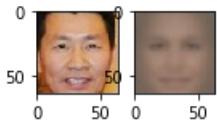
75%|###| ETA: 0:00:32 Epoch: 7 Iter: 2298 Class Loss: 0.00 Loss: 0.03



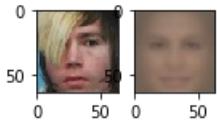
78%|###| ETA: 0:00:28 Epoch: 7 Iter: 2398 Class Loss: 0.00 Loss: 0.04



81%|###| ETA: 0:00:24 Epoch: 7 Iter: 2497 Class Loss: 0.00 Loss: 0.03



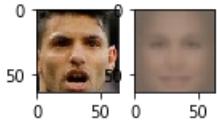
85%|###| ETA: 0:00:19 Epoch: 7 Iter: 2597 Class Loss: 0.00 Loss: 0.03



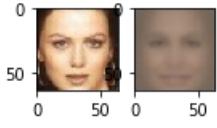
88%|###| ETA: 0:00:15 Epoch: 7 Iter: 2699 Class Loss: 0.00 Loss: 0.03



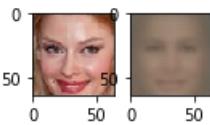
91%|###| ETA: 0:00:11 Epoch: 7 Iter: 2797 Class Loss: 0.00 Loss: 0.04



94%|###| ETA: 0:00:06 Epoch: 7 Iter: 2898 Class Loss: 0.00 Loss: 0.03



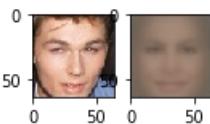
98%|###| ETA: 0:00:02 Epoch: 7 Iter: 2998 Class Loss: 0.00 Loss: 0.03



100%|####|Time: 0:02:25 Epoch: 7 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

N/A%| |ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



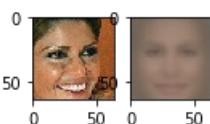
3%| |ETA: 0:02:27 Epoch: 8 Iter: 98 Class Loss: 0.00 Loss: 0.03



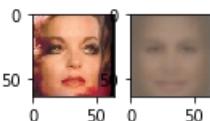
6%| |ETA: 0:02:22 Epoch: 8 Iter: 199 Class Loss: 0.00 Loss: 0.03



9%| |ETA: 0:02:00 Epoch: 8 Iter: 297 Class Loss: 0.00 Loss: 0.03



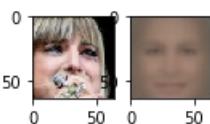
13%| |ETA: 0:01:55 Epoch: 8 Iter: 398 Class Loss: 0.00 Loss: 0.03



16%| |ETA: 0:01:52 Epoch: 8 Iter: 498 Class Loss: 0.00 Loss: 0.03



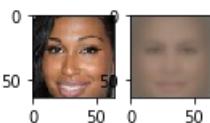
19%| |ETA: 0:01:47 Epoch: 8 Iter: 597 Class Loss: 0.00 Loss: 0.03



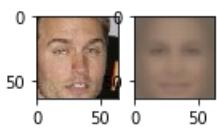
22%|# |ETA: 0:01:43 Epoch: 8 Iter: 698 Class Loss: 0.00 Loss: 0.04



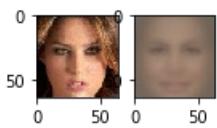
26%|# |ETA: 0:01:38 Epoch: 8 Iter: 799 Class Loss: 0.00 Loss: 0.04



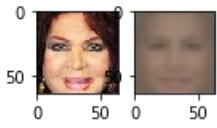
29%|# |ETA: 0:01:33 Epoch: 8 Iter: 897 Class Loss: 0.00 Loss: 0.03



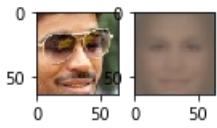
32%|# |ETA: 0:01:30 Epoch: 8 Iter: 998 Class Loss: 0.00 Loss: 0.03



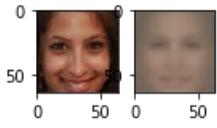
35%|# |ETA: 0:01:26 Epoch: 8 Iter: 1097 Class Loss: 0.00 Loss: 0.03



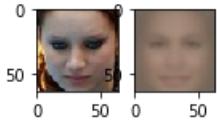
39%|# |ETA: 0:01:21 Epoch: 8 Iter: 1198 Class Loss: 0.02 Loss: 0.04



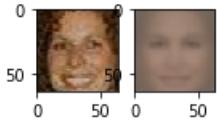
42%|# |ETA: 0:01:17 Epoch: 8 Iter: 1298 Class Loss: 0.00 Loss: 0.03



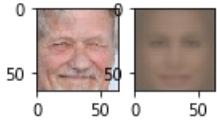
45%|# |ETA: 0:01:12 Epoch: 8 Iter: 1397 Class Loss: 0.00 Loss: 0.03



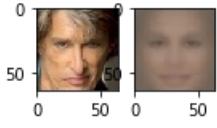
49%|# |ETA: 0:01:09 Epoch: 8 Iter: 1497 Class Loss: 0.00 Loss: 0.03



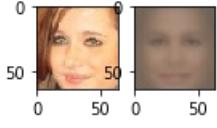
52%|## |ETA: 0:01:03 Epoch: 8 Iter: 1598 Class Loss: 0.00 Loss: 0.03



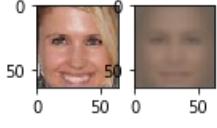
55%|## |ETA: 0:00:59 Epoch: 8 Iter: 1697 Class Loss: 0.00 Loss: 0.03



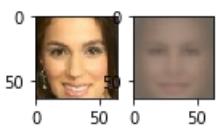
58%|## |ETA: 0:00:59 Epoch: 8 Iter: 1797 Class Loss: 0.00 Loss: 0.03



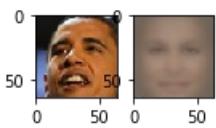
62%|## |ETA: 0:00:57 Epoch: 8 Iter: 1898 Class Loss: 0.00 Loss: 0.03



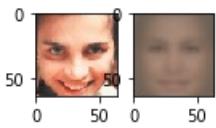
65%|## |ETA: 0:00:55 Epoch: 8 Iter: 1998 Class Loss: 0.00 Loss: 0.03



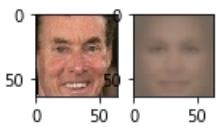
68%|##| ETA: 0:00:50 Epoch: 8 Iter: 2098 Class Loss: 0.00 Loss: 0.03



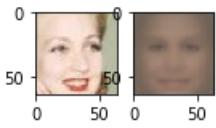
71%|##| ETA: 0:00:37 Epoch: 8 Iter: 2197 Class Loss: 0.00 Loss: 0.03



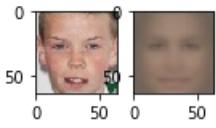
75%|###| ETA: 0:00:33 Epoch: 8 Iter: 2297 Class Loss: 0.00 Loss: 0.03



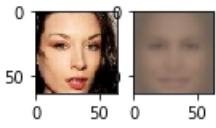
78%|###| ETA: 0:00:28 Epoch: 8 Iter: 2399 Class Loss: 0.00 Loss: 0.03



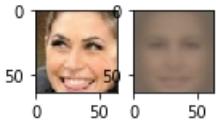
81%|###| ETA: 0:00:24 Epoch: 8 Iter: 2498 Class Loss: 0.00 Loss: 0.03



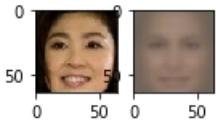
85%|###| ETA: 0:00:19 Epoch: 8 Iter: 2598 Class Loss: 0.00 Loss: 0.04



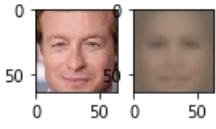
88%|###| ETA: 0:00:15 Epoch: 8 Iter: 2699 Class Loss: 0.00 Loss: 0.03



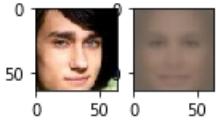
91%|###| ETA: 0:00:11 Epoch: 8 Iter: 2797 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 8 Iter: 2898 Class Loss: 0.00 Loss: 0.03



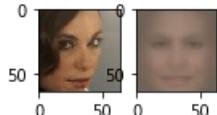
98%|###| ETA: 0:00:02 Epoch: 8 Iter: 2998 Class Loss: 0.00 Loss: 0.03



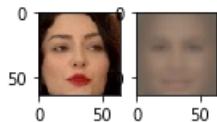
100%|####| Time: 0:02:22 Epoch: 8 Iter: 3052 Class Loss: 0.01 Loss: 0.04

Recomputing the sampling probabilities

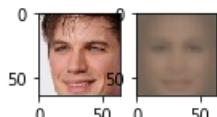
N/A%| | ETA: --:--- Epoch: nan Iter: nan Class Loss: nan Loss: nan



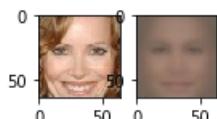
3%| | ETA: 0:02:09 Epoch: 9 Iter: 97 Class Loss: 0.00 Loss: 0.03



6%| | ETA: 0:02:05 Epoch: 9 Iter: 198 Class Loss: 0.00 Loss: 0.03



9%| | ETA: 0:02:01 Epoch: 9 Iter: 297 Class Loss: 0.00 Loss: 0.03



13%| | ETA: 0:02:00 Epoch: 9 Iter: 399 Class Loss: 0.08 Loss: 0.12



16%| | ETA: 0:02:07 Epoch: 9 Iter: 499 Class Loss: 0.00 Loss: 0.03



19%| | ETA: 0:01:56 Epoch: 9 Iter: 599 Class Loss: 0.00 Loss: 0.03



22%|# | ETA: 0:01:43 Epoch: 9 Iter: 698 Class Loss: 0.00 Loss: 0.03



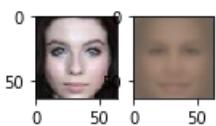
26%|# | ETA: 0:01:39 Epoch: 9 Iter: 798 Class Loss: 0.00 Loss: 0.03



29%|# | ETA: 0:01:36 Epoch: 9 Iter: 899 Class Loss: 0.00 Loss: 0.03



32%|# | ETA: 0:01:29 Epoch: 9 Iter: 999 Class Loss: 0.00 Loss: 0.03



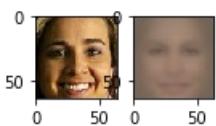
36%|# |ETA: 0:01:25 Epoch: 9 Iter: 1099 Class Loss: 0.00 Loss: 0.03



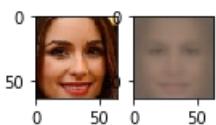
39%|# |ETA: 0:01:22 Epoch: 9 Iter: 1197 Class Loss: 0.00 Loss: 0.03



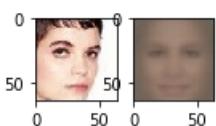
42%|# |ETA: 0:01:17 Epoch: 9 Iter: 1299 Class Loss: 0.01 Loss: 0.04



45%|# |ETA: 0:01:12 Epoch: 9 Iter: 1397 Class Loss: 0.00 Loss: 0.03



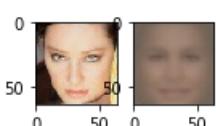
49%|# |ETA: 0:01:08 Epoch: 9 Iter: 1497 Class Loss: 0.00 Loss: 0.03



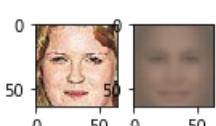
52%|## |ETA: 0:01:04 Epoch: 9 Iter: 1598 Class Loss: 0.00 Loss: 0.03



55%|## |ETA: 0:01:00 Epoch: 9 Iter: 1697 Class Loss: 0.00 Loss: 0.03



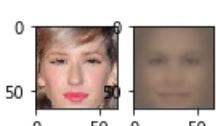
58%|## |ETA: 0:00:56 Epoch: 9 Iter: 1797 Class Loss: 0.00 Loss: 0.03



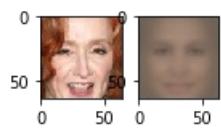
62%|## |ETA: 0:00:52 Epoch: 9 Iter: 1899 Class Loss: 0.00 Loss: 0.03



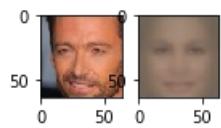
65%|## |ETA: 0:00:47 Epoch: 9 Iter: 1997 Class Loss: 0.00 Loss: 0.03



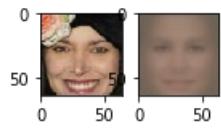
68%|## |ETA: 0:00:42 Epoch: 9 Iter: 2098 Class Loss: 0.00 Loss: 0.04



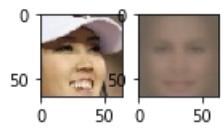
71%|##| ETA: 0:00:43 Epoch: 9 Iter: 2197 Class Loss: 0.00 Loss: 0.04



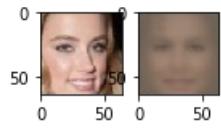
75%|###| ETA: 0:00:38 Epoch: 9 Iter: 2298 Class Loss: 0.00 Loss: 0.04



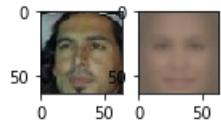
78%|###| ETA: 0:00:28 Epoch: 9 Iter: 2399 Class Loss: 0.00 Loss: 0.03



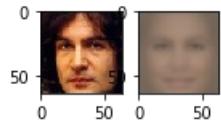
81%|###| ETA: 0:00:30 Epoch: 9 Iter: 2499 Class Loss: 0.00 Loss: 0.03



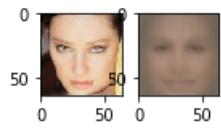
85%|###| ETA: 0:00:20 Epoch: 9 Iter: 2598 Class Loss: 0.00 Loss: 0.04



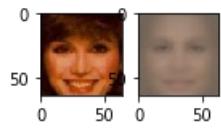
88%|###| ETA: 0:00:15 Epoch: 9 Iter: 2698 Class Loss: 0.00 Loss: 0.03



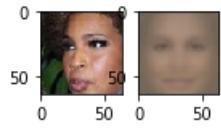
91%|###| ETA: 0:00:11 Epoch: 9 Iter: 2797 Class Loss: 0.00 Loss: 0.04



94%|###| ETA: 0:00:06 Epoch: 9 Iter: 2898 Class Loss: 0.00 Loss: 0.03



98%|###| ETA: 0:00:02 Epoch: 9 Iter: 2997 Class Loss: 0.00 Loss: 0.03



100%|####| Time: 0:02:24 Epoch: 9 Iter: 3052 Class Loss: 0.00 Loss: 0.03
100% (97 of 97) |#####| Elapsed Time: 0:01:09 Time: 0:01:09
N/A% (0 of 72) | Elapsed Time: 0:00:00 ETA: --:--:--

male lighter: 0.835051546392

100% (72 of 72) |#####| Elapsed Time: 0:00:50 Time: 0:00:50
N/A% (0 of 78) | Elapsed Time: 0:00:00 ETA: --:--:--

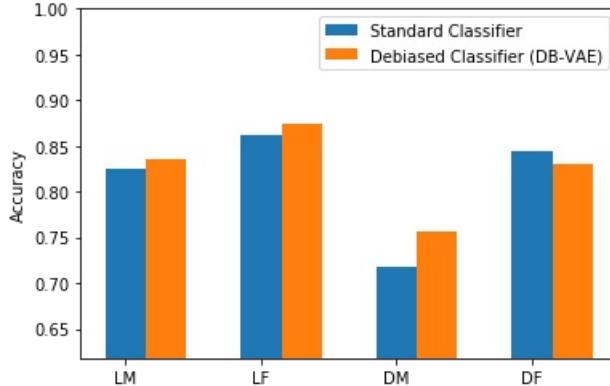
female lighter: 0.875

100% (78 of 78) |#####| Elapsed Time: 0:00:52 Time: 0:00:52
N/A% (0 of 71) | Elapsed Time: 0:00:00 ETA: --:--:--

male darker: 0.75641025641

100% (71 of 71) |#####| Elapsed Time: 0:00:48 Time: 0:00:48

female darker: 0.830985915493



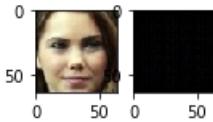
Smooth Factor of 0.5

In [62]:

```
run_smooth_factor(0.5)
```

Recomputing the sampling probabilities

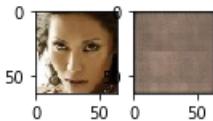
N/A% | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



3% | ETA: 0:02:12 Epoch: 0 Iter: 98 Class Loss: 0.16 Loss: 0.19



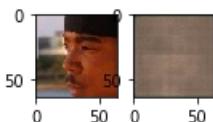
6% | ETA: 0:02:07 Epoch: 0 Iter: 198 Class Loss: 0.06 Loss: 0.10



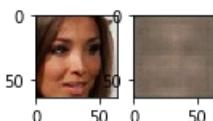
9% | ETA: 0:02:02 Epoch: 0 Iter: 299 Class Loss: 0.07 Loss: 0.11



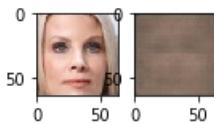
13% | ETA: 0:01:58 Epoch: 0 Iter: 397 Class Loss: 0.13 Loss: 0.17



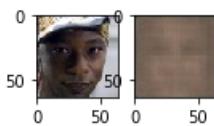
16% | ETA: 0:01:53 Epoch: 0 Iter: 499 Class Loss: 0.21 Loss: 0.25



19% | ETA: 0:01:49 Epoch: 0 Iter: 599 Class Loss: 0.01 Loss: 0.04



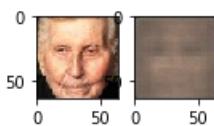
22%|# |ETA: 0:01:44 Epoch: 0 Iter: 697 Class Loss: 0.16 Loss: 0.19



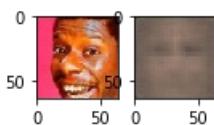
26%|# |ETA: 0:01:42 Epoch: 0 Iter: 798 Class Loss: 0.08 Loss: 0.11



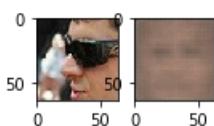
29%|# |ETA: 0:01:37 Epoch: 0 Iter: 897 Class Loss: 0.05 Loss: 0.09



32%|# |ETA: 0:01:32 Epoch: 0 Iter: 997 Class Loss: 0.01 Loss: 0.04



36%|# |ETA: 0:01:27 Epoch: 0 Iter: 1099 Class Loss: 0.03 Loss: 0.06



39%|# |ETA: 0:01:38 Epoch: 0 Iter: 1198 Class Loss: 0.01 Loss: 0.05



42%|# |ETA: 0:01:29 Epoch: 0 Iter: 1298 Class Loss: 0.01 Loss: 0.05



45%|# |ETA: 0:01:14 Epoch: 0 Iter: 1399 Class Loss: 0.00 Loss: 0.04



49%|# |ETA: 0:01:09 Epoch: 0 Iter: 1497 Class Loss: 0.08 Loss: 0.12



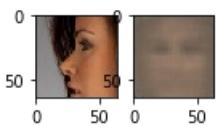
52%|## |ETA: 0:01:04 Epoch: 0 Iter: 1598 Class Loss: 0.01 Loss: 0.05



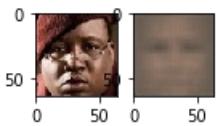
55%|## |ETA: 0:01:00 Epoch: 0 Iter: 1698 Class Loss: 0.01 Loss: 0.05



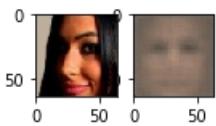
58%|##| ETA: 0:00:56 Epoch: 0 Iter: 1798 Class Loss: 0.00 Loss: 0.03



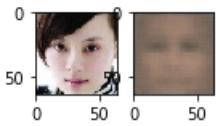
62%|##| ETA: 0:00:51 Epoch: 0 Iter: 1899 Class Loss: 0.01 Loss: 0.05



65%|##| ETA: 0:00:46 Epoch: 0 Iter: 1999 Class Loss: 0.00 Loss: 0.04



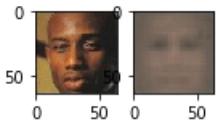
68%|##| ETA: 0:00:42 Epoch: 0 Iter: 2098 Class Loss: 0.00 Loss: 0.04



72%|##| ETA: 0:00:38 Epoch: 0 Iter: 2198 Class Loss: 0.00 Loss: 0.03



75%|###| ETA: 0:00:33 Epoch: 0 Iter: 2298 Class Loss: 0.02 Loss: 0.06



78%|###| ETA: 0:00:29 Epoch: 0 Iter: 2399 Class Loss: 0.01 Loss: 0.05



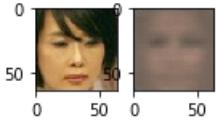
81%|###| ETA: 0:00:24 Epoch: 0 Iter: 2499 Class Loss: 0.00 Loss: 0.04



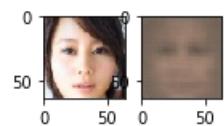
85%|###| ETA: 0:00:20 Epoch: 0 Iter: 2598 Class Loss: 0.00 Loss: 0.04



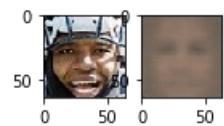
88%|###| ETA: 0:00:15 Epoch: 0 Iter: 2697 Class Loss: 0.00 Loss: 0.04



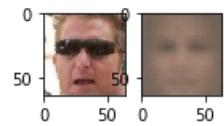
91%|###| ETA: 0:00:11 Epoch: 0 Iter: 2797 Class Loss: 0.05 Loss: 0.09



94%|###| ETA: 0:00:07 Epoch: 0 Iter: 2898 Class Loss: 0.00 Loss: 0.04



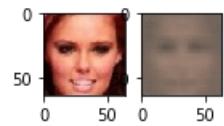
98%|###| ETA: 0:00:02 Epoch: 0 Iter: 2998 Class Loss: 0.02 Loss: 0.05



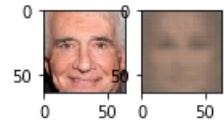
100%|####| Time: 0:02:23 Epoch: 0 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

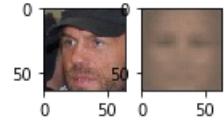
N/A%| | ETA: --:--- Epoch: nan Iter: nan Class Loss: nan Loss: nan



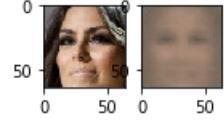
3%| | ETA: 0:02:12 Epoch: 1 Iter: 97 Class Loss: 0.00 Loss: 0.03



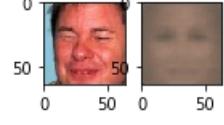
6%| | ETA: 0:02:07 Epoch: 1 Iter: 199 Class Loss: 0.00 Loss: 0.03



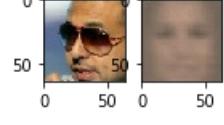
9%| | ETA: 0:02:02 Epoch: 1 Iter: 299 Class Loss: 0.00 Loss: 0.04



13%| | ETA: 0:01:58 Epoch: 1 Iter: 397 Class Loss: 0.01 Loss: 0.04



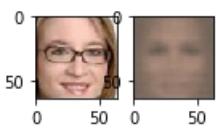
16%| | ETA: 0:01:54 Epoch: 1 Iter: 498 Class Loss: 0.00 Loss: 0.04



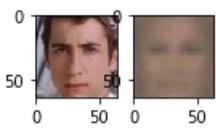
19%| | ETA: 0:01:49 Epoch: 1 Iter: 597 Class Loss: 0.00 Loss: 0.03



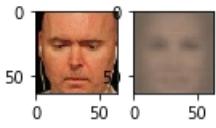
22%|#| ETA: 0:01:44 Epoch: 1 Iter: 697 Class Loss: 0.04 Loss: 0.07



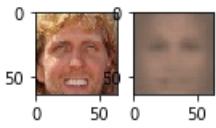
26%|# |ETA: 0:01:39 Epoch: 1 Iter: 799 Class Loss: 0.00 Loss: 0.04



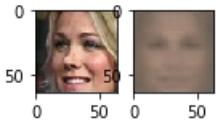
29%|# |ETA: 0:01:37 Epoch: 1 Iter: 897 Class Loss: 0.21 Loss: 0.24



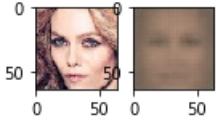
32%|# |ETA: 0:01:31 Epoch: 1 Iter: 998 Class Loss: 0.04 Loss: 0.07



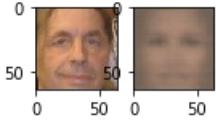
36%|# |ETA: 0:01:27 Epoch: 1 Iter: 1099 Class Loss: 0.00 Loss: 0.03



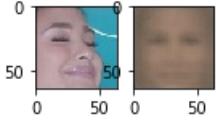
39%|# |ETA: 0:01:39 Epoch: 1 Iter: 1198 Class Loss: 0.01 Loss: 0.04



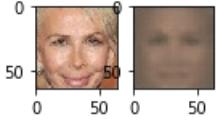
42%|# |ETA: 0:01:28 Epoch: 1 Iter: 1299 Class Loss: 0.00 Loss: 0.04



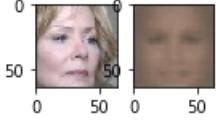
45%|# |ETA: 0:01:14 Epoch: 1 Iter: 1399 Class Loss: 0.01 Loss: 0.04



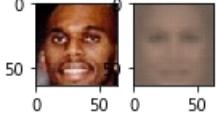
49%|# |ETA: 0:01:19 Epoch: 1 Iter: 1499 Class Loss: 0.00 Loss: 0.03



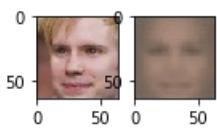
52%|## |ETA: 0:01:12 Epoch: 1 Iter: 1599 Class Loss: 0.00 Loss: 0.04



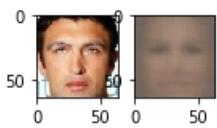
55%|## |ETA: 0:01:00 Epoch: 1 Iter: 1697 Class Loss: 0.01 Loss: 0.05



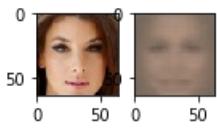
58%|## |ETA: 0:00:56 Epoch: 1 Iter: 1797 Class Loss: 0.02 Loss: 0.06



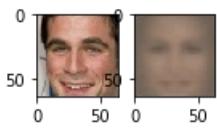
62%|##| ETA: 0:00:51 Epoch: 1 Iter: 1897 Class Loss: 0.00 Loss: 0.04



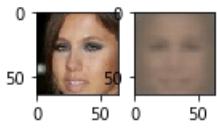
65%|##| ETA: 0:00:46 Epoch: 1 Iter: 1999 Class Loss: 0.00 Loss: 0.03



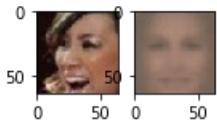
68%|##| ETA: 0:00:42 Epoch: 1 Iter: 2098 Class Loss: 0.00 Loss: 0.03



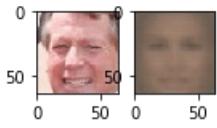
72%|##| ETA: 0:00:38 Epoch: 1 Iter: 2199 Class Loss: 0.01 Loss: 0.04



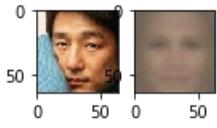
75%|###| ETA: 0:00:33 Epoch: 1 Iter: 2297 Class Loss: 0.00 Loss: 0.04



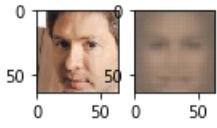
78%|###| ETA: 0:00:29 Epoch: 1 Iter: 2399 Class Loss: 0.00 Loss: 0.03



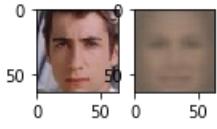
81%|###| ETA: 0:00:24 Epoch: 1 Iter: 2499 Class Loss: 0.00 Loss: 0.03



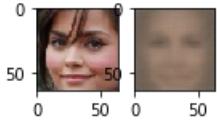
85%|###| ETA: 0:00:20 Epoch: 1 Iter: 2598 Class Loss: 0.01 Loss: 0.04



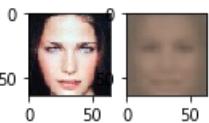
88%|###| ETA: 0:00:16 Epoch: 1 Iter: 2698 Class Loss: 0.00 Loss: 0.04



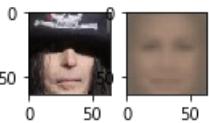
91%|###| ETA: 0:00:11 Epoch: 1 Iter: 2797 Class Loss: 0.00 Loss: 0.04



94%|###| ETA: 0:00:06 Epoch: 1 Iter: 2898 Class Loss: 0.00 Loss: 0.03



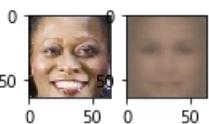
98%|###| ETA: 0:00:02 Epoch: 1 Iter: 2999 Class Loss: 0.00 Loss: 0.04



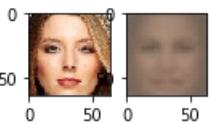
100%|####| Time: 0:02:23 Epoch: 1 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

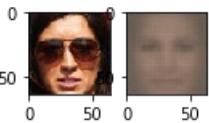
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



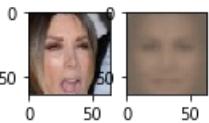
3%| | ETA: 0:02:28 Epoch: 2 Iter: 97 Class Loss: 0.00 Loss: 0.04



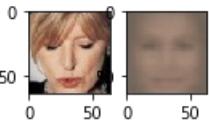
6%| | ETA: 0:02:24 Epoch: 2 Iter: 197 Class Loss: 0.01 Loss: 0.05



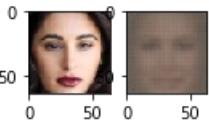
9%| | ETA: 0:02:03 Epoch: 2 Iter: 299 Class Loss: 0.00 Loss: 0.03



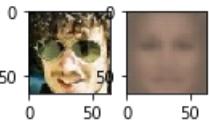
13%| | ETA: 0:01:58 Epoch: 2 Iter: 397 Class Loss: 0.00 Loss: 0.04



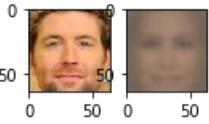
16%| | ETA: 0:01:53 Epoch: 2 Iter: 498 Class Loss: 0.00 Loss: 0.03



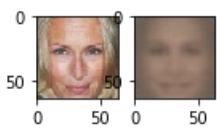
19%| | ETA: 0:01:49 Epoch: 2 Iter: 598 Class Loss: 0.00 Loss: 0.04



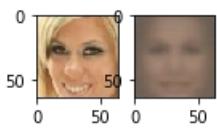
22%|#| ETA: 0:01:44 Epoch: 2 Iter: 697 Class Loss: 0.00 Loss: 0.03



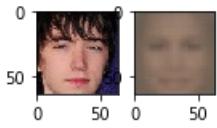
26%|#| ETA: 0:01:40 Epoch: 2 Iter: 797 Class Loss: 0.00 Loss: 0.03



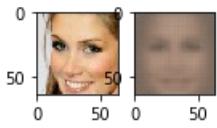
29%|# |ETA: 0:01:37 Epoch: 2 Iter: 899 Class Loss: 0.00 Loss: 0.04



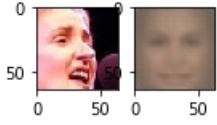
32%|# |ETA: 0:01:30 Epoch: 2 Iter: 999 Class Loss: 0.00 Loss: 0.03



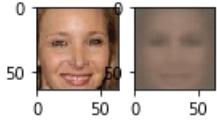
36%|# |ETA: 0:01:26 Epoch: 2 Iter: 1099 Class Loss: 0.01 Loss: 0.04



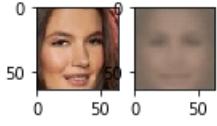
39%|# |ETA: 0:01:22 Epoch: 2 Iter: 1197 Class Loss: 0.00 Loss: 0.04



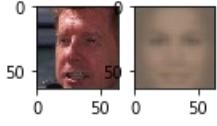
42%|# |ETA: 0:01:18 Epoch: 2 Iter: 1298 Class Loss: 0.00 Loss: 0.03



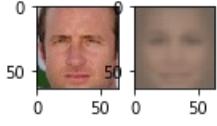
45%|# |ETA: 0:01:26 Epoch: 2 Iter: 1399 Class Loss: 0.00 Loss: 0.03



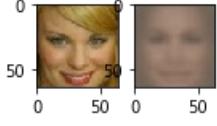
49%|# |ETA: 0:01:09 Epoch: 2 Iter: 1498 Class Loss: 0.00 Loss: 0.03



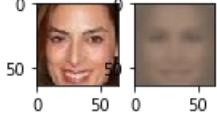
52%## |ETA: 0:01:04 Epoch: 2 Iter: 1599 Class Loss: 0.00 Loss: 0.04



55%## |ETA: 0:01:01 Epoch: 2 Iter: 1699 Class Loss: 0.00 Loss: 0.03



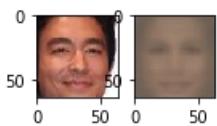
58%## |ETA: 0:01:04 Epoch: 2 Iter: 1798 Class Loss: 0.00 Loss: 0.04



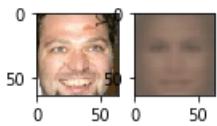
62%## |ETA: 0:00:58 Epoch: 2 Iter: 1899 Class Loss: 0.00 Loss: 0.03



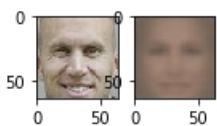
65%|##| ETA: 0:00:47 Epoch: 2 Iter: 1999 Class Loss: 0.00 Loss: 0.04



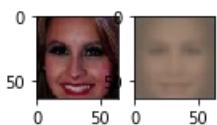
68%|##| ETA: 0:00:42 Epoch: 2 Iter: 2097 Class Loss: 0.04 Loss: 0.07



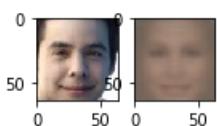
72%|##| ETA: 0:00:38 Epoch: 2 Iter: 2199 Class Loss: 0.00 Loss: 0.03



75%|###| ETA: 0:00:33 Epoch: 2 Iter: 2298 Class Loss: 0.00 Loss: 0.04



78%|###| ETA: 0:00:29 Epoch: 2 Iter: 2399 Class Loss: 0.00 Loss: 0.03



81%|###| ETA: 0:00:24 Epoch: 2 Iter: 2497 Class Loss: 0.00 Loss: 0.03



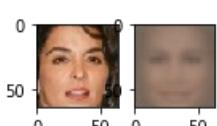
85%|###| ETA: 0:00:20 Epoch: 2 Iter: 2597 Class Loss: 0.00 Loss: 0.03



88%|###| ETA: 0:00:15 Epoch: 2 Iter: 2699 Class Loss: 0.00 Loss: 0.03



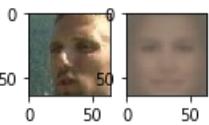
91%|###| ETA: 0:00:11 Epoch: 2 Iter: 2799 Class Loss: 0.00 Loss: 0.04



94%|###| ETA: 0:00:06 Epoch: 2 Iter: 2898 Class Loss: 0.00 Loss: 0.04



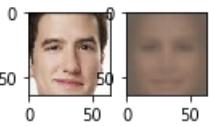
98%|###| ETA: 0:00:02 Epoch: 2 Iter: 2999 Class Loss: 0.00 Loss: 0.03



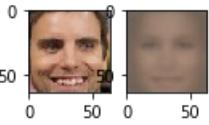
100%|####|Time: 0:02:24 Epoch: 2 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

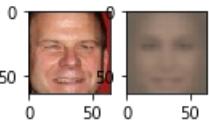
N/A%| |ETA: ---- Epoch: nan Iter: nan Class Loss: nan Loss: nan



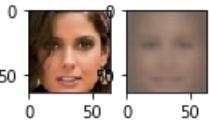
3%| |ETA: 0:02:12 Epoch: 3 Iter: 98 Class Loss: 0.00 Loss: 0.03



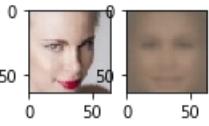
6%| |ETA: 0:02:06 Epoch: 3 Iter: 198 Class Loss: 0.00 Loss: 0.03



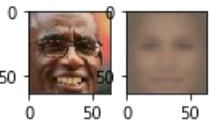
9%| |ETA: 0:02:02 Epoch: 3 Iter: 299 Class Loss: 0.00 Loss: 0.03



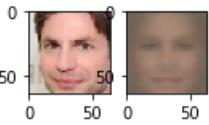
13%| |ETA: 0:02:14 Epoch: 3 Iter: 397 Class Loss: 0.00 Loss: 0.04



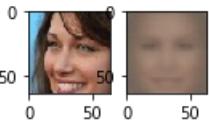
16%| |ETA: 0:02:11 Epoch: 3 Iter: 498 Class Loss: 0.00 Loss: 0.03



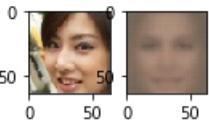
19%| |ETA: 0:01:49 Epoch: 3 Iter: 597 Class Loss: 0.00 Loss: 0.04



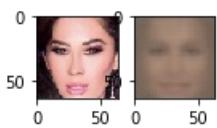
22%|# |ETA: 0:01:44 Epoch: 3 Iter: 698 Class Loss: 0.00 Loss: 0.04



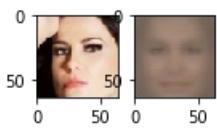
26%|# |ETA: 0:01:40 Epoch: 3 Iter: 799 Class Loss: 0.01 Loss: 0.04



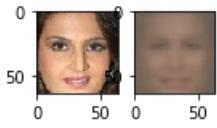
29%|# |ETA: 0:01:37 Epoch: 3 Iter: 899 Class Loss: 0.00 Loss: 0.03



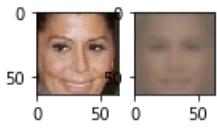
32%|# |ETA: 0:01:31 Epoch: 3 Iter: 999 Class Loss: 0.00 Loss: 0.03



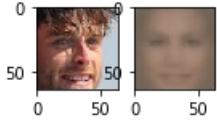
36%|# |ETA: 0:01:26 Epoch: 3 Iter: 1099 Class Loss: 0.00 Loss: 0.03



39%|# |ETA: 0:01:22 Epoch: 3 Iter: 1197 Class Loss: 0.00 Loss: 0.03



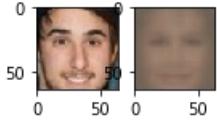
42%|# |ETA: 0:01:18 Epoch: 3 Iter: 1298 Class Loss: 0.00 Loss: 0.03



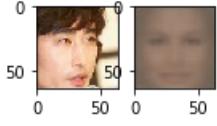
45%|# |ETA: 0:01:14 Epoch: 3 Iter: 1397 Class Loss: 0.00 Loss: 0.04



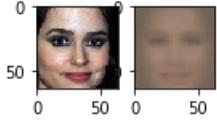
49%|# |ETA: 0:01:19 Epoch: 3 Iter: 1499 Class Loss: 0.00 Loss: 0.03



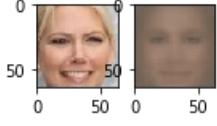
52%|## |ETA: 0:01:18 Epoch: 3 Iter: 1598 Class Loss: 0.00 Loss: 0.04



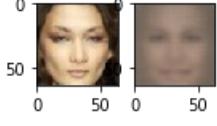
55%|## |ETA: 0:01:00 Epoch: 3 Iter: 1699 Class Loss: 0.02 Loss: 0.05



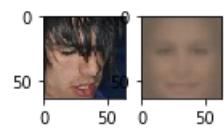
58%|## |ETA: 0:00:56 Epoch: 3 Iter: 1798 Class Loss: 0.00 Loss: 0.03



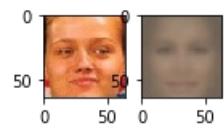
62%|## |ETA: 0:00:51 Epoch: 3 Iter: 1899 Class Loss: 0.00 Loss: 0.03



65%|## |ETA: 0:00:46 Epoch: 3 Iter: 1997 Class Loss: 0.00 Loss: 0.03



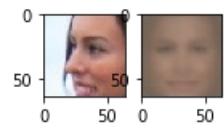
68%|##| ETA: 0:00:47 Epoch: 3 Iter: 2098 Class Loss: 0.00 Loss: 0.03



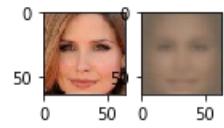
72%|##| ETA: 0:00:43 Epoch: 3 Iter: 2198 Class Loss: 0.00 Loss: 0.03



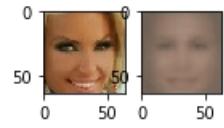
75%|###| ETA: 0:00:33 Epoch: 3 Iter: 2297 Class Loss: 0.00 Loss: 0.03



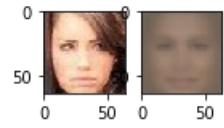
78%|###| ETA: 0:00:29 Epoch: 3 Iter: 2398 Class Loss: 0.04 Loss: 0.07



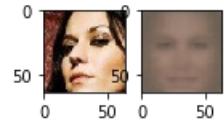
81%|###| ETA: 0:00:24 Epoch: 3 Iter: 2499 Class Loss: 0.00 Loss: 0.03



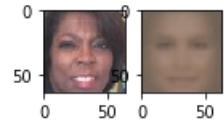
85%|###| ETA: 0:00:20 Epoch: 3 Iter: 2597 Class Loss: 0.04 Loss: 0.08



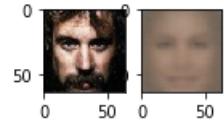
88%|###| ETA: 0:00:15 Epoch: 3 Iter: 2699 Class Loss: 0.00 Loss: 0.03



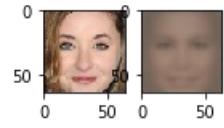
91%|###| ETA: 0:00:11 Epoch: 3 Iter: 2797 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 3 Iter: 2898 Class Loss: 0.02 Loss: 0.05



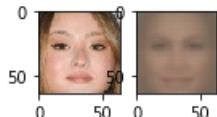
98%|###| ETA: 0:00:02 Epoch: 3 Iter: 2998 Class Loss: 0.00 Loss: 0.04



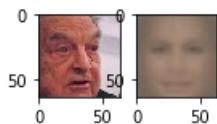
100%|####| Time: 0:02:25 Epoch: 3 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

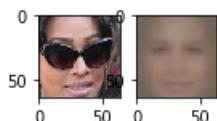
N/A%| |ETA: --:--- Epoch: nan Iter: nan Class Loss: nan Loss: nan



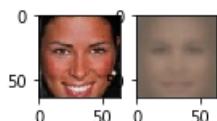
3%| |ETA: 0:02:10 Epoch: 4 Iter: 97 Class Loss: 0.00 Loss: 0.03



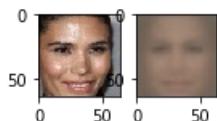
6%| |ETA: 0:02:06 Epoch: 4 Iter: 199 Class Loss: 0.00 Loss: 0.03



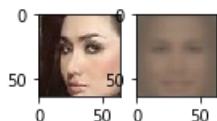
9%| |ETA: 0:02:02 Epoch: 4 Iter: 297 Class Loss: 0.00 Loss: 0.03



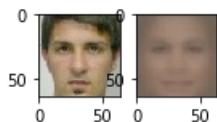
13%| |ETA: 0:01:57 Epoch: 4 Iter: 397 Class Loss: 0.00 Loss: 0.04



16%| |ETA: 0:01:53 Epoch: 4 Iter: 498 Class Loss: 0.00 Loss: 0.03



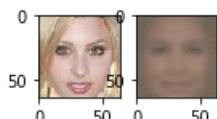
19%| |ETA: 0:01:48 Epoch: 4 Iter: 598 Class Loss: 0.00 Loss: 0.03



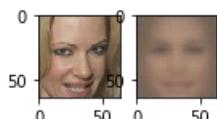
22%|# |ETA: 0:01:59 Epoch: 4 Iter: 699 Class Loss: 0.01 Loss: 0.04



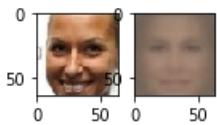
26%|# |ETA: 0:01:52 Epoch: 4 Iter: 799 Class Loss: 0.00 Loss: 0.03



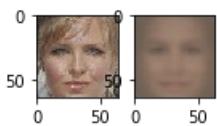
29%|# |ETA: 0:01:36 Epoch: 4 Iter: 897 Class Loss: 0.08 Loss: 0.11



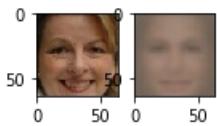
32%|# |ETA: 0:01:31 Epoch: 4 Iter: 997 Class Loss: 0.00 Loss: 0.04



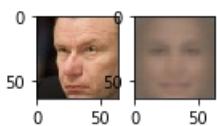
36%|# |ETA: 0:01:26 Epoch: 4 Iter: 1099 Class Loss: 0.00 Loss: 0.03



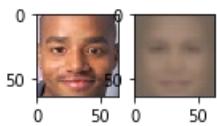
39%|# |ETA: 0:01:22 Epoch: 4 Iter: 1198 Class Loss: 0.00 Loss: 0.03



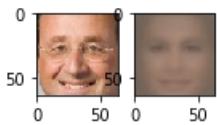
42%|# |ETA: 0:01:18 Epoch: 4 Iter: 1298 Class Loss: 0.00 Loss: 0.03



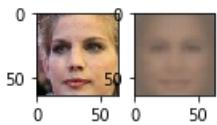
45%|# |ETA: 0:01:13 Epoch: 4 Iter: 1398 Class Loss: 0.00 Loss: 0.03



49%|# |ETA: 0:01:08 Epoch: 4 Iter: 1499 Class Loss: 0.01 Loss: 0.04



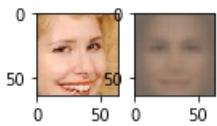
52%|## |ETA: 0:01:04 Epoch: 4 Iter: 1599 Class Loss: 0.00 Loss: 0.03



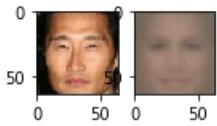
55%|## |ETA: 0:01:10 Epoch: 4 Iter: 1699 Class Loss: 0.00 Loss: 0.03



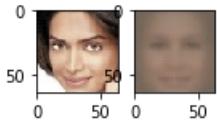
58%|## |ETA: 0:00:56 Epoch: 4 Iter: 1797 Class Loss: 0.00 Loss: 0.04



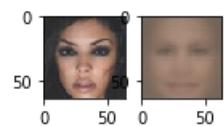
62%|## |ETA: 0:00:50 Epoch: 4 Iter: 1899 Class Loss: 0.00 Loss: 0.03



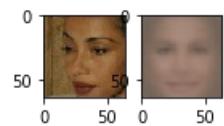
65%|## |ETA: 0:00:46 Epoch: 4 Iter: 1998 Class Loss: 0.00 Loss: 0.04



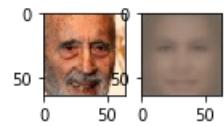
68%|## |ETA: 0:00:42 Epoch: 4 Iter: 2097 Class Loss: 0.00 Loss: 0.03



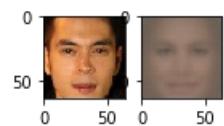
72%|##| ETA: 0:00:38 Epoch: 4 Iter: 2198 Class Loss: 0.00 Loss: 0.04



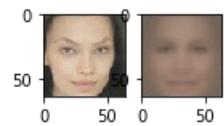
75%|###| ETA: 0:00:33 Epoch: 4 Iter: 2298 Class Loss: 0.00 Loss: 0.03



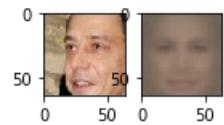
78%|###| ETA: 0:00:33 Epoch: 4 Iter: 2398 Class Loss: 0.00 Loss: 0.03



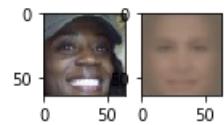
81%|###| ETA: 0:00:27 Epoch: 4 Iter: 2498 Class Loss: 0.00 Loss: 0.03



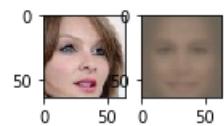
85%|###| ETA: 0:00:20 Epoch: 4 Iter: 2597 Class Loss: 0.00 Loss: 0.04



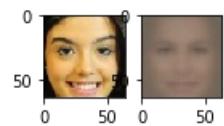
88%|###| ETA: 0:00:15 Epoch: 4 Iter: 2697 Class Loss: 0.00 Loss: 0.03



91%|###| ETA: 0:00:11 Epoch: 4 Iter: 2798 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 4 Iter: 2898 Class Loss: 0.00 Loss: 0.03



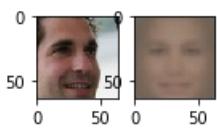
98%|###| ETA: 0:00:02 Epoch: 4 Iter: 2998 Class Loss: 0.00 Loss: 0.04



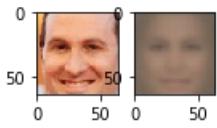
100%|####| Time: 0:02:23 Epoch: 4 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

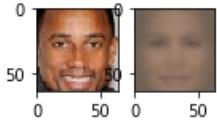
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



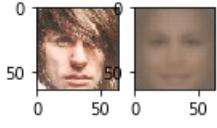
3% | ETA: 0:02:10 Epoch: 5 Iter: 98 Class Loss: 0.00 Loss: 0.03



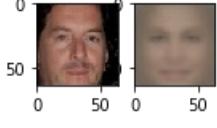
6% | ETA: 0:02:07 Epoch: 5 Iter: 197 Class Loss: 0.00 Loss: 0.03



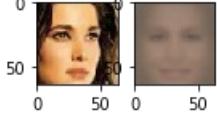
9% | ETA: 0:02:02 Epoch: 5 Iter: 299 Class Loss: 0.00 Loss: 0.03



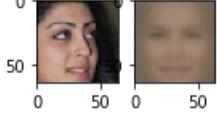
13% | ETA: 0:01:57 Epoch: 5 Iter: 399 Class Loss: 0.00 Loss: 0.03



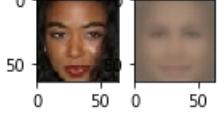
16% | ETA: 0:01:54 Epoch: 5 Iter: 499 Class Loss: 0.00 Loss: 0.04



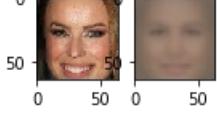
19% | ETA: 0:01:48 Epoch: 5 Iter: 597 Class Loss: 0.04 Loss: 0.07



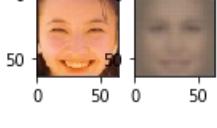
22% | # ETA: 0:01:44 Epoch: 5 Iter: 697 Class Loss: 0.00 Loss: 0.04



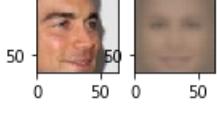
26% | # ETA: 0:01:40 Epoch: 5 Iter: 798 Class Loss: 0.00 Loss: 0.03



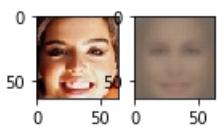
29% | # ETA: 0:01:36 Epoch: 5 Iter: 897 Class Loss: 0.00 Loss: 0.03



32% | # ETA: 0:01:44 Epoch: 5 Iter: 998 Class Loss: 0.00 Loss: 0.03



35% | # ETA: 0:01:38 Epoch: 5 Iter: 1098 Class Loss: 0.00 Loss: 0.03



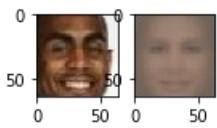
39%|# |ETA: 0:01:23 Epoch: 5 Iter: 1198 Class Loss: 0.00 Loss: 0.03



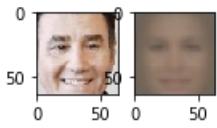
42%|# |ETA: 0:01:17 Epoch: 5 Iter: 1298 Class Loss: 0.00 Loss: 0.04



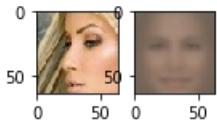
45%|# |ETA: 0:01:14 Epoch: 5 Iter: 1397 Class Loss: 0.00 Loss: 0.04



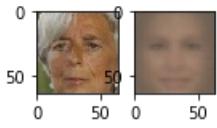
49%|# |ETA: 0:01:09 Epoch: 5 Iter: 1497 Class Loss: 0.00 Loss: 0.03



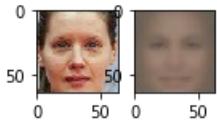
52%|## |ETA: 0:01:04 Epoch: 5 Iter: 1599 Class Loss: 0.00 Loss: 0.04



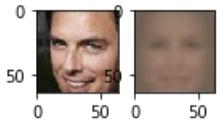
55%|## |ETA: 0:01:00 Epoch: 5 Iter: 1697 Class Loss: 0.00 Loss: 0.03



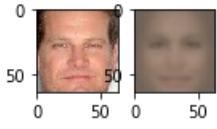
58%|## |ETA: 0:00:56 Epoch: 5 Iter: 1799 Class Loss: 0.00 Loss: 0.03



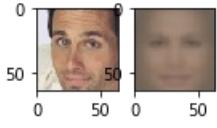
62%|## |ETA: 0:01:02 Epoch: 5 Iter: 1899 Class Loss: 0.00 Loss: 0.04



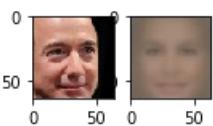
65%|## |ETA: 0:00:56 Epoch: 5 Iter: 1998 Class Loss: 0.00 Loss: 0.03



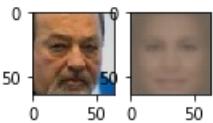
68%|## |ETA: 0:00:42 Epoch: 5 Iter: 2098 Class Loss: 0.00 Loss: 0.03



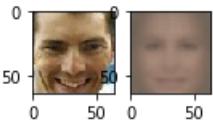
72%|## |ETA: 0:00:38 Epoch: 5 Iter: 2199 Class Loss: 0.02 Loss: 0.05



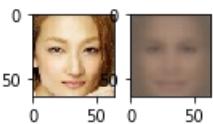
75%|###| ETA: 0:00:33 Epoch: 5 Iter: 2297 Class Loss: 0.00 Loss: 0.03



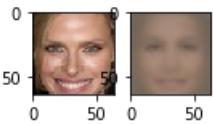
78%|###| ETA: 0:00:29 Epoch: 5 Iter: 2398 Class Loss: 0.00 Loss: 0.03



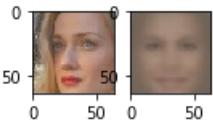
81%|###| ETA: 0:00:24 Epoch: 5 Iter: 2498 Class Loss: 0.00 Loss: 0.04



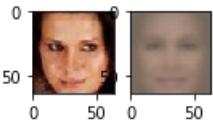
85%|###| ETA: 0:00:20 Epoch: 5 Iter: 2599 Class Loss: 0.01 Loss: 0.04



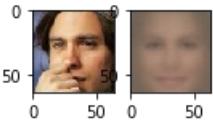
88%|###| ETA: 0:00:18 Epoch: 5 Iter: 2698 Class Loss: 0.00 Loss: 0.03



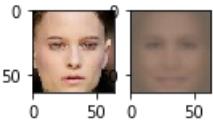
91%|###| ETA: 0:00:12 Epoch: 5 Iter: 2797 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 5 Iter: 2898 Class Loss: 0.00 Loss: 0.03



98%|###| ETA: 0:00:02 Epoch: 5 Iter: 2998 Class Loss: 0.00 Loss: 0.03



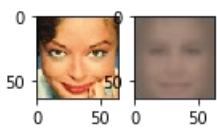
100%|####| Time: 0:02:25 Epoch: 5 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

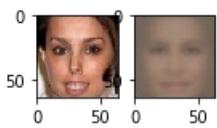
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



3%| | ETA: 0:02:11 Epoch: 6 Iter: 98 Class Loss: 0.00 Loss: 0.04



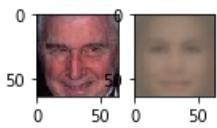
6% | ETA: 0:02:06 Epoch: 6 Iter: 198 Class Loss: 0.01 Loss: 0.04



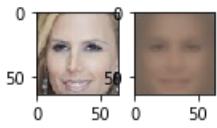
9% | ETA: 0:02:02 Epoch: 6 Iter: 298 Class Loss: 0.00 Loss: 0.04



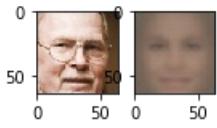
13% | ETA: 0:01:58 Epoch: 6 Iter: 398 Class Loss: 0.00 Loss: 0.04



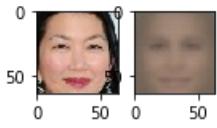
16% | ETA: 0:01:54 Epoch: 6 Iter: 497 Class Loss: 0.00 Loss: 0.03



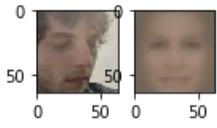
19% | ETA: 0:01:48 Epoch: 6 Iter: 597 Class Loss: 0.00 Loss: 0.03



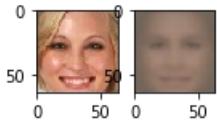
22% | # ETA: 0:01:44 Epoch: 6 Iter: 697 Class Loss: 0.10 Loss: 0.13



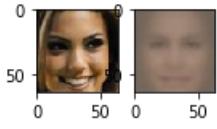
26% | # ETA: 0:01:39 Epoch: 6 Iter: 799 Class Loss: 0.00 Loss: 0.03



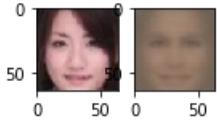
29% | # ETA: 0:01:36 Epoch: 6 Iter: 898 Class Loss: 0.00 Loss: 0.04



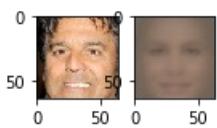
32% | # ETA: 0:01:31 Epoch: 6 Iter: 998 Class Loss: 0.00 Loss: 0.03



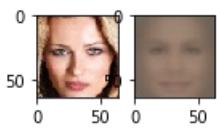
36% | # ETA: 0:01:26 Epoch: 6 Iter: 1099 Class Loss: 0.00 Loss: 0.03



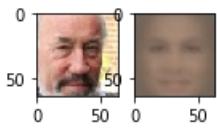
39% | # ETA: 0:01:22 Epoch: 6 Iter: 1198 Class Loss: 0.00 Loss: 0.04



42%|# |ETA: 0:01:29 Epoch: 6 Iter: 1299 Class Loss: 0.11 Loss: 0.14



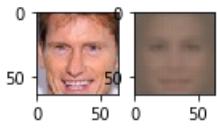
45%|# |ETA: 0:01:24 Epoch: 6 Iter: 1399 Class Loss: 0.00 Loss: 0.04



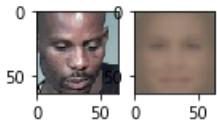
49%|# |ETA: 0:01:08 Epoch: 6 Iter: 1497 Class Loss: 0.00 Loss: 0.03



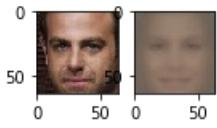
52%|## |ETA: 0:01:04 Epoch: 6 Iter: 1599 Class Loss: 0.00 Loss: 0.03



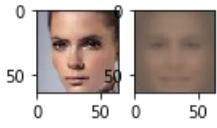
55%|## |ETA: 0:01:00 Epoch: 6 Iter: 1699 Class Loss: 0.00 Loss: 0.03



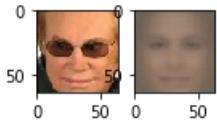
58%|## |ETA: 0:00:55 Epoch: 6 Iter: 1798 Class Loss: 0.00 Loss: 0.03



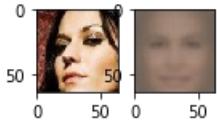
62%|## |ETA: 0:00:55 Epoch: 6 Iter: 1899 Class Loss: 0.00 Loss: 0.03



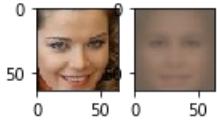
65%|## |ETA: 0:00:52 Epoch: 6 Iter: 1997 Class Loss: 0.00 Loss: 0.03



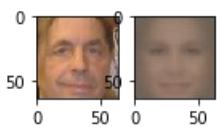
68%|## |ETA: 0:00:42 Epoch: 6 Iter: 2098 Class Loss: 0.00 Loss: 0.04



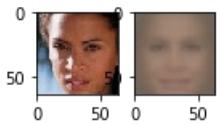
71%|## |ETA: 0:00:37 Epoch: 6 Iter: 2197 Class Loss: 0.02 Loss: 0.05



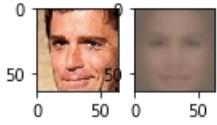
75%|### |ETA: 0:00:34 Epoch: 6 Iter: 2297 Class Loss: 0.00 Loss: 0.03



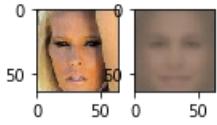
78%|###| ETA: 0:00:29 Epoch: 6 Iter: 2398 Class Loss: 0.05 Loss: 0.08



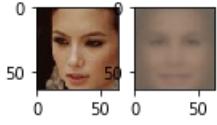
81%|###| ETA: 0:00:24 Epoch: 6 Iter: 2499 Class Loss: 0.00 Loss: 0.03



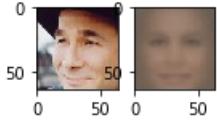
85%|###| ETA: 0:00:20 Epoch: 6 Iter: 2598 Class Loss: 0.00 Loss: 0.03



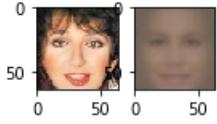
88%|###| ETA: 0:00:15 Epoch: 6 Iter: 2698 Class Loss: 0.00 Loss: 0.03



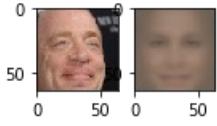
91%|###| ETA: 0:00:11 Epoch: 6 Iter: 2797 Class Loss: 0.08 Loss: 0.11



94%|###| ETA: 0:00:06 Epoch: 6 Iter: 2898 Class Loss: 0.00 Loss: 0.03



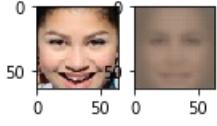
98%|###| ETA: 0:00:02 Epoch: 6 Iter: 2997 Class Loss: 0.00 Loss: 0.04



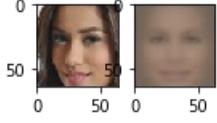
100%|####| Time: 0:02:23 Epoch: 6 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

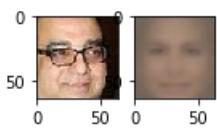
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



3%| | ETA: 0:02:13 Epoch: 7 Iter: 98 Class Loss: 0.00 Loss: 0.03



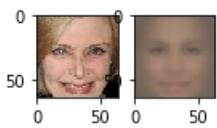
6%| | ETA: 0:02:06 Epoch: 7 Iter: 199 Class Loss: 0.00 Loss: 0.03



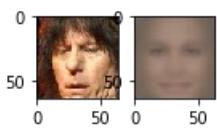
9% | ETA: 0:02:01 Epoch: 7 Iter: 299 Class Loss: 0.00 Loss: 0.04



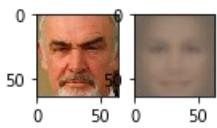
13% | ETA: 0:01:58 Epoch: 7 Iter: 398 Class Loss: 0.02 Loss: 0.05



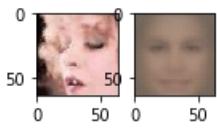
16% | ETA: 0:01:53 Epoch: 7 Iter: 499 Class Loss: 0.00 Loss: 0.04



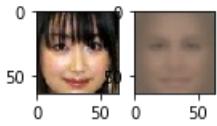
19% | ETA: 0:01:48 Epoch: 7 Iter: 597 Class Loss: 0.00 Loss: 0.04



22% | # ETA: 0:01:45 Epoch: 7 Iter: 698 Class Loss: 0.00 Loss: 0.03



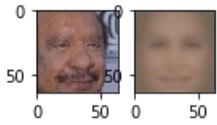
26% | # ETA: 0:01:40 Epoch: 7 Iter: 797 Class Loss: 0.00 Loss: 0.03



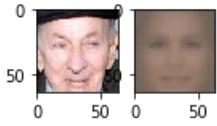
29% | # ETA: 0:01:35 Epoch: 7 Iter: 899 Class Loss: 0.00 Loss: 0.03



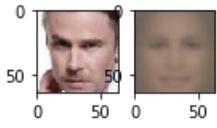
32% | # ETA: 0:01:33 Epoch: 7 Iter: 998 Class Loss: 0.00 Loss: 0.03



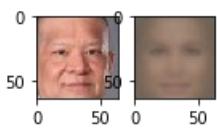
36% | # ETA: 0:01:26 Epoch: 7 Iter: 1099 Class Loss: 0.00 Loss: 0.03



39% | # ETA: 0:01:22 Epoch: 7 Iter: 1199 Class Loss: 0.00 Loss: 0.04



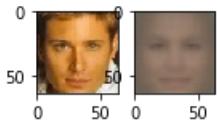
42% | # ETA: 0:01:18 Epoch: 7 Iter: 1298 Class Loss: 0.00 Loss: 0.03



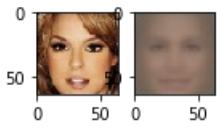
45%|# |ETA: 0:01:13 Epoch: 7 Iter: 1399 Class Loss: 0.00 Loss: 0.03



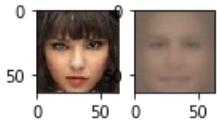
49%|# |ETA: 0:01:09 Epoch: 7 Iter: 1498 Class Loss: 0.00 Loss: 0.03



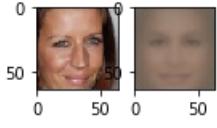
52%|## |ETA: 0:01:10 Epoch: 7 Iter: 1597 Class Loss: 0.00 Loss: 0.03



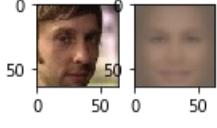
55%|## |ETA: 0:01:08 Epoch: 7 Iter: 1697 Class Loss: 0.00 Loss: 0.03



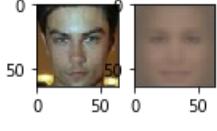
58%|## |ETA: 0:00:57 Epoch: 7 Iter: 1797 Class Loss: 0.00 Loss: 0.03



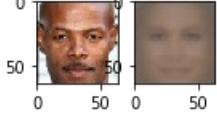
62%|## |ETA: 0:00:51 Epoch: 7 Iter: 1898 Class Loss: 0.00 Loss: 0.03



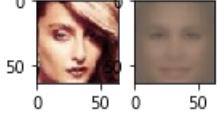
65%|## |ETA: 0:00:47 Epoch: 7 Iter: 1998 Class Loss: 0.00 Loss: 0.03



68%|## |ETA: 0:00:42 Epoch: 7 Iter: 2098 Class Loss: 0.00 Loss: 0.04



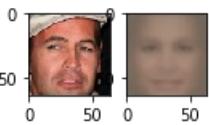
71%|## |ETA: 0:00:37 Epoch: 7 Iter: 2197 Class Loss: 0.00 Loss: 0.03



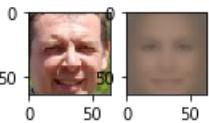
75%|### |ETA: 0:00:40 Epoch: 7 Iter: 2298 Class Loss: 0.00 Loss: 0.03



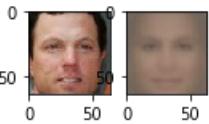
78%|### |ETA: 0:00:34 Epoch: 7 Iter: 2398 Class Loss: 0.01 Loss: 0.05



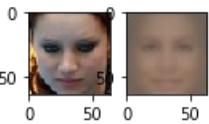
81%|###|ETA: 0:00:24 Epoch: 7 Iter: 2499 Class Loss: 0.00 Loss: 0.03



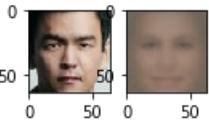
85%|###|ETA: 0:00:20 Epoch: 7 Iter: 2598 Class Loss: 0.00 Loss: 0.04



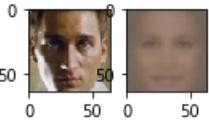
88%|###|ETA: 0:00:15 Epoch: 7 Iter: 2698 Class Loss: 0.00 Loss: 0.03



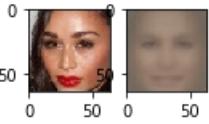
91%|###|ETA: 0:00:11 Epoch: 7 Iter: 2799 Class Loss: 0.00 Loss: 0.03



94%|###|ETA: 0:00:06 Epoch: 7 Iter: 2898 Class Loss: 0.00 Loss: 0.03



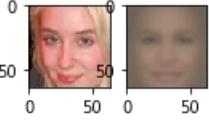
98%|###|ETA: 0:00:02 Epoch: 7 Iter: 2999 Class Loss: 0.00 Loss: 0.03



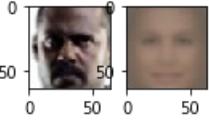
100%|####|Time: 0:02:23 Epoch: 7 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

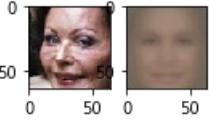
N/A%| |ETA: ---- Epoch: nan Iter: nan Class Loss: nan Loss: nan



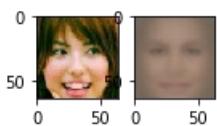
3%| |ETA: 0:02:12 Epoch: 8 Iter: 99 Class Loss: 0.00 Loss: 0.03



6%| |ETA: 0:02:23 Epoch: 8 Iter: 197 Class Loss: 0.00 Loss: 0.03



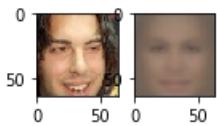
9%| |ETA: 0:02:19 Epoch: 8 Iter: 299 Class Loss: 0.00 Loss: 0.03



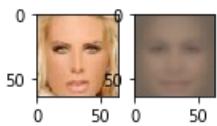
13% | ETA: 0:01:58 Epoch: 8 Iter: 397 Class Loss: 0.00 Loss: 0.03



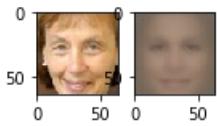
16% | ETA: 0:01:53 Epoch: 8 Iter: 498 Class Loss: 0.00 Loss: 0.04



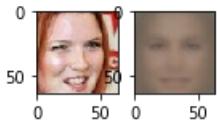
19% | ETA: 0:01:49 Epoch: 8 Iter: 599 Class Loss: 0.00 Loss: 0.03



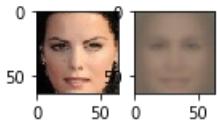
22% | # ETA: 0:01:45 Epoch: 8 Iter: 697 Class Loss: 0.00 Loss: 0.04



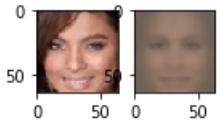
26% | # ETA: 0:01:40 Epoch: 8 Iter: 799 Class Loss: 0.00 Loss: 0.03



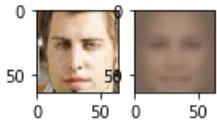
29% | # ETA: 0:01:35 Epoch: 8 Iter: 899 Class Loss: 0.00 Loss: 0.03



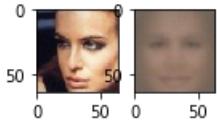
32% | # ETA: 0:01:32 Epoch: 8 Iter: 997 Class Loss: 0.00 Loss: 0.03



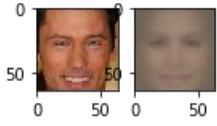
35% | # ETA: 0:01:27 Epoch: 8 Iter: 1098 Class Loss: 0.00 Loss: 0.03



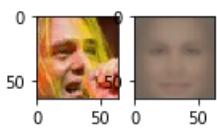
39% | # ETA: 0:01:22 Epoch: 8 Iter: 1198 Class Loss: 0.00 Loss: 0.04



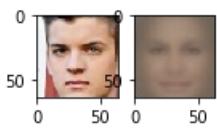
42% | # ETA: 0:01:17 Epoch: 8 Iter: 1297 Class Loss: 0.00 Loss: 0.03



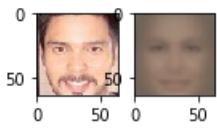
45% | # ETA: 0:01:13 Epoch: 8 Iter: 1399 Class Loss: 0.00 Loss: 0.03



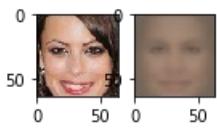
49%|# |ETA: 0:01:08 Epoch: 8 Iter: 1498 Class Loss: 0.00 Loss: 0.04



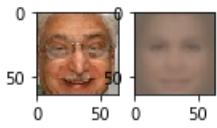
52%|## |ETA: 0:01:04 Epoch: 8 Iter: 1598 Class Loss: 0.00 Loss: 0.04



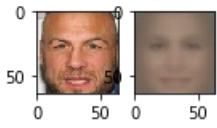
55%|## |ETA: 0:01:00 Epoch: 8 Iter: 1699 Class Loss: 0.00 Loss: 0.03



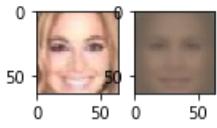
58%|## |ETA: 0:00:55 Epoch: 8 Iter: 1798 Class Loss: 0.00 Loss: 0.04



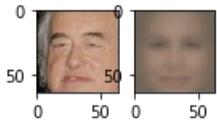
62%|## |ETA: 0:00:58 Epoch: 8 Iter: 1898 Class Loss: 0.00 Loss: 0.04



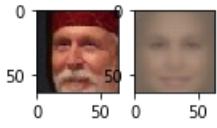
65%|## |ETA: 0:00:53 Epoch: 8 Iter: 1998 Class Loss: 0.03 Loss: 0.07



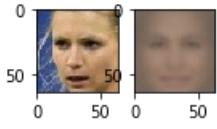
68%|## |ETA: 0:00:43 Epoch: 8 Iter: 2098 Class Loss: 0.08 Loss: 0.11



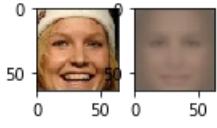
72%|## |ETA: 0:00:44 Epoch: 8 Iter: 2198 Class Loss: 0.00 Loss: 0.03



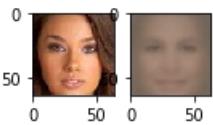
75%|### |ETA: 0:00:34 Epoch: 8 Iter: 2297 Class Loss: 0.00 Loss: 0.03



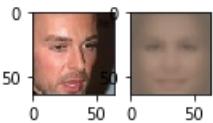
78%|### |ETA: 0:00:29 Epoch: 8 Iter: 2398 Class Loss: 0.08 Loss: 0.12



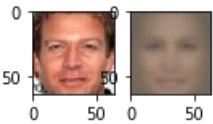
81%|### |ETA: 0:00:24 Epoch: 8 Iter: 2498 Class Loss: 0.00 Loss: 0.04



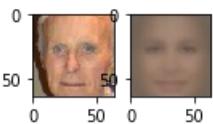
85%|###|ETA: 0:00:20 Epoch: 8 Iter: 2598 Class Loss: 0.00 Loss: 0.04



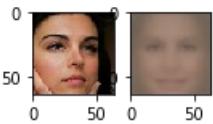
88%|###|ETA: 0:00:15 Epoch: 8 Iter: 2698 Class Loss: 0.00 Loss: 0.03



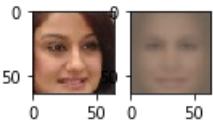
91%|###|ETA: 0:00:11 Epoch: 8 Iter: 2799 Class Loss: 0.00 Loss: 0.04



94%|###|ETA: 0:00:06 Epoch: 8 Iter: 2898 Class Loss: 0.01 Loss: 0.03



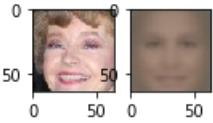
98%|###|ETA: 0:00:02 Epoch: 8 Iter: 2999 Class Loss: 0.00 Loss: 0.04



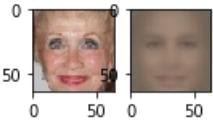
100%|####|Time: 0:02:24 Epoch: 8 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

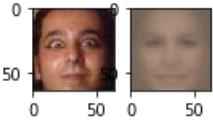
N/A%| |ETA: ---- Epoch: nan Iter: nan Class Loss: nan Loss: nan



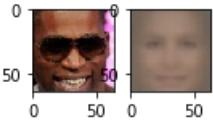
3%| |ETA: 0:02:12 Epoch: 9 Iter: 98 Class Loss: 0.01 Loss: 0.04



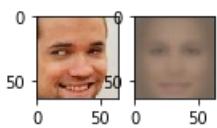
6%| |ETA: 0:02:06 Epoch: 9 Iter: 198 Class Loss: 0.00 Loss: 0.03



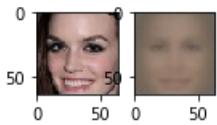
9%| |ETA: 0:02:02 Epoch: 9 Iter: 297 Class Loss: 0.01 Loss: 0.04



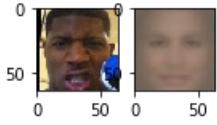
13%| |ETA: 0:01:57 Epoch: 9 Iter: 397 Class Loss: 0.00 Loss: 0.04



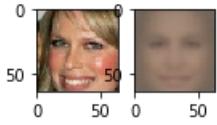
16% | ETA: 0:02:07 Epoch: 9 Iter: 498 Class Loss: 0.00 Loss: 0.03



19% | ETA: 0:02:06 Epoch: 9 Iter: 598 Class Loss: 0.01 Loss: 0.04



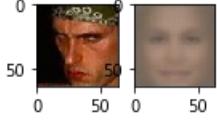
22% |# ETA: 0:01:44 Epoch: 9 Iter: 698 Class Loss: 0.00 Loss: 0.04



26% |# ETA: 0:01:40 Epoch: 9 Iter: 799 Class Loss: 0.00 Loss: 0.03



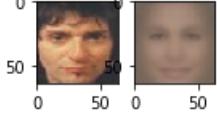
29% |# ETA: 0:01:35 Epoch: 9 Iter: 897 Class Loss: 0.00 Loss: 0.03



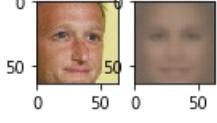
32% |# ETA: 0:01:33 Epoch: 9 Iter: 998 Class Loss: 0.00 Loss: 0.03



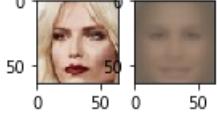
35% |# ETA: 0:01:27 Epoch: 9 Iter: 1097 Class Loss: 0.00 Loss: 0.04



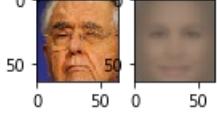
39% |# ETA: 0:01:22 Epoch: 9 Iter: 1198 Class Loss: 0.00 Loss: 0.04



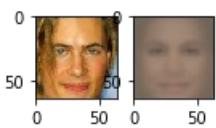
42% |# ETA: 0:01:18 Epoch: 9 Iter: 1299 Class Loss: 0.00 Loss: 0.03



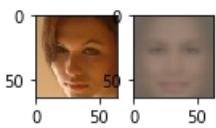
45% |# ETA: 0:01:14 Epoch: 9 Iter: 1398 Class Loss: 0.00 Loss: 0.03



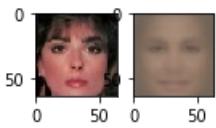
49% |# ETA: 0:01:10 Epoch: 9 Iter: 1499 Class Loss: 0.00 Loss: 0.03



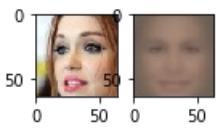
52%|##| ETA: 0:01:04 Epoch: 9 Iter: 1599 Class Loss: 0.00 Loss: 0.03



55%|##| ETA: 0:01:00 Epoch: 9 Iter: 1697 Class Loss: 0.00 Loss: 0.03



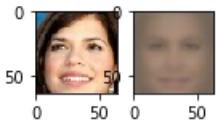
58%|##| ETA: 0:00:55 Epoch: 9 Iter: 1798 Class Loss: 0.01 Loss: 0.04



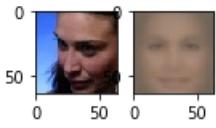
62%|##| ETA: 0:00:52 Epoch: 9 Iter: 1899 Class Loss: 0.00 Loss: 0.04



65%|##| ETA: 0:00:46 Epoch: 9 Iter: 1997 Class Loss: 0.00 Loss: 0.03



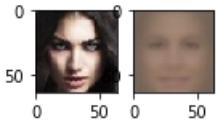
68%|##| ETA: 0:00:42 Epoch: 9 Iter: 2097 Class Loss: 0.00 Loss: 0.03



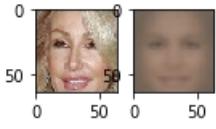
72%|##| ETA: 0:00:40 Epoch: 9 Iter: 2198 Class Loss: 0.00 Loss: 0.03



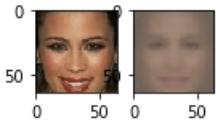
75%|###| ETA: 0:00:38 Epoch: 9 Iter: 2297 Class Loss: 0.00 Loss: 0.03



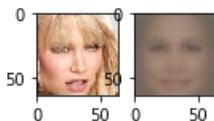
78%|###| ETA: 0:00:30 Epoch: 9 Iter: 2399 Class Loss: 0.00 Loss: 0.03



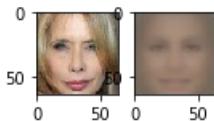
81%|###| ETA: 0:00:24 Epoch: 9 Iter: 2499 Class Loss: 0.00 Loss: 0.03



85%|###| ETA: 0:00:21 Epoch: 9 Iter: 2599 Class Loss: 0.00 Loss: 0.03



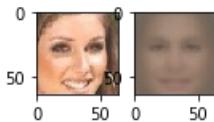
88%|###| ETA: 0:00:18 Epoch: 9 Iter: 2698 Class Loss: 0.00 Loss: 0.03



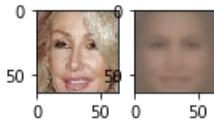
91%|###| ETA: 0:00:12 Epoch: 9 Iter: 2797 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 9 Iter: 2898 Class Loss: 0.00 Loss: 0.03



98%|###| ETA: 0:00:02 Epoch: 9 Iter: 2998 Class Loss: 0.00 Loss: 0.03



100%|####| Time: 0:02:24 Epoch: 9 Iter: 3052 Class Loss: 0.00 Loss: 0.04
100% (97 of 97) |#####| Elapsed Time: 0:01:09 Time: 0:01:09
N/A% (0 of 72) | Elapsed Time: 0:00:00 ETA: --:--:--

male lighter: 0.979381443299

100% (72 of 72) |#####| Elapsed Time: 0:00:49 Time: 0:00:49
N/A% (0 of 78) | Elapsed Time: 0:00:00 ETA: --:--:--

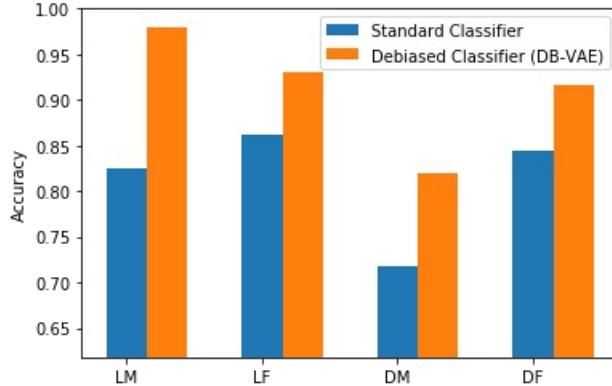
female lighter: 0.930555555556

100% (78 of 78) |#####| Elapsed Time: 0:00:52 Time: 0:00:52
N/A% (0 of 71) | Elapsed Time: 0:00:00 ETA: --:--:--

male darker: 0.820512820513

100% (71 of 71) |#####| Elapsed Time: 0:00:48 Time: 0:00:48

female darker: 0.915492957746



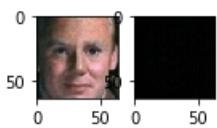
Smooth Factor of 1

In [63]:

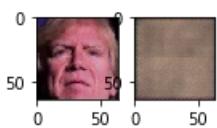
```
run_smooth_factor(1)
```

Recomputing the sampling probabilities

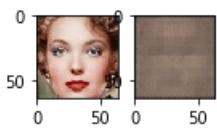
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



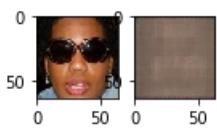
3% | ETA: 0:02:09 Epoch: 0 Iter: 97 Class Loss: 0.13 Loss: 0.17



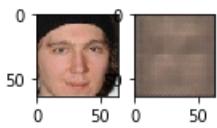
6% | ETA: 0:02:06 Epoch: 0 Iter: 199 Class Loss: 0.08 Loss: 0.11



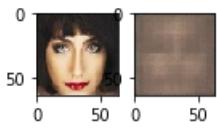
9% | ETA: 0:02:00 Epoch: 0 Iter: 297 Class Loss: 0.18 Loss: 0.22



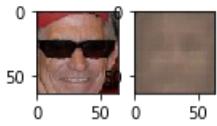
13% | ETA: 0:01:55 Epoch: 0 Iter: 398 Class Loss: 0.13 Loss: 0.17



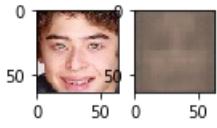
16% | ETA: 0:01:51 Epoch: 0 Iter: 499 Class Loss: 0.16 Loss: 0.20



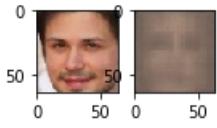
19% | ETA: 0:01:47 Epoch: 0 Iter: 599 Class Loss: 0.01 Loss: 0.06



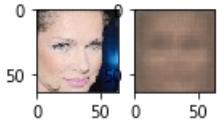
22% | # ETA: 0:01:43 Epoch: 0 Iter: 698 Class Loss: 0.03 Loss: 0.06



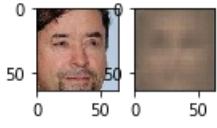
26% | # ETA: 0:01:38 Epoch: 0 Iter: 799 Class Loss: 0.15 Loss: 0.18



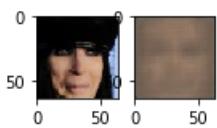
29% | # ETA: 0:01:35 Epoch: 0 Iter: 897 Class Loss: 0.04 Loss: 0.08



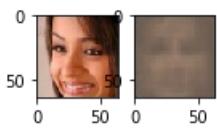
32% | # ETA: 0:01:30 Epoch: 0 Iter: 999 Class Loss: 0.01 Loss: 0.04



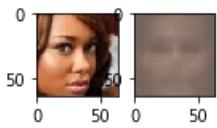
36% | # ETA: 0:01:26 Epoch: 0 Iter: 1099 Class Loss: 0.00 Loss: 0.04



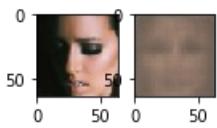
39%|# |ETA: 0:01:21 Epoch: 0 Iter: 1197 Class Loss: 0.01 Loss: 0.04



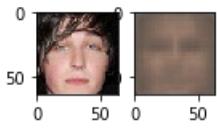
42%|# |ETA: 0:01:26 Epoch: 0 Iter: 1297 Class Loss: 0.00 Loss: 0.04



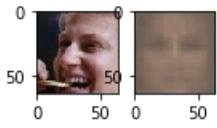
45%|# |ETA: 0:01:27 Epoch: 0 Iter: 1399 Class Loss: 0.00 Loss: 0.04



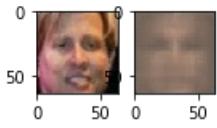
49%|# |ETA: 0:01:22 Epoch: 0 Iter: 1498 Class Loss: 0.01 Loss: 0.05



52%|## |ETA: 0:01:04 Epoch: 0 Iter: 1598 Class Loss: 0.00 Loss: 0.04



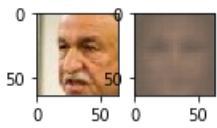
55%|## |ETA: 0:00:59 Epoch: 0 Iter: 1699 Class Loss: 0.00 Loss: 0.04



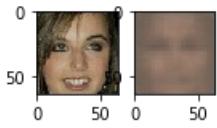
58%|## |ETA: 0:00:54 Epoch: 0 Iter: 1798 Class Loss: 0.00 Loss: 0.04



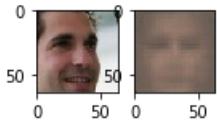
62%|## |ETA: 0:00:50 Epoch: 0 Iter: 1899 Class Loss: 0.00 Loss: 0.04



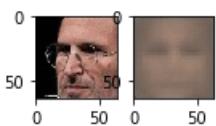
65%|## |ETA: 0:00:46 Epoch: 0 Iter: 1997 Class Loss: 0.00 Loss: 0.04



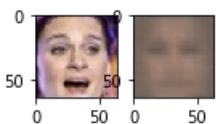
68%|## |ETA: 0:00:41 Epoch: 0 Iter: 2098 Class Loss: 0.01 Loss: 0.05



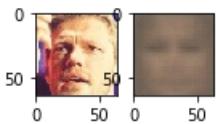
72%|## |ETA: 0:00:37 Epoch: 0 Iter: 2199 Class Loss: 0.00 Loss: 0.03



75%|###| ETA: 0:00:32 Epoch: 0 Iter: 2297 Class Loss: 0.01 Loss: 0.06



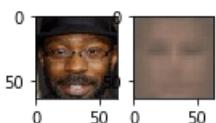
78%|###| ETA: 0:00:28 Epoch: 0 Iter: 2398 Class Loss: 0.00 Loss: 0.04



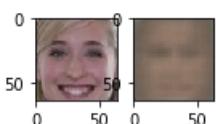
81%|###| ETA: 0:00:24 Epoch: 0 Iter: 2497 Class Loss: 0.00 Loss: 0.03



85%|###| ETA: 0:00:19 Epoch: 0 Iter: 2598 Class Loss: 0.00 Loss: 0.04



88%|###| ETA: 0:00:15 Epoch: 0 Iter: 2698 Class Loss: 0.00 Loss: 0.04



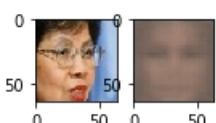
91%|###| ETA: 0:00:11 Epoch: 0 Iter: 2797 Class Loss: 0.07 Loss: 0.10



94%|###| ETA: 0:00:06 Epoch: 0 Iter: 2899 Class Loss: 0.00 Loss: 0.03



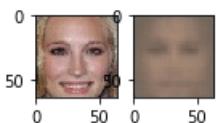
98%|###| ETA: 0:00:02 Epoch: 0 Iter: 2998 Class Loss: 0.00 Loss: 0.04



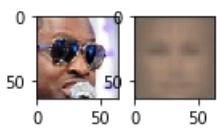
100%|####| Time: 0:02:21 Epoch: 0 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

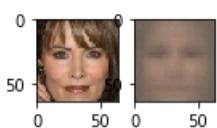
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



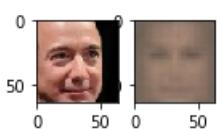
3%| | ETA: 0:02:07 Epoch: 1 Iter: 98 Class Loss: 0.00 Loss: 0.04



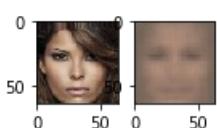
6% | ETA: 0:02:04 Epoch: 1 Iter: 199 Class Loss: 0.00 Loss: 0.03



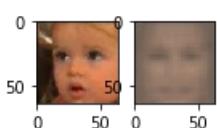
9% | ETA: 0:01:59 Epoch: 1 Iter: 299 Class Loss: 0.00 Loss: 0.03



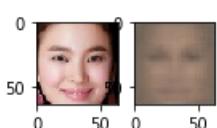
13% | ETA: 0:01:55 Epoch: 1 Iter: 398 Class Loss: 0.00 Loss: 0.03



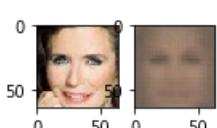
16% | ETA: 0:01:51 Epoch: 1 Iter: 499 Class Loss: 0.01 Loss: 0.05



19% | ETA: 0:01:47 Epoch: 1 Iter: 597 Class Loss: 0.00 Loss: 0.04



22% |# ETA: 0:01:42 Epoch: 1 Iter: 698 Class Loss: 0.09 Loss: 0.12



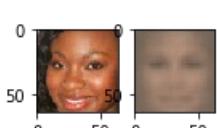
26% |# ETA: 0:01:37 Epoch: 1 Iter: 798 Class Loss: 0.00 Loss: 0.03



29% |# ETA: 0:01:34 Epoch: 1 Iter: 897 Class Loss: 0.00 Loss: 0.03



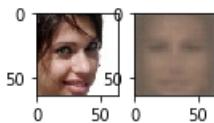
32% |# ETA: 0:01:29 Epoch: 1 Iter: 999 Class Loss: 0.00 Loss: 0.04



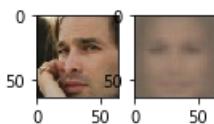
36% |# ETA: 0:01:24 Epoch: 1 Iter: 1099 Class Loss: 0.01 Loss: 0.04



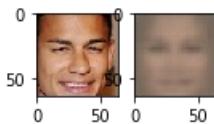
39% |# ETA: 0:01:36 Epoch: 1 Iter: 1198 Class Loss: 0.00 Loss: 0.03



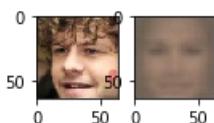
42%|# |ETA: 0:01:16 Epoch: 1 Iter: 1299 Class Loss: 0.00 Loss: 0.04



45%|# |ETA: 0:01:12 Epoch: 1 Iter: 1397 Class Loss: 0.06 Loss: 0.10



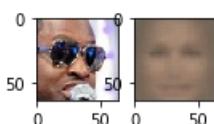
49%|# |ETA: 0:01:07 Epoch: 1 Iter: 1497 Class Loss: 0.00 Loss: 0.03



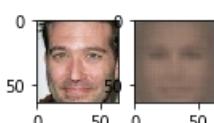
52%|## |ETA: 0:01:12 Epoch: 1 Iter: 1599 Class Loss: 0.00 Loss: 0.03



55%|## |ETA: 0:01:07 Epoch: 1 Iter: 1697 Class Loss: 0.06 Loss: 0.10



58%|## |ETA: 0:00:54 Epoch: 1 Iter: 1797 Class Loss: 0.02 Loss: 0.05



62%|## |ETA: 0:00:50 Epoch: 1 Iter: 1898 Class Loss: 0.00 Loss: 0.03



65%|## |ETA: 0:00:46 Epoch: 1 Iter: 1997 Class Loss: 0.00 Loss: 0.04



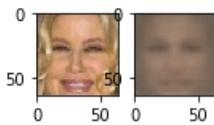
68%|## |ETA: 0:00:41 Epoch: 1 Iter: 2097 Class Loss: 0.00 Loss: 0.03



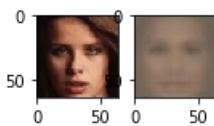
71%|## |ETA: 0:00:39 Epoch: 1 Iter: 2197 Class Loss: 0.00 Loss: 0.03



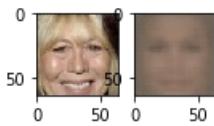
75%|### |ETA: 0:00:33 Epoch: 1 Iter: 2297 Class Loss: 0.00 Loss: 0.03



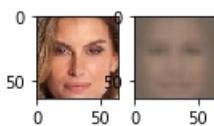
78%|###| ETA: 0:00:28 Epoch: 1 Iter: 2399 Class Loss: 0.02 Loss: 0.05



81%|###| ETA: 0:00:24 Epoch: 1 Iter: 2499 Class Loss: 0.00 Loss: 0.03



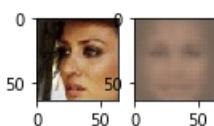
85%|###| ETA: 0:00:19 Epoch: 1 Iter: 2598 Class Loss: 0.00 Loss: 0.04



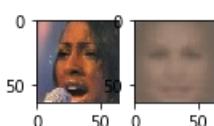
88%|###| ETA: 0:00:15 Epoch: 1 Iter: 2698 Class Loss: 0.02 Loss: 0.05



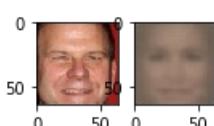
91%|###| ETA: 0:00:11 Epoch: 1 Iter: 2797 Class Loss: 0.16 Loss: 0.19



94%|###| ETA: 0:00:06 Epoch: 1 Iter: 2897 Class Loss: 0.00 Loss: 0.03



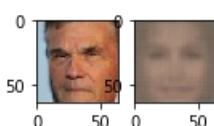
98%|###| ETA: 0:00:02 Epoch: 1 Iter: 2999 Class Loss: 0.00 Loss: 0.03



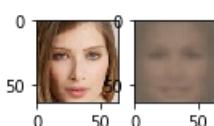
100%|####| Time: 0:02:20 Epoch: 1 Iter: 3052 Class Loss: 0.01 Loss: 0.04

Recomputing the sampling probabilities

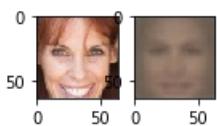
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



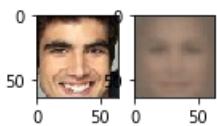
3%| | ETA: 0:02:08 Epoch: 2 Iter: 98 Class Loss: 0.00 Loss: 0.03



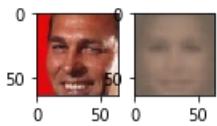
6%| | ETA: 0:02:04 Epoch: 2 Iter: 199 Class Loss: 0.00 Loss: 0.03



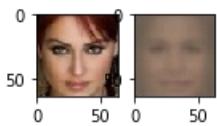
9% | ETA: 0:02:16 Epoch: 2 Iter: 299 Class Loss: 0.00 Loss: 0.04



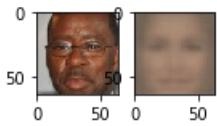
13% | ETA: 0:02:12 Epoch: 2 Iter: 397 Class Loss: 0.00 Loss: 0.03



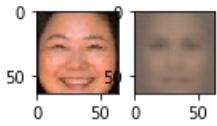
16% | ETA: 0:01:52 Epoch: 2 Iter: 499 Class Loss: 0.00 Loss: 0.03



19% | ETA: 0:01:47 Epoch: 2 Iter: 597 Class Loss: 0.00 Loss: 0.03



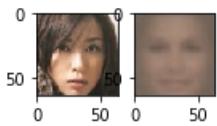
22% | # ETA: 0:01:43 Epoch: 2 Iter: 698 Class Loss: 0.00 Loss: 0.04



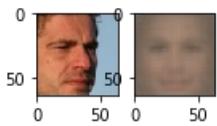
26% | # ETA: 0:01:38 Epoch: 2 Iter: 799 Class Loss: 0.00 Loss: 0.03



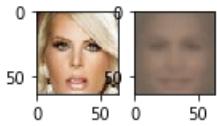
29% | # ETA: 0:01:34 Epoch: 2 Iter: 897 Class Loss: 0.00 Loss: 0.03



32% | # ETA: 0:01:30 Epoch: 2 Iter: 999 Class Loss: 0.00 Loss: 0.03



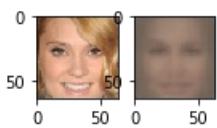
35% | # ETA: 0:01:25 Epoch: 2 Iter: 1097 Class Loss: 0.02 Loss: 0.05



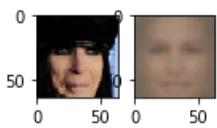
39% | # ETA: 0:01:21 Epoch: 2 Iter: 1198 Class Loss: 0.00 Loss: 0.03



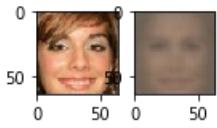
42% | # ETA: 0:01:17 Epoch: 2 Iter: 1297 Class Loss: 0.00 Loss: 0.03



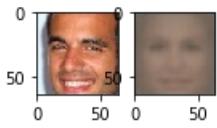
45%|# |ETA: 0:01:12 Epoch: 2 Iter: 1397 Class Loss: 0.00 Loss: 0.03



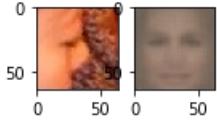
49%|# |ETA: 0:01:08 Epoch: 2 Iter: 1497 Class Loss: 0.00 Loss: 0.04



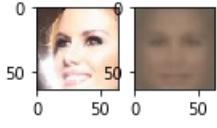
52%|## |ETA: 0:01:03 Epoch: 2 Iter: 1599 Class Loss: 0.00 Loss: 0.03



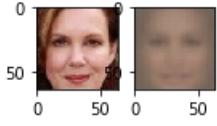
55%|## |ETA: 0:00:59 Epoch: 2 Iter: 1697 Class Loss: 0.00 Loss: 0.03



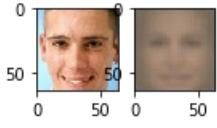
58%|## |ETA: 0:00:55 Epoch: 2 Iter: 1798 Class Loss: 0.00 Loss: 0.04



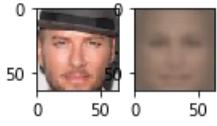
62%|## |ETA: 0:00:53 Epoch: 2 Iter: 1899 Class Loss: 0.12 Loss: 0.16



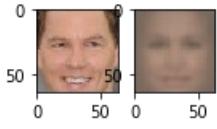
65%|## |ETA: 0:00:55 Epoch: 2 Iter: 1998 Class Loss: 0.00 Loss: 0.03



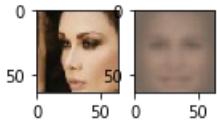
68%|## |ETA: 0:00:48 Epoch: 2 Iter: 2098 Class Loss: 0.00 Loss: 0.03



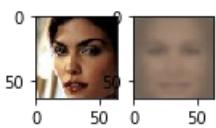
71%|## |ETA: 0:00:37 Epoch: 2 Iter: 2197 Class Loss: 0.00 Loss: 0.03



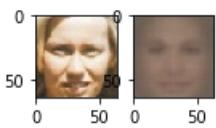
75%|### |ETA: 0:00:32 Epoch: 2 Iter: 2298 Class Loss: 0.00 Loss: 0.04



78%|### |ETA: 0:00:28 Epoch: 2 Iter: 2399 Class Loss: 0.00 Loss: 0.04



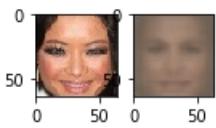
81%|###| ETA: 0:00:24 Epoch: 2 Iter: 2497 Class Loss: 0.00 Loss: 0.03



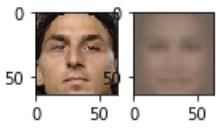
85%|###| ETA: 0:00:19 Epoch: 2 Iter: 2598 Class Loss: 0.00 Loss: 0.03



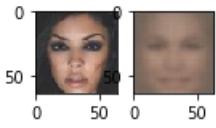
88%|###| ETA: 0:00:15 Epoch: 2 Iter: 2698 Class Loss: 0.00 Loss: 0.04



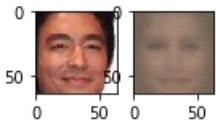
91%|###| ETA: 0:00:11 Epoch: 2 Iter: 2797 Class Loss: 0.02 Loss: 0.05



94%|###| ETA: 0:00:06 Epoch: 2 Iter: 2897 Class Loss: 0.00 Loss: 0.04



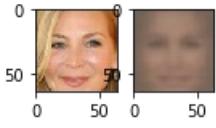
98%|###| ETA: 0:00:02 Epoch: 2 Iter: 2998 Class Loss: 0.01 Loss: 0.04



100%|####| Time: 0:02:21 Epoch: 2 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

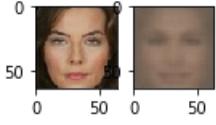
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



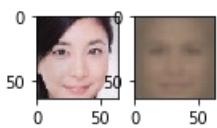
3%| | ETA: 0:02:10 Epoch: 3 Iter: 98 Class Loss: 0.00 Loss: 0.03



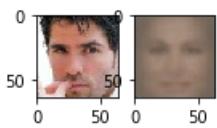
6%| | ETA: 0:02:04 Epoch: 3 Iter: 199 Class Loss: 0.00 Loss: 0.03



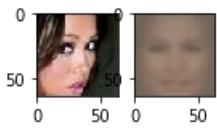
9%| | ETA: 0:02:01 Epoch: 3 Iter: 298 Class Loss: 0.00 Loss: 0.03



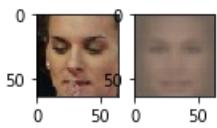
13% | ETA: 0:01:56 Epoch: 3 Iter: 398 Class Loss: 0.00 Loss: 0.04



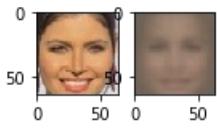
16% | ETA: 0:01:54 Epoch: 3 Iter: 497 Class Loss: 0.02 Loss: 0.05



19% | ETA: 0:01:56 Epoch: 3 Iter: 599 Class Loss: 0.00 Loss: 0.03



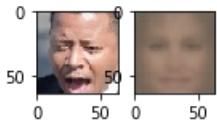
22% | # ETA: 0:01:57 Epoch: 3 Iter: 698 Class Loss: 0.00 Loss: 0.03



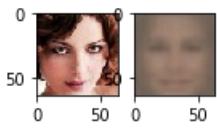
26% | # ETA: 0:01:42 Epoch: 3 Iter: 799 Class Loss: 0.00 Loss: 0.03



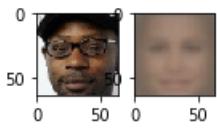
29% | # ETA: 0:01:34 Epoch: 3 Iter: 897 Class Loss: 0.00 Loss: 0.03



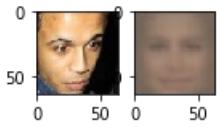
32% | # ETA: 0:01:30 Epoch: 3 Iter: 998 Class Loss: 0.00 Loss: 0.04



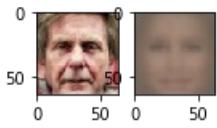
36% | # ETA: 0:01:25 Epoch: 3 Iter: 1099 Class Loss: 0.02 Loss: 0.05



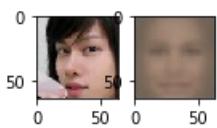
39% | # ETA: 0:01:21 Epoch: 3 Iter: 1197 Class Loss: 0.00 Loss: 0.03



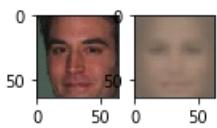
42% | # ETA: 0:01:16 Epoch: 3 Iter: 1299 Class Loss: 0.00 Loss: 0.03



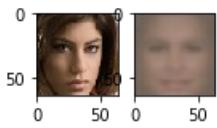
45% | # ETA: 0:01:12 Epoch: 3 Iter: 1397 Class Loss: 0.02 Loss: 0.06



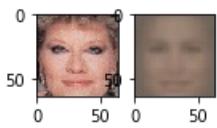
49%|#| ETA: 0:01:08 Epoch: 3 Iter: 1497 Class Loss: 0.00 Loss: 0.04



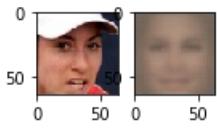
52%|##| ETA: 0:01:15 Epoch: 3 Iter: 1599 Class Loss: 0.00 Loss: 0.03



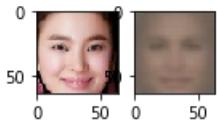
55%|##| ETA: 0:00:59 Epoch: 3 Iter: 1697 Class Loss: 0.00 Loss: 0.03



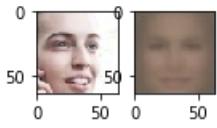
58%|##| ETA: 0:00:55 Epoch: 3 Iter: 1798 Class Loss: 0.00 Loss: 0.03



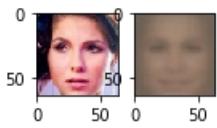
62%|##| ETA: 0:00:50 Epoch: 3 Iter: 1899 Class Loss: 0.00 Loss: 0.03



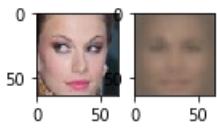
65%|##| ETA: 0:00:46 Epoch: 3 Iter: 1999 Class Loss: 0.00 Loss: 0.03



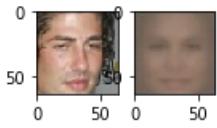
68%|##| ETA: 0:00:41 Epoch: 3 Iter: 2098 Class Loss: 0.13 Loss: 0.16



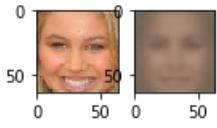
71%|##| ETA: 0:00:37 Epoch: 3 Iter: 2197 Class Loss: 0.00 Loss: 0.03



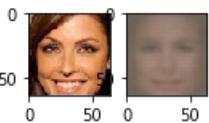
75%|###| ETA: 0:00:33 Epoch: 3 Iter: 2297 Class Loss: 0.00 Loss: 0.03



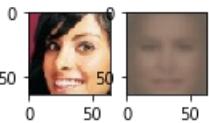
78%|###| ETA: 0:00:32 Epoch: 3 Iter: 2398 Class Loss: 0.00 Loss: 0.03



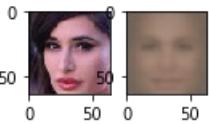
81%|###| ETA: 0:00:27 Epoch: 3 Iter: 2499 Class Loss: 0.00 Loss: 0.04



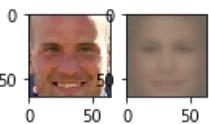
85%|###| ETA: 0:00:19 Epoch: 3 Iter: 2598 Class Loss: 0.00 Loss: 0.04



88%|###| ETA: 0:00:15 Epoch: 3 Iter: 2697 Class Loss: 0.00 Loss: 0.03



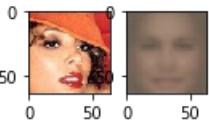
91%|###| ETA: 0:00:11 Epoch: 3 Iter: 2797 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 3 Iter: 2898 Class Loss: 0.00 Loss: 0.03



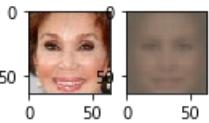
98%|###| ETA: 0:00:02 Epoch: 3 Iter: 2999 Class Loss: 0.00 Loss: 0.04



100%|####| Time: 0:02:21 Epoch: 3 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

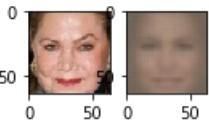
N/A%| | ETA: ---- Epoch: nan Iter: nan Class Loss: nan Loss: nan



3%| | ETA: 0:02:10 Epoch: 4 Iter: 97 Class Loss: 0.16 Loss: 0.20



6%| | ETA: 0:02:05 Epoch: 4 Iter: 199 Class Loss: 0.00 Loss: 0.03



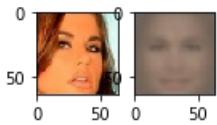
9%| | ETA: 0:02:00 Epoch: 4 Iter: 299 Class Loss: 0.02 Loss: 0.05



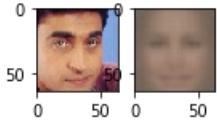
13%| | ETA: 0:01:56 Epoch: 4 Iter: 399 Class Loss: 0.00 Loss: 0.03



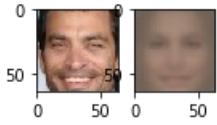
16% | ETA: 0:01:52 Epoch: 4 Iter: 499 Class Loss: 0.00 Loss: 0.03



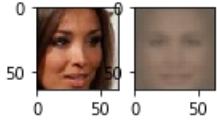
19% | ETA: 0:01:47 Epoch: 4 Iter: 597 Class Loss: 0.00 Loss: 0.04



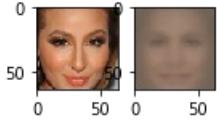
22% |# ETA: 0:01:43 Epoch: 4 Iter: 698 Class Loss: 0.00 Loss: 0.03



26% |# ETA: 0:01:40 Epoch: 4 Iter: 798 Class Loss: 0.00 Loss: 0.03



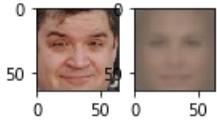
29% |# ETA: 0:01:34 Epoch: 4 Iter: 897 Class Loss: 0.00 Loss: 0.03



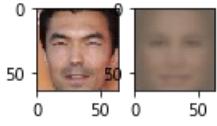
32% |# ETA: 0:01:43 Epoch: 4 Iter: 997 Class Loss: 0.00 Loss: 0.03



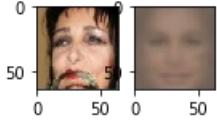
36% |# ETA: 0:01:38 Epoch: 4 Iter: 1099 Class Loss: 0.00 Loss: 0.04



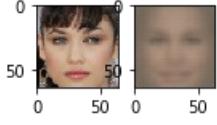
39% |# ETA: 0:01:21 Epoch: 4 Iter: 1198 Class Loss: 0.00 Loss: 0.03



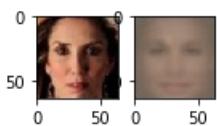
42% |# ETA: 0:01:17 Epoch: 4 Iter: 1298 Class Loss: 0.00 Loss: 0.03



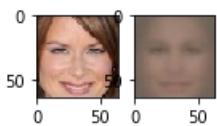
45% |# ETA: 0:01:12 Epoch: 4 Iter: 1397 Class Loss: 0.00 Loss: 0.03



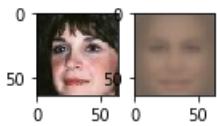
49% |# ETA: 0:01:08 Epoch: 4 Iter: 1498 Class Loss: 0.00 Loss: 0.03



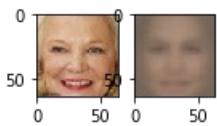
52%|##| ETA: 0:01:03 Epoch: 4 Iter: 1598 Class Loss: 0.00 Loss: 0.03



55%|##| ETA: 0:00:58 Epoch: 4 Iter: 1699 Class Loss: 0.00 Loss: 0.04



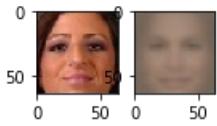
58%|##| ETA: 0:00:54 Epoch: 4 Iter: 1799 Class Loss: 0.00 Loss: 0.04



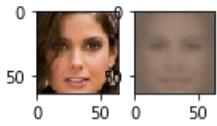
62%|##| ETA: 0:00:50 Epoch: 4 Iter: 1899 Class Loss: 0.00 Loss: 0.04



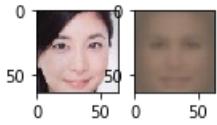
65%|##| ETA: 0:00:46 Epoch: 4 Iter: 1997 Class Loss: 0.00 Loss: 0.03



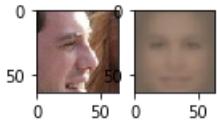
68%|##| ETA: 0:00:41 Epoch: 4 Iter: 2098 Class Loss: 0.00 Loss: 0.03



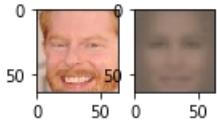
71%|##| ETA: 0:00:37 Epoch: 4 Iter: 2197 Class Loss: 0.00 Loss: 0.03



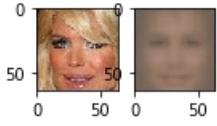
75%|###| ETA: 0:00:32 Epoch: 4 Iter: 2299 Class Loss: 0.00 Loss: 0.03



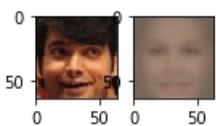
78%|###| ETA: 0:00:31 Epoch: 4 Iter: 2398 Class Loss: 0.00 Loss: 0.03



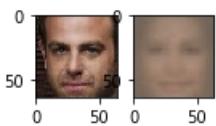
81%|###| ETA: 0:00:29 Epoch: 4 Iter: 2499 Class Loss: 0.00 Loss: 0.03



85%|###| ETA: 0:00:21 Epoch: 4 Iter: 2597 Class Loss: 0.00 Loss: 0.03



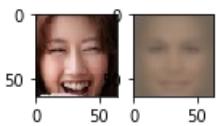
88%|###|ETA: 0:00:17 Epoch: 4 Iter: 2698 Class Loss: 0.00 Loss: 0.04



91%|###|ETA: 0:00:12 Epoch: 4 Iter: 2798 Class Loss: 0.00 Loss: 0.03



94%|###|ETA: 0:00:06 Epoch: 4 Iter: 2898 Class Loss: 0.00 Loss: 0.04



98%|###|ETA: 0:00:02 Epoch: 4 Iter: 2999 Class Loss: 0.00 Loss: 0.03



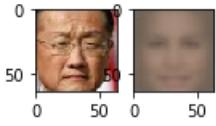
100%|####|Time: 0:02:22 Epoch: 4 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

N/A%| |ETA: --:--- Epoch: nan Iter: nan Class Loss: nan Loss: nan



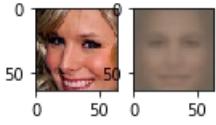
3%| |ETA: 0:02:10 Epoch: 5 Iter: 97 Class Loss: 0.00 Loss: 0.03



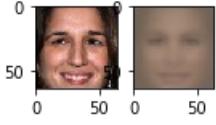
6%| |ETA: 0:02:04 Epoch: 5 Iter: 199 Class Loss: 0.00 Loss: 0.03



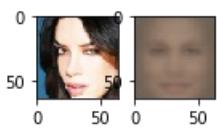
9%| |ETA: 0:02:01 Epoch: 5 Iter: 298 Class Loss: 0.00 Loss: 0.03



13%| |ETA: 0:01:57 Epoch: 5 Iter: 398 Class Loss: 0.00 Loss: 0.03



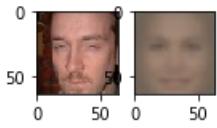
16%| |ETA: 0:01:52 Epoch: 5 Iter: 499 Class Loss: 0.00 Loss: 0.04



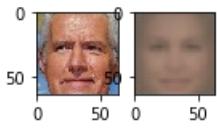
19% | # ETA: 0:01:49 Epoch: 5 Iter: 599 Class Loss: 0.00 Loss: 0.03



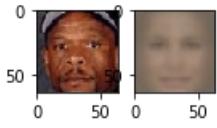
22% | # ETA: 0:01:43 Epoch: 5 Iter: 697 Class Loss: 0.00 Loss: 0.03



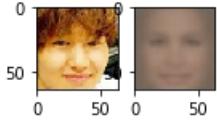
26% | # ETA: 0:01:38 Epoch: 5 Iter: 799 Class Loss: 0.00 Loss: 0.03



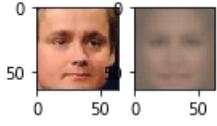
29% | # ETA: 0:01:35 Epoch: 5 Iter: 897 Class Loss: 0.00 Loss: 0.03



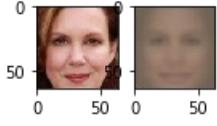
32% | # ETA: 0:01:30 Epoch: 5 Iter: 997 Class Loss: 0.00 Loss: 0.03



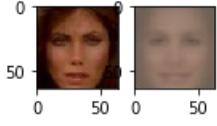
36% | # ETA: 0:01:25 Epoch: 5 Iter: 1099 Class Loss: 0.00 Loss: 0.04



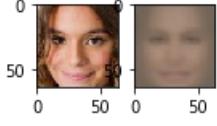
39% | # ETA: 0:01:20 Epoch: 5 Iter: 1198 Class Loss: 0.00 Loss: 0.03



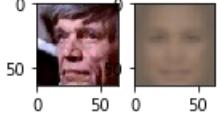
42% | # ETA: 0:01:16 Epoch: 5 Iter: 1298 Class Loss: 0.00 Loss: 0.03



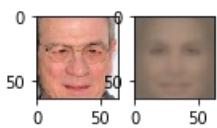
45% | # ETA: 0:01:22 Epoch: 5 Iter: 1399 Class Loss: 0.00 Loss: 0.03



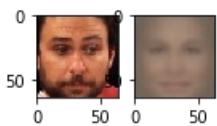
49% | # ETA: 0:01:17 Epoch: 5 Iter: 1499 Class Loss: 0.00 Loss: 0.03



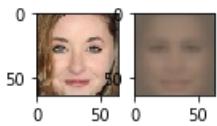
52% | ## ETA: 0:01:04 Epoch: 5 Iter: 1599 Class Loss: 0.00 Loss: 0.03



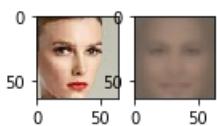
55%|##| ETA: 0:00:59 Epoch: 5 Iter: 1697 Class Loss: 0.00 Loss: 0.03



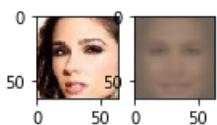
58%|##| ETA: 0:00:54 Epoch: 5 Iter: 1798 Class Loss: 0.02 Loss: 0.06



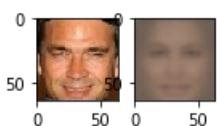
62%|##| ETA: 0:00:51 Epoch: 5 Iter: 1898 Class Loss: 0.00 Loss: 0.04



65%|##| ETA: 0:00:54 Epoch: 5 Iter: 1998 Class Loss: 0.00 Loss: 0.03



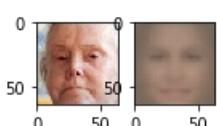
68%|##| ETA: 0:00:42 Epoch: 5 Iter: 2098 Class Loss: 0.00 Loss: 0.03



71%|##| ETA: 0:00:37 Epoch: 5 Iter: 2197 Class Loss: 0.00 Loss: 0.03



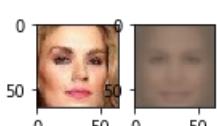
75%|###| ETA: 0:00:33 Epoch: 5 Iter: 2297 Class Loss: 0.00 Loss: 0.03



78%|###| ETA: 0:00:28 Epoch: 5 Iter: 2399 Class Loss: 0.00 Loss: 0.03



81%|###| ETA: 0:00:24 Epoch: 5 Iter: 2497 Class Loss: 0.00 Loss: 0.03



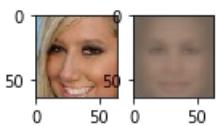
85%|###| ETA: 0:00:19 Epoch: 5 Iter: 2598 Class Loss: 0.00 Loss: 0.03



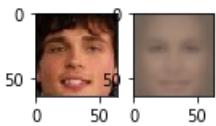
88%|###| ETA: 0:00:15 Epoch: 5 Iter: 2698 Class Loss: 0.00 Loss: 0.03



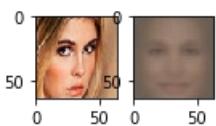
91%|###| ETA: 0:00:11 Epoch: 5 Iter: 2799 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 5 Iter: 2898 Class Loss: 0.01 Loss: 0.05



98%|###| ETA: 0:00:02 Epoch: 5 Iter: 2999 Class Loss: 0.00 Loss: 0.04



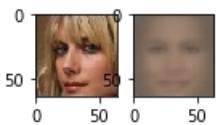
100%|####| Time: 0:02:20 Epoch: 5 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

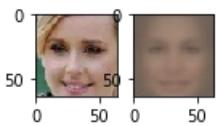
N/A%| | ETA: --:--- Epoch: nan Iter: nan Class Loss: nan Loss: nan



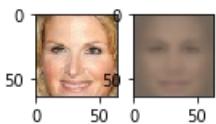
3%| | ETA: 0:02:29 Epoch: 6 Iter: 98 Class Loss: 0.00 Loss: 0.04



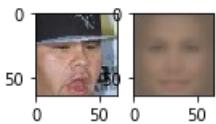
6%| | ETA: 0:02:06 Epoch: 6 Iter: 199 Class Loss: 0.00 Loss: 0.04



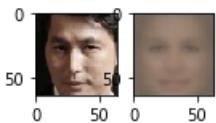
9%| | ETA: 0:02:01 Epoch: 6 Iter: 297 Class Loss: 0.00 Loss: 0.03



13%| | ETA: 0:01:56 Epoch: 6 Iter: 398 Class Loss: 0.00 Loss: 0.03



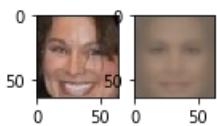
16%| | ETA: 0:01:53 Epoch: 6 Iter: 497 Class Loss: 0.00 Loss: 0.03



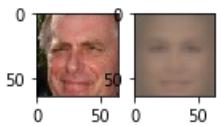
19%| | ETA: 0:01:47 Epoch: 6 Iter: 599 Class Loss: 0.00 Loss: 0.04



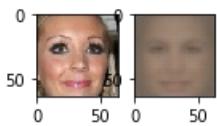
22%|# |ETA: 0:01:43 Epoch: 6 Iter: 698 Class Loss: 0.00 Loss: 0.03



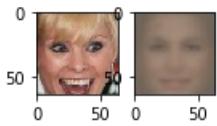
26%|# |ETA: 0:01:38 Epoch: 6 Iter: 798 Class Loss: 0.00 Loss: 0.03



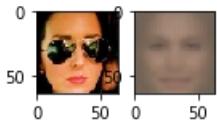
29%|# |ETA: 0:01:34 Epoch: 6 Iter: 897 Class Loss: 0.00 Loss: 0.03



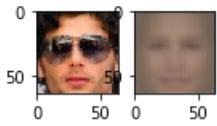
32%|# |ETA: 0:01:29 Epoch: 6 Iter: 998 Class Loss: 0.00 Loss: 0.04



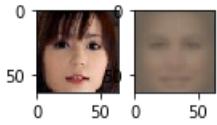
35%|# |ETA: 0:01:25 Epoch: 6 Iter: 1097 Class Loss: 0.00 Loss: 0.03



39%|# |ETA: 0:01:21 Epoch: 6 Iter: 1198 Class Loss: 0.00 Loss: 0.04



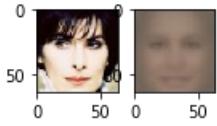
42%|# |ETA: 0:01:16 Epoch: 6 Iter: 1298 Class Loss: 0.00 Loss: 0.03



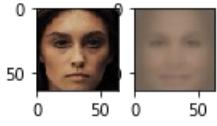
45%|# |ETA: 0:01:12 Epoch: 6 Iter: 1399 Class Loss: 0.00 Loss: 0.03



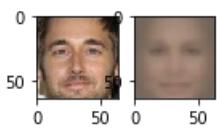
49%|# |ETA: 0:01:07 Epoch: 6 Iter: 1497 Class Loss: 0.00 Loss: 0.03



52%|## |ETA: 0:01:03 Epoch: 6 Iter: 1599 Class Loss: 0.00 Loss: 0.03



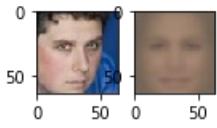
55%|## |ETA: 0:01:04 Epoch: 6 Iter: 1699 Class Loss: 0.00 Loss: 0.03



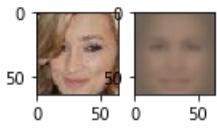
58%|##| ETA: 0:01:03 Epoch: 6 Iter: 1798 Class Loss: 0.00 Loss: 0.04



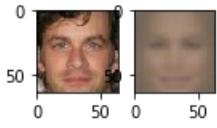
62%|##| ETA: 0:00:52 Epoch: 6 Iter: 1898 Class Loss: 0.00 Loss: 0.03



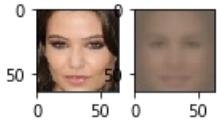
65%|##| ETA: 0:00:46 Epoch: 6 Iter: 1999 Class Loss: 0.00 Loss: 0.03



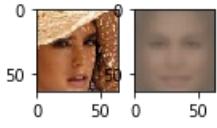
68%|##| ETA: 0:00:41 Epoch: 6 Iter: 2098 Class Loss: 0.00 Loss: 0.03



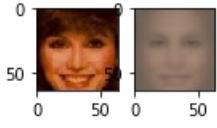
72%|##| ETA: 0:00:37 Epoch: 6 Iter: 2199 Class Loss: 0.00 Loss: 0.04



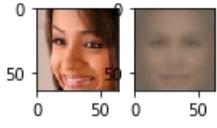
75%|###| ETA: 0:00:33 Epoch: 6 Iter: 2298 Class Loss: 0.00 Loss: 0.03



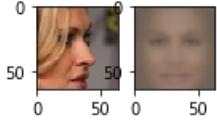
78%|###| ETA: 0:00:28 Epoch: 6 Iter: 2399 Class Loss: 0.00 Loss: 0.03



81%|###| ETA: 0:00:24 Epoch: 6 Iter: 2497 Class Loss: 0.00 Loss: 0.03



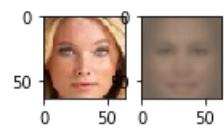
85%|###| ETA: 0:00:19 Epoch: 6 Iter: 2598 Class Loss: 0.00 Loss: 0.03



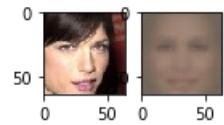
88%|###| ETA: 0:00:15 Epoch: 6 Iter: 2698 Class Loss: 0.00 Loss: 0.03



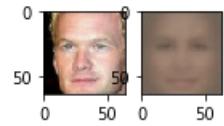
91%|###| ETA: 0:00:11 Epoch: 6 Iter: 2799 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:07 Epoch: 6 Iter: 2898 Class Loss: 0.00 Loss: 0.03



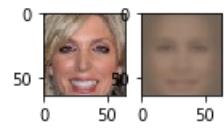
98%|###| ETA: 0:00:02 Epoch: 6 Iter: 2999 Class Loss: 0.11 Loss: 0.15



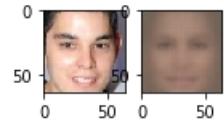
100%|####| Time: 0:02:22 Epoch: 6 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

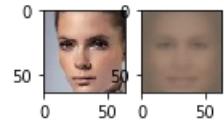
N/A%| | ETA: --:--- Epoch: nan Iter: nan Class Loss: nan Loss: nan



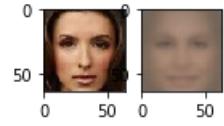
3%| | ETA: 0:02:08 Epoch: 7 Iter: 99 Class Loss: 0.00 Loss: 0.03



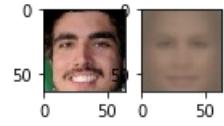
6%| | ETA: 0:02:04 Epoch: 7 Iter: 198 Class Loss: 0.00 Loss: 0.04



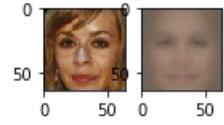
9%| | ETA: 0:02:02 Epoch: 7 Iter: 299 Class Loss: 0.00 Loss: 0.03



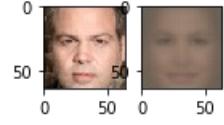
13%| | ETA: 0:02:11 Epoch: 7 Iter: 398 Class Loss: 0.00 Loss: 0.03



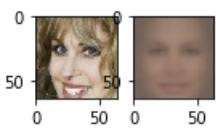
16%| | ETA: 0:02:02 Epoch: 7 Iter: 497 Class Loss: 0.00 Loss: 0.03



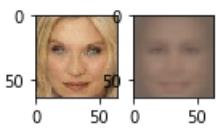
19%| | ETA: 0:01:47 Epoch: 7 Iter: 597 Class Loss: 0.00 Loss: 0.03



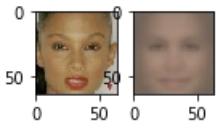
22%|#| ETA: 0:01:43 Epoch: 7 Iter: 697 Class Loss: 0.00 Loss: 0.03



26%|# |ETA: 0:01:38 Epoch: 7 Iter: 798 Class Loss: 0.00 Loss: 0.04



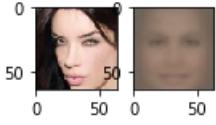
29%|# |ETA: 0:01:33 Epoch: 7 Iter: 897 Class Loss: 0.00 Loss: 0.03



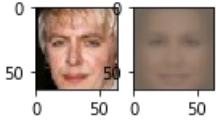
32%|# |ETA: 0:01:28 Epoch: 7 Iter: 998 Class Loss: 0.00 Loss: 0.03



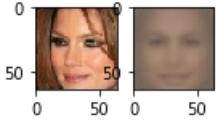
35%|# |ETA: 0:01:24 Epoch: 7 Iter: 1097 Class Loss: 0.00 Loss: 0.03



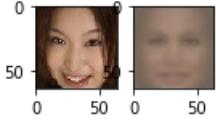
39%|# |ETA: 0:01:21 Epoch: 7 Iter: 1198 Class Loss: 0.00 Loss: 0.03



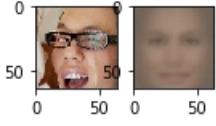
42%|# |ETA: 0:01:16 Epoch: 7 Iter: 1299 Class Loss: 0.00 Loss: 0.03



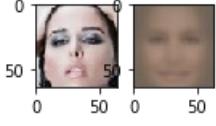
45%|# |ETA: 0:01:12 Epoch: 7 Iter: 1399 Class Loss: 0.00 Loss: 0.03



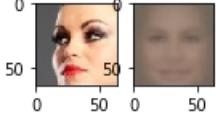
49%|# |ETA: 0:01:08 Epoch: 7 Iter: 1498 Class Loss: 0.00 Loss: 0.03



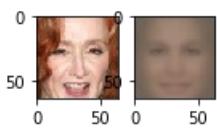
52%|## |ETA: 0:01:03 Epoch: 7 Iter: 1599 Class Loss: 0.00 Loss: 0.03



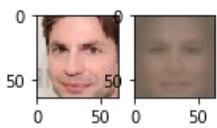
55%|## |ETA: 0:00:59 Epoch: 7 Iter: 1697 Class Loss: 0.00 Loss: 0.04



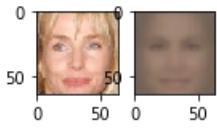
58%|## |ETA: 0:00:54 Epoch: 7 Iter: 1798 Class Loss: 0.00 Loss: 0.04



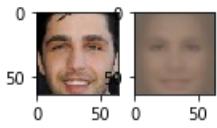
62%|##| ETA: 0:00:50 Epoch: 7 Iter: 1899 Class Loss: 0.00 Loss: 0.03



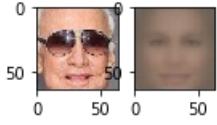
65%|##| ETA: 0:00:46 Epoch: 7 Iter: 1997 Class Loss: 0.00 Loss: 0.04



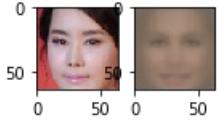
68%|##| ETA: 0:00:47 Epoch: 7 Iter: 2097 Class Loss: 0.00 Loss: 0.04



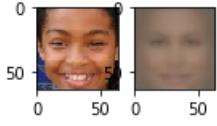
72%|##| ETA: 0:00:42 Epoch: 7 Iter: 2199 Class Loss: 0.00 Loss: 0.03



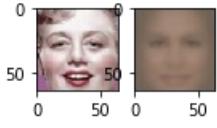
75%|###| ETA: 0:00:34 Epoch: 7 Iter: 2299 Class Loss: 0.00 Loss: 0.03



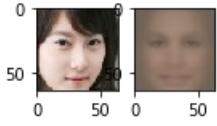
78%|###| ETA: 0:00:31 Epoch: 7 Iter: 2399 Class Loss: 0.00 Loss: 0.03



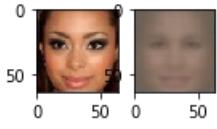
81%|###| ETA: 0:00:24 Epoch: 7 Iter: 2499 Class Loss: 0.00 Loss: 0.03



85%|###| ETA: 0:00:19 Epoch: 7 Iter: 2598 Class Loss: 0.00 Loss: 0.03



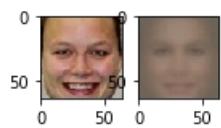
88%|###| ETA: 0:00:15 Epoch: 7 Iter: 2699 Class Loss: 0.00 Loss: 0.03



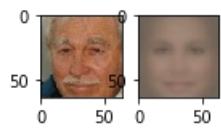
91%|###| ETA: 0:00:11 Epoch: 7 Iter: 2799 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 7 Iter: 2898 Class Loss: 0.00 Loss: 0.03



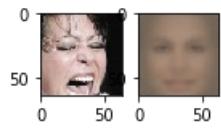
98%|###| ETA: 0:00:02 Epoch: 7 Iter: 2998 Class Loss: 0.00 Loss: 0.03



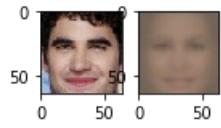
100%|####| Time: 0:02:20 Epoch: 7 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

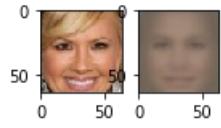
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



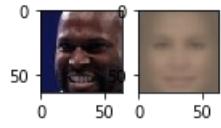
3%| | ETA: 0:02:08 Epoch: 8 Iter: 99 Class Loss: 0.00 Loss: 0.03



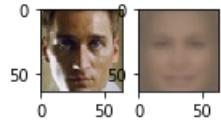
6%| | ETA: 0:02:05 Epoch: 8 Iter: 198 Class Loss: 0.00 Loss: 0.04



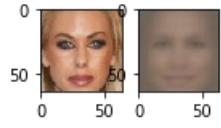
9%| | ETA: 0:01:59 Epoch: 8 Iter: 299 Class Loss: 0.00 Loss: 0.03



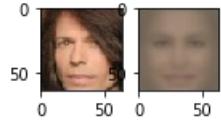
13%| | ETA: 0:01:55 Epoch: 8 Iter: 397 Class Loss: 0.00 Loss: 0.03



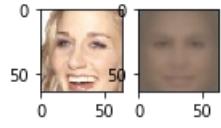
16%| | ETA: 0:01:51 Epoch: 8 Iter: 499 Class Loss: 0.00 Loss: 0.04



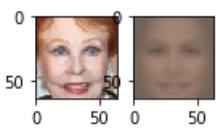
19%| | ETA: 0:01:46 Epoch: 8 Iter: 599 Class Loss: 0.00 Loss: 0.03



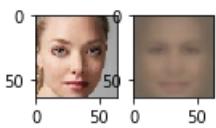
22%|#| ETA: 0:01:47 Epoch: 8 Iter: 698 Class Loss: 0.00 Loss: 0.03



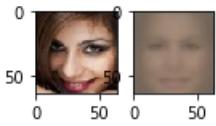
26%|#| ETA: 0:01:51 Epoch: 8 Iter: 799 Class Loss: 0.11 Loss: 0.15



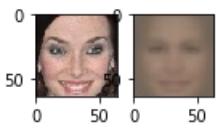
29%|# |ETA: 0:01:38 Epoch: 8 Iter: 898 Class Loss: 0.00 Loss: 0.03



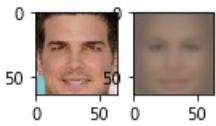
32%|# |ETA: 0:01:29 Epoch: 8 Iter: 999 Class Loss: 0.00 Loss: 0.03



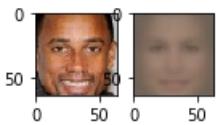
35%|# |ETA: 0:01:24 Epoch: 8 Iter: 1097 Class Loss: 0.00 Loss: 0.04



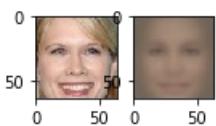
39%|# |ETA: 0:01:21 Epoch: 8 Iter: 1197 Class Loss: 0.00 Loss: 0.03



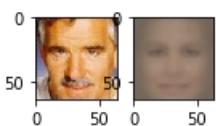
42%|# |ETA: 0:01:16 Epoch: 8 Iter: 1298 Class Loss: 0.01 Loss: 0.04



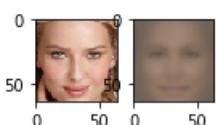
45%|# |ETA: 0:01:12 Epoch: 8 Iter: 1397 Class Loss: 0.00 Loss: 0.03



49%|# |ETA: 0:01:08 Epoch: 8 Iter: 1497 Class Loss: 0.00 Loss: 0.03



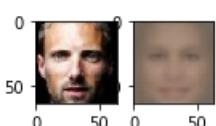
52%## |ETA: 0:01:03 Epoch: 8 Iter: 1598 Class Loss: 0.00 Loss: 0.03



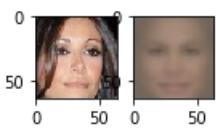
55%## |ETA: 0:00:58 Epoch: 8 Iter: 1697 Class Loss: 0.00 Loss: 0.03



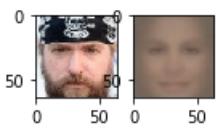
58%## |ETA: 0:00:54 Epoch: 8 Iter: 1798 Class Loss: 0.00 Loss: 0.03



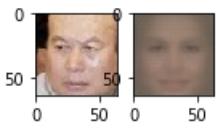
62%## |ETA: 0:00:49 Epoch: 8 Iter: 1898 Class Loss: 0.00 Loss: 0.03



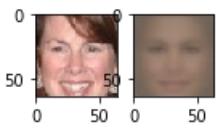
65%|##| ETA: 0:00:45 Epoch: 8 Iter: 1997 Class Loss: 0.00 Loss: 0.03



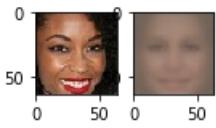
68%|##| ETA: 0:00:41 Epoch: 8 Iter: 2098 Class Loss: 0.00 Loss: 0.03



72%|##| ETA: 0:00:37 Epoch: 8 Iter: 2199 Class Loss: 0.02 Loss: 0.04



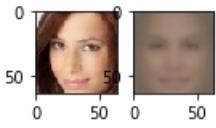
75%|###| ETA: 0:00:32 Epoch: 8 Iter: 2297 Class Loss: 0.00 Loss: 0.04



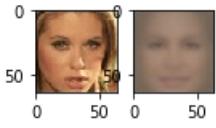
78%|###| ETA: 0:00:28 Epoch: 8 Iter: 2398 Class Loss: 0.00 Loss: 0.03



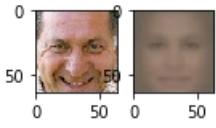
81%|###| ETA: 0:00:27 Epoch: 8 Iter: 2498 Class Loss: 0.00 Loss: 0.03



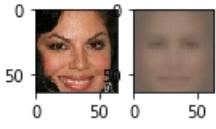
85%|###| ETA: 0:00:22 Epoch: 8 Iter: 2599 Class Loss: 0.00 Loss: 0.03



88%|###| ETA: 0:00:15 Epoch: 8 Iter: 2698 Class Loss: 0.00 Loss: 0.03



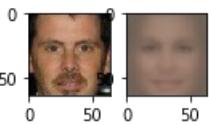
91%|###| ETA: 0:00:11 Epoch: 8 Iter: 2799 Class Loss: 0.00 Loss: 0.04



94%|###| ETA: 0:00:06 Epoch: 8 Iter: 2898 Class Loss: 0.00 Loss: 0.03



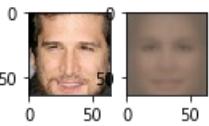
98%|###| ETA: 0:00:02 Epoch: 8 Iter: 2998 Class Loss: 0.00 Loss: 0.04



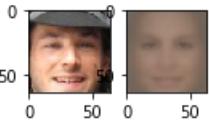
100%|####|Time: 0:02:20 Epoch: 8 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

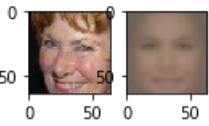
N/A%| |ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



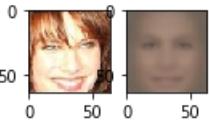
3%| |ETA: 0:02:09 Epoch: 9 Iter: 97 Class Loss: 0.02 Loss: 0.05



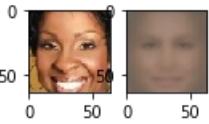
6%| |ETA: 0:02:04 Epoch: 9 Iter: 199 Class Loss: 0.00 Loss: 0.03



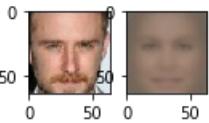
9%| |ETA: 0:02:09 Epoch: 9 Iter: 298 Class Loss: 0.00 Loss: 0.03



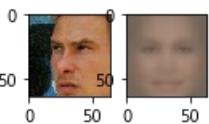
13%| |ETA: 0:02:20 Epoch: 9 Iter: 398 Class Loss: 0.00 Loss: 0.02



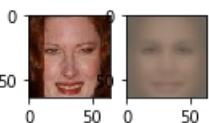
16%| |ETA: 0:01:59 Epoch: 9 Iter: 498 Class Loss: 0.00 Loss: 0.04



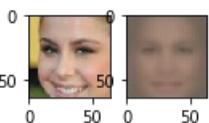
19%| |ETA: 0:01:46 Epoch: 9 Iter: 599 Class Loss: 0.00 Loss: 0.04



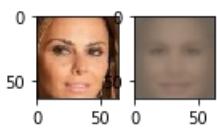
22%|# |ETA: 0:01:42 Epoch: 9 Iter: 698 Class Loss: 0.00 Loss: 0.03



26%|# |ETA: 0:01:38 Epoch: 9 Iter: 799 Class Loss: 0.00 Loss: 0.03



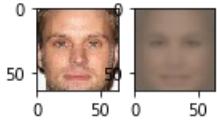
29%|# |ETA: 0:01:33 Epoch: 9 Iter: 897 Class Loss: 0.00 Loss: 0.03



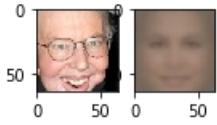
32%|# |ETA: 0:01:29 Epoch: 9 Iter: 997 Class Loss: 0.00 Loss: 0.03



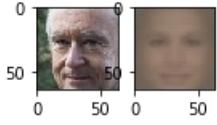
35%|# |ETA: 0:01:37 Epoch: 9 Iter: 1098 Class Loss: 0.00 Loss: 0.03



39%|# |ETA: 0:01:32 Epoch: 9 Iter: 1198 Class Loss: 0.00 Loss: 0.03



42%|# |ETA: 0:01:16 Epoch: 9 Iter: 1299 Class Loss: 0.00 Loss: 0.03



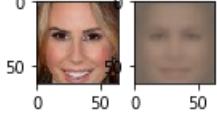
45%|# |ETA: 0:01:11 Epoch: 9 Iter: 1399 Class Loss: 0.00 Loss: 0.03



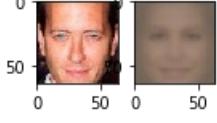
49%|# |ETA: 0:01:07 Epoch: 9 Iter: 1499 Class Loss: 0.00 Loss: 0.03



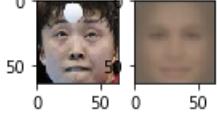
52%|## |ETA: 0:01:03 Epoch: 9 Iter: 1599 Class Loss: 0.00 Loss: 0.04



55%|## |ETA: 0:00:59 Epoch: 9 Iter: 1697 Class Loss: 0.00 Loss: 0.03



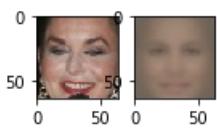
58%|## |ETA: 0:00:54 Epoch: 9 Iter: 1798 Class Loss: 0.00 Loss: 0.03



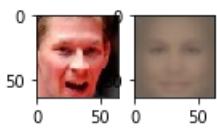
62%|## |ETA: 0:00:50 Epoch: 9 Iter: 1898 Class Loss: 0.00 Loss: 0.03



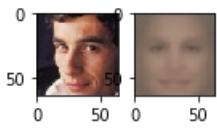
65%|## |ETA: 0:00:45 Epoch: 9 Iter: 1997 Class Loss: 0.00 Loss: 0.03



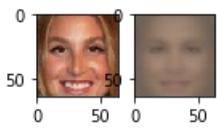
68%|##| ETA: 0:00:41 Epoch: 9 Iter: 2098 Class Loss: 0.00 Loss: 0.04



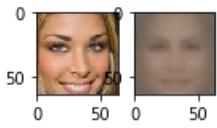
71%|##| ETA: 0:00:37 Epoch: 9 Iter: 2197 Class Loss: 0.00 Loss: 0.03



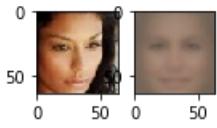
75%|###| ETA: 0:00:32 Epoch: 9 Iter: 2298 Class Loss: 0.00 Loss: 0.03



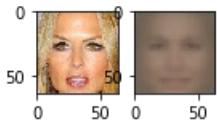
78%|###| ETA: 0:00:28 Epoch: 9 Iter: 2399 Class Loss: 0.00 Loss: 0.03



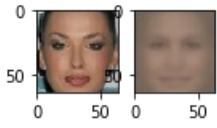
81%|###| ETA: 0:00:23 Epoch: 9 Iter: 2499 Class Loss: 0.00 Loss: 0.04



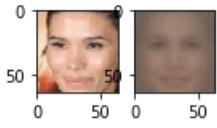
85%|###| ETA: 0:00:19 Epoch: 9 Iter: 2598 Class Loss: 0.00 Loss: 0.03



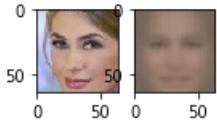
88%|###| ETA: 0:00:15 Epoch: 9 Iter: 2698 Class Loss: 0.00 Loss: 0.03



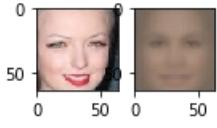
91%|###| ETA: 0:00:13 Epoch: 9 Iter: 2798 Class Loss: 0.00 Loss: 0.04



94%|###| ETA: 0:00:07 Epoch: 9 Iter: 2898 Class Loss: 0.00 Loss: 0.04



98%|###| ETA: 0:00:02 Epoch: 9 Iter: 2999 Class Loss: 0.00 Loss: 0.03



```
100% #####|Time: 0:02:21 Epoch: 9 Iter: 3052 Class Loss: 0.00 Loss: 0.03  
100% (97 of 97) |##### Elapsed Time: 0:01:06 Time: 0:01:06  
N/A% (0 of 72) | Elapsed Time: 0:00:00 ETA: -----
```

male lighter: 0.886597938144

```
100% (72 of 72) |##### Elapsed Time: 0:00:50 Time: 0:00:50  
N/A% (0 of 78) | Elapsed Time: 0:00:00 ETA: -----
```

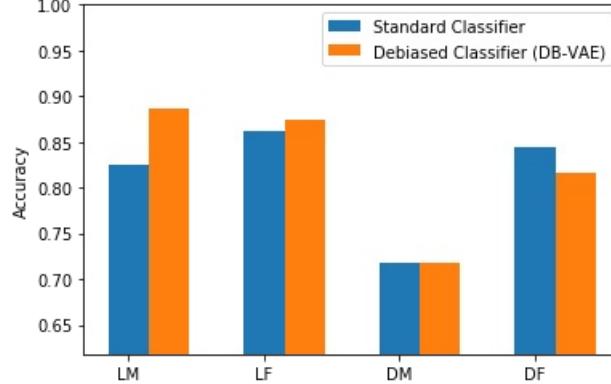
female lighter: 0.875

```
100% (78 of 78) |##### Elapsed Time: 0:00:52 Time: 0:00:52  
N/A% (0 of 71) | Elapsed Time: 0:00:00 ETA: -----
```

male darker: 0.717948717949

```
100% (71 of 71) |##### Elapsed Time: 0:00:47 Time: 0:00:47
```

female darker: 0.816901408451



Smooth Factor of 2

In [64]:

```
run_smooth_factor(2)
```

Recomputing the sampling probabilities

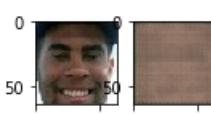
```
N/A%| |ETA: ----- Epoch: nan Iter: nan Class Loss: nan Loss: nan
```



```
3%| |ETA: 0:02:07 Epoch: 0 Iter: 98 Class Loss: 0.02 Loss: 0.06
```



```
6%| |ETA: 0:02:11 Epoch: 0 Iter: 199 Class Loss: 0.02 Loss: 0.06
```



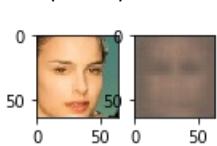
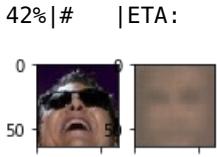
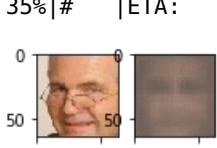
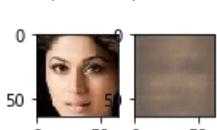
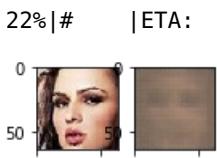
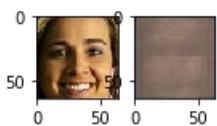
```
9%| |ETA: 0:02:18 Epoch: 0 Iter: 297 Class Loss: 0.03 Loss: 0.07
```

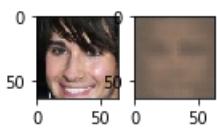


```
13%| |ETA: 0:02:04 Epoch: 0 Iter: 398 Class Loss: 0.08 Loss: 0.12
```

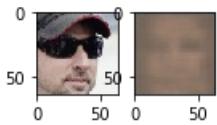


```
16%| |ETA: 0:01:51 Epoch: 0 Iter: 499 Class Loss: 0.02 Loss: 0.05
```

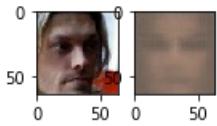




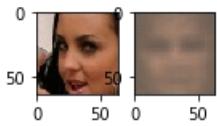
55%|##| ETA: 0:00:58 Epoch: 0 Iter: 1697 Class Loss: 0.00 Loss: 0.04



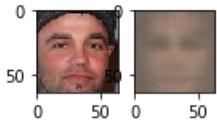
58%|##| ETA: 0:00:54 Epoch: 0 Iter: 1798 Class Loss: 0.00 Loss: 0.04



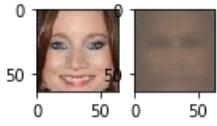
62%|##| ETA: 0:00:50 Epoch: 0 Iter: 1899 Class Loss: 0.00 Loss: 0.04



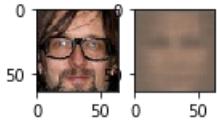
65%|##| ETA: 0:00:52 Epoch: 0 Iter: 1997 Class Loss: 0.00 Loss: 0.03



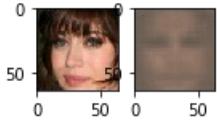
68%|##| ETA: 0:00:47 Epoch: 0 Iter: 2097 Class Loss: 0.00 Loss: 0.04



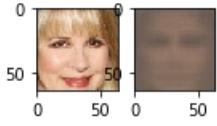
71%|##| ETA: 0:00:37 Epoch: 0 Iter: 2197 Class Loss: 0.00 Loss: 0.04



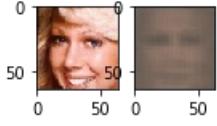
75%|###| ETA: 0:00:32 Epoch: 0 Iter: 2298 Class Loss: 0.00 Loss: 0.04



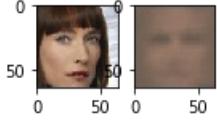
78%|###| ETA: 0:00:28 Epoch: 0 Iter: 2398 Class Loss: 0.01 Loss: 0.04



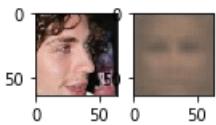
81%|###| ETA: 0:00:24 Epoch: 0 Iter: 2499 Class Loss: 0.00 Loss: 0.04



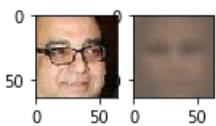
85%|###| ETA: 0:00:19 Epoch: 0 Iter: 2598 Class Loss: 0.03 Loss: 0.07



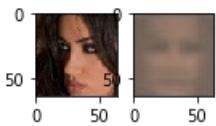
88%|###| ETA: 0:00:15 Epoch: 0 Iter: 2698 Class Loss: 0.02 Loss: 0.06



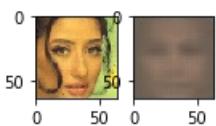
91%|###|ETA: 0:00:11 Epoch: 0 Iter: 2797 Class Loss: 0.00 Loss: 0.04



94%|###|ETA: 0:00:06 Epoch: 0 Iter: 2898 Class Loss: 0.00 Loss: 0.03



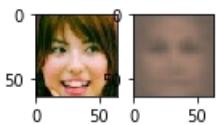
98%|###|ETA: 0:00:02 Epoch: 0 Iter: 2999 Class Loss: 0.00 Loss: 0.04



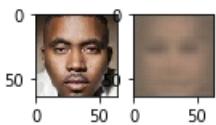
100%|####|Time: 0:02:20 Epoch: 0 Iter: 3052 Class Loss: 0.01 Loss: 0.04

Recomputing the sampling probabilities

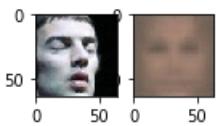
N/A%| |ETA: --- Epoch: nan Iter: nan Class Loss: nan Loss: nan



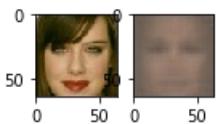
3%| |ETA: 0:02:07 Epoch: 1 Iter: 98 Class Loss: 0.00 Loss: 0.03



6%| |ETA: 0:02:04 Epoch: 1 Iter: 199 Class Loss: 0.00 Loss: 0.03



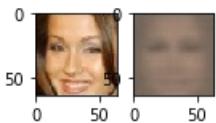
9%| |ETA: 0:01:59 Epoch: 1 Iter: 297 Class Loss: 0.00 Loss: 0.04



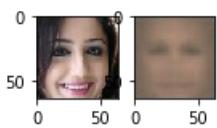
13%| |ETA: 0:01:55 Epoch: 1 Iter: 398 Class Loss: 0.00 Loss: 0.03



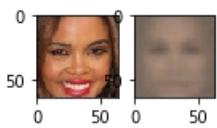
16%| |ETA: 0:01:50 Epoch: 1 Iter: 499 Class Loss: 0.00 Loss: 0.03



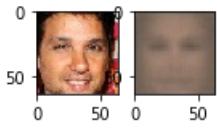
19%| |ETA: 0:01:53 Epoch: 1 Iter: 597 Class Loss: 0.02 Loss: 0.05



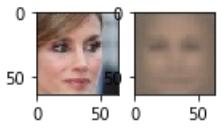
22%|# |ETA: 0:01:58 Epoch: 1 Iter: 698 Class Loss: 0.00 Loss: 0.04



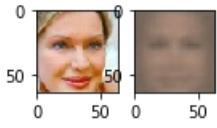
26%|# |ETA: 0:01:44 Epoch: 1 Iter: 799 Class Loss: 0.01 Loss: 0.04



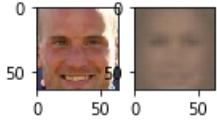
29%|# |ETA: 0:01:32 Epoch: 1 Iter: 897 Class Loss: 0.00 Loss: 0.03



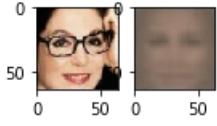
32%|# |ETA: 0:01:28 Epoch: 1 Iter: 997 Class Loss: 0.05 Loss: 0.08



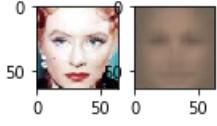
36%|# |ETA: 0:01:24 Epoch: 1 Iter: 1099 Class Loss: 0.01 Loss: 0.04



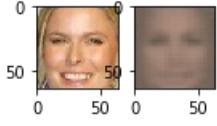
39%|# |ETA: 0:01:20 Epoch: 1 Iter: 1198 Class Loss: 0.00 Loss: 0.04



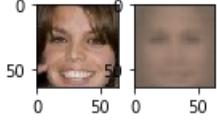
42%|# |ETA: 0:01:16 Epoch: 1 Iter: 1299 Class Loss: 0.01 Loss: 0.03



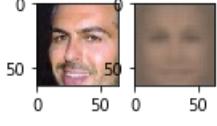
45%|# |ETA: 0:01:11 Epoch: 1 Iter: 1399 Class Loss: 0.00 Loss: 0.03



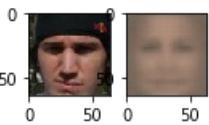
49%|# |ETA: 0:01:07 Epoch: 1 Iter: 1497 Class Loss: 0.03 Loss: 0.06



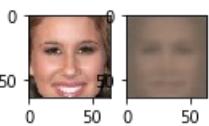
52%|## |ETA: 0:01:03 Epoch: 1 Iter: 1599 Class Loss: 0.00 Loss: 0.03



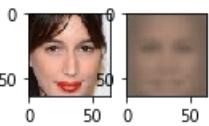
55%|## |ETA: 0:00:58 Epoch: 1 Iter: 1697 Class Loss: 0.00 Loss: 0.03



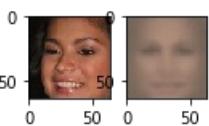
58%|##| ETA: 0:00:54 Epoch: 1 Iter: 1798 Class Loss: 0.00 Loss: 0.03



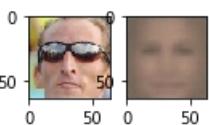
62%|##| ETA: 0:00:50 Epoch: 1 Iter: 1899 Class Loss: 0.15 Loss: 0.18



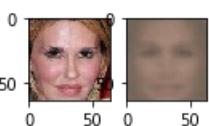
65%|##| ETA: 0:00:45 Epoch: 1 Iter: 1999 Class Loss: 0.00 Loss: 0.03



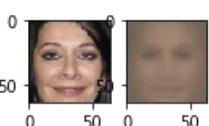
68%|##| ETA: 0:00:41 Epoch: 1 Iter: 2098 Class Loss: 0.00 Loss: 0.03



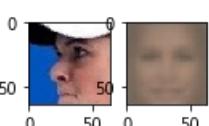
72%|##| ETA: 0:00:36 Epoch: 1 Iter: 2199 Class Loss: 0.00 Loss: 0.03



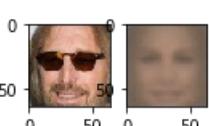
75%|###| ETA: 0:00:32 Epoch: 1 Iter: 2297 Class Loss: 0.00 Loss: 0.04



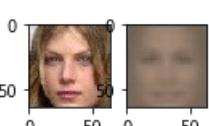
78%|###| ETA: 0:00:32 Epoch: 1 Iter: 2399 Class Loss: 0.00 Loss: 0.04



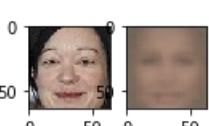
81%|###| ETA: 0:00:28 Epoch: 1 Iter: 2497 Class Loss: 0.00 Loss: 0.03



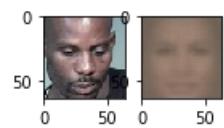
85%|###| ETA: 0:00:19 Epoch: 1 Iter: 2597 Class Loss: 0.00 Loss: 0.03



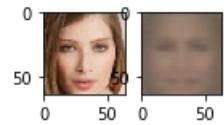
88%|###| ETA: 0:00:15 Epoch: 1 Iter: 2698 Class Loss: 0.03 Loss: 0.07



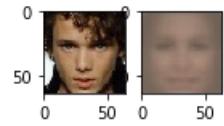
91%|###| ETA: 0:00:11 Epoch: 1 Iter: 2797 Class Loss: 0.00 Loss: 0.04



94%|###| ETA: 0:00:08 Epoch: 1 Iter: 2898 Class Loss: 0.00 Loss: 0.03



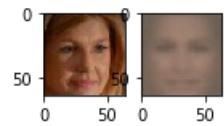
98%|###| ETA: 0:00:02 Epoch: 1 Iter: 2997 Class Loss: 0.00 Loss: 0.03



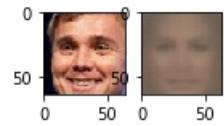
100%|####| Time: 0:02:21 Epoch: 1 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

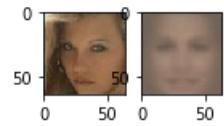
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



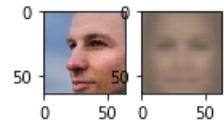
3%| | ETA: 0:02:06 Epoch: 2 Iter: 98 Class Loss: 0.00 Loss: 0.03



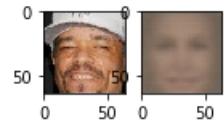
6%| | ETA: 0:02:03 Epoch: 2 Iter: 199 Class Loss: 0.00 Loss: 0.04



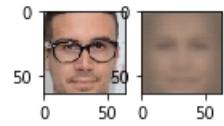
9%| | ETA: 0:01:58 Epoch: 2 Iter: 297 Class Loss: 0.00 Loss: 0.04



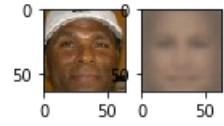
13%| | ETA: 0:01:53 Epoch: 2 Iter: 398 Class Loss: 0.00 Loss: 0.03



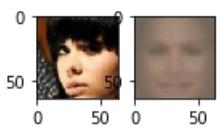
16%| | ETA: 0:01:49 Epoch: 2 Iter: 499 Class Loss: 0.00 Loss: 0.03



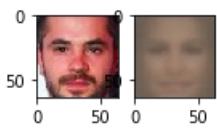
19%| | ETA: 0:01:45 Epoch: 2 Iter: 597 Class Loss: 0.01 Loss: 0.04



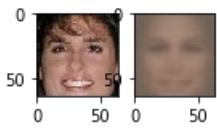
22%|#| ETA: 0:01:41 Epoch: 2 Iter: 698 Class Loss: 0.00 Loss: 0.03



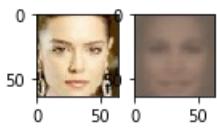
26%|# |ETA: 0:01:36 Epoch: 2 Iter: 798 Class Loss: 0.00 Loss: 0.03



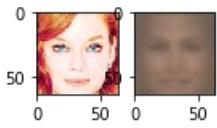
29%|# |ETA: 0:01:32 Epoch: 2 Iter: 897 Class Loss: 0.01 Loss: 0.04



32%|# |ETA: 0:01:31 Epoch: 2 Iter: 997 Class Loss: 0.00 Loss: 0.03



35%|# |ETA: 0:01:37 Epoch: 2 Iter: 1098 Class Loss: 0.00 Loss: 0.03



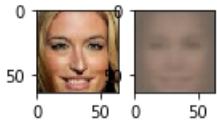
39%|# |ETA: 0:01:28 Epoch: 2 Iter: 1197 Class Loss: 0.00 Loss: 0.03



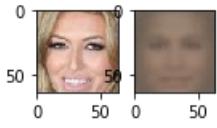
42%|# |ETA: 0:01:14 Epoch: 2 Iter: 1298 Class Loss: 0.10 Loss: 0.14



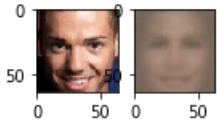
45%|# |ETA: 0:01:10 Epoch: 2 Iter: 1399 Class Loss: 0.00 Loss: 0.04



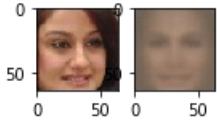
49%|# |ETA: 0:01:06 Epoch: 2 Iter: 1497 Class Loss: 0.00 Loss: 0.03



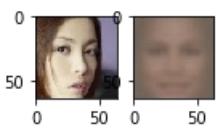
52%|## |ETA: 0:01:02 Epoch: 2 Iter: 1599 Class Loss: 0.00 Loss: 0.03



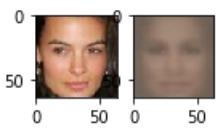
55%|## |ETA: 0:00:57 Epoch: 2 Iter: 1697 Class Loss: 0.00 Loss: 0.03



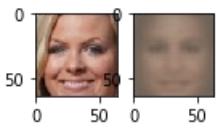
58%|## |ETA: 0:00:54 Epoch: 2 Iter: 1798 Class Loss: 0.00 Loss: 0.04



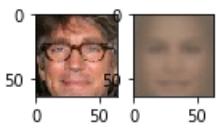
62%|##| ETA: 0:00:49 Epoch: 2 Iter: 1899 Class Loss: 0.00 Loss: 0.03



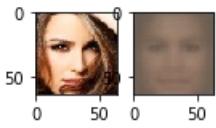
65%|##| ETA: 0:00:54 Epoch: 2 Iter: 1998 Class Loss: 0.00 Loss: 0.04



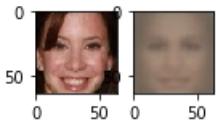
68%|##| ETA: 0:00:41 Epoch: 2 Iter: 2098 Class Loss: 0.04 Loss: 0.07



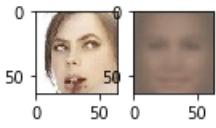
71%|##| ETA: 0:00:36 Epoch: 2 Iter: 2197 Class Loss: 0.02 Loss: 0.05



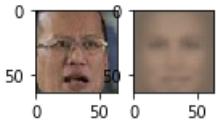
75%|###| ETA: 0:00:32 Epoch: 2 Iter: 2298 Class Loss: 0.00 Loss: 0.03



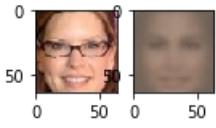
78%|###| ETA: 0:00:28 Epoch: 2 Iter: 2399 Class Loss: 0.00 Loss: 0.04



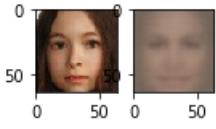
81%|###| ETA: 0:00:24 Epoch: 2 Iter: 2497 Class Loss: 0.00 Loss: 0.04



85%|###| ETA: 0:00:19 Epoch: 2 Iter: 2598 Class Loss: 0.00 Loss: 0.03



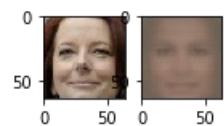
88%|###| ETA: 0:00:15 Epoch: 2 Iter: 2697 Class Loss: 0.00 Loss: 0.04



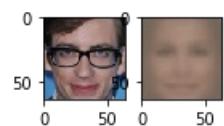
91%|###| ETA: 0:00:12 Epoch: 2 Iter: 2798 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:07 Epoch: 2 Iter: 2898 Class Loss: 0.00 Loss: 0.03



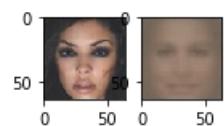
98%|###| ETA: 0:00:02 Epoch: 2 Iter: 2999 Class Loss: 0.00 Loss: 0.03



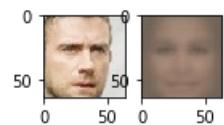
100%|####| Time: 0:02:18 Epoch: 2 Iter: 3052 Class Loss: 0.01 Loss: 0.04

Recomputing the sampling probabilities

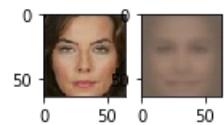
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



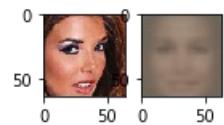
3%| | ETA: 0:02:06 Epoch: 3 Iter: 98 Class Loss: 0.00 Loss: 0.04



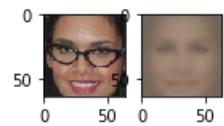
6%| | ETA: 0:02:02 Epoch: 3 Iter: 198 Class Loss: 0.02 Loss: 0.05



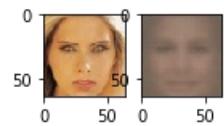
9%| | ETA: 0:01:58 Epoch: 3 Iter: 299 Class Loss: 0.01 Loss: 0.04



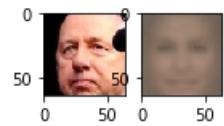
13%| | ETA: 0:01:53 Epoch: 3 Iter: 398 Class Loss: 0.00 Loss: 0.03



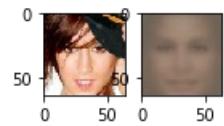
16%| | ETA: 0:01:49 Epoch: 3 Iter: 499 Class Loss: 0.00 Loss: 0.03



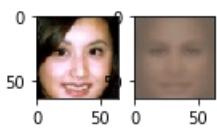
19%| | ETA: 0:01:45 Epoch: 3 Iter: 597 Class Loss: 0.00 Loss: 0.03



22%|#| ETA: 0:01:41 Epoch: 3 Iter: 698 Class Loss: 0.00 Loss: 0.04



26%|#| ETA: 0:01:37 Epoch: 3 Iter: 798 Class Loss: 0.00 Loss: 0.03



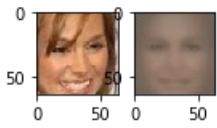
29%|# |ETA: 0:01:32 Epoch: 3 Iter: 897 Class Loss: 0.00 Loss: 0.03



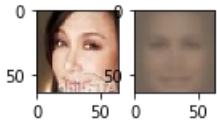
32%|# |ETA: 0:01:28 Epoch: 3 Iter: 997 Class Loss: 0.00 Loss: 0.03



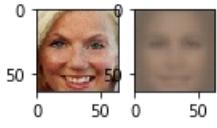
36%|# |ETA: 0:01:24 Epoch: 3 Iter: 1099 Class Loss: 0.05 Loss: 0.08



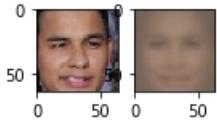
39%|# |ETA: 0:01:19 Epoch: 3 Iter: 1198 Class Loss: 0.01 Loss: 0.04



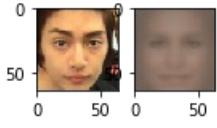
42%|# |ETA: 0:01:15 Epoch: 3 Iter: 1299 Class Loss: 0.00 Loss: 0.04



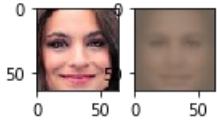
45%|# |ETA: 0:01:11 Epoch: 3 Iter: 1397 Class Loss: 0.00 Loss: 0.04



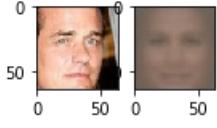
49%|# |ETA: 0:01:17 Epoch: 3 Iter: 1497 Class Loss: 0.00 Loss: 0.04



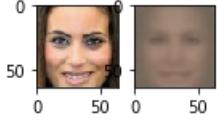
52%## |ETA: 0:01:12 Epoch: 3 Iter: 1598 Class Loss: 0.10 Loss: 0.13



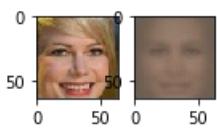
55%## |ETA: 0:00:58 Epoch: 3 Iter: 1697 Class Loss: 0.00 Loss: 0.03



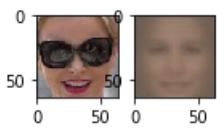
58%## |ETA: 0:00:54 Epoch: 3 Iter: 1798 Class Loss: 0.00 Loss: 0.03



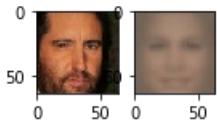
62%## |ETA: 0:00:49 Epoch: 3 Iter: 1899 Class Loss: 0.00 Loss: 0.03



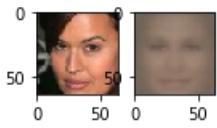
65%|##| ETA: 0:00:45 Epoch: 3 Iter: 1997 Class Loss: 0.00 Loss: 0.03



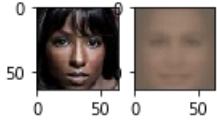
68%|##| ETA: 0:00:41 Epoch: 3 Iter: 2098 Class Loss: 0.00 Loss: 0.04



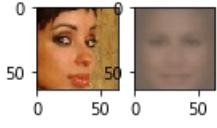
71%|##| ETA: 0:00:36 Epoch: 3 Iter: 2197 Class Loss: 0.00 Loss: 0.04



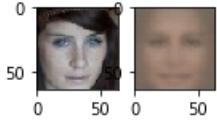
75%|###| ETA: 0:00:32 Epoch: 3 Iter: 2298 Class Loss: 0.00 Loss: 0.03



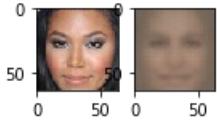
78%|###| ETA: 0:00:27 Epoch: 3 Iter: 2398 Class Loss: 0.00 Loss: 0.03



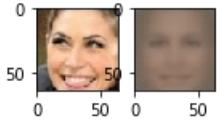
81%|###| ETA: 0:00:23 Epoch: 3 Iter: 2497 Class Loss: 0.00 Loss: 0.03



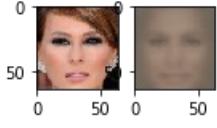
85%|###| ETA: 0:00:19 Epoch: 3 Iter: 2598 Class Loss: 0.00 Loss: 0.04



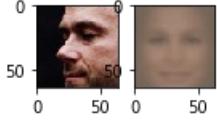
88%|###| ETA: 0:00:15 Epoch: 3 Iter: 2698 Class Loss: 0.00 Loss: 0.03



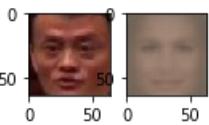
91%|###| ETA: 0:00:10 Epoch: 3 Iter: 2799 Class Loss: 0.01 Loss: 0.04



94%|###| ETA: 0:00:06 Epoch: 3 Iter: 2898 Class Loss: 0.00 Loss: 0.03



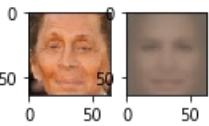
98%|###| ETA: 0:00:02 Epoch: 3 Iter: 2999 Class Loss: 0.00 Loss: 0.04



100%|####|Time: 0:02:17 Epoch: 3 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

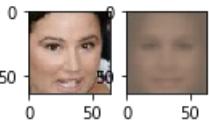
N/A%| |ETA: ---- Epoch: nan Iter: nan Class Loss: nan Loss: nan



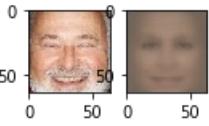
3%| |ETA: 0:02:07 Epoch: 4 Iter: 98 Class Loss: 0.00 Loss: 0.04



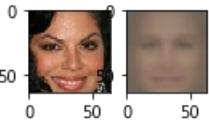
6%| |ETA: 0:02:20 Epoch: 4 Iter: 199 Class Loss: 0.00 Loss: 0.04



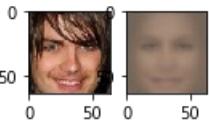
9%| |ETA: 0:02:16 Epoch: 4 Iter: 298 Class Loss: 0.00 Loss: 0.03



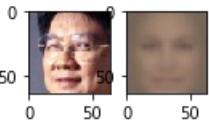
13%| |ETA: 0:02:18 Epoch: 4 Iter: 398 Class Loss: 0.00 Loss: 0.03



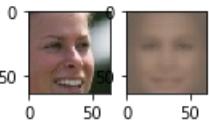
16%| |ETA: 0:02:05 Epoch: 4 Iter: 497 Class Loss: 0.00 Loss: 0.03



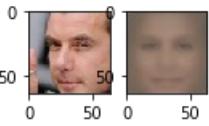
19%| |ETA: 0:01:44 Epoch: 4 Iter: 597 Class Loss: 0.00 Loss: 0.03



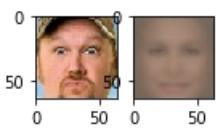
22%|# |ETA: 0:01:40 Epoch: 4 Iter: 698 Class Loss: 0.00 Loss: 0.03



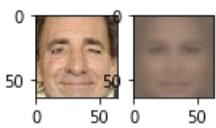
26%|# |ETA: 0:01:36 Epoch: 4 Iter: 799 Class Loss: 0.00 Loss: 0.04



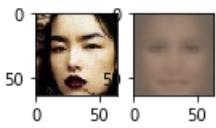
29%|# |ETA: 0:01:32 Epoch: 4 Iter: 897 Class Loss: 0.00 Loss: 0.03



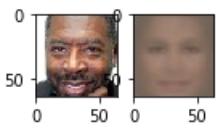
32%|# |ETA: 0:01:27 Epoch: 4 Iter: 999 Class Loss: 0.00 Loss: 0.03



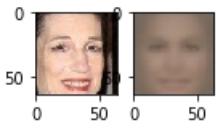
36%|# |ETA: 0:01:23 Epoch: 4 Iter: 1099 Class Loss: 0.00 Loss: 0.03



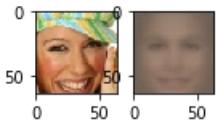
39%|# |ETA: 0:01:19 Epoch: 4 Iter: 1197 Class Loss: 0.00 Loss: 0.03



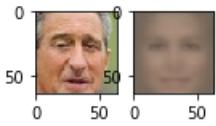
42%|# |ETA: 0:01:15 Epoch: 4 Iter: 1298 Class Loss: 0.00 Loss: 0.04



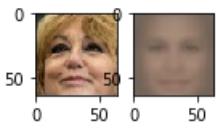
45%|# |ETA: 0:01:11 Epoch: 4 Iter: 1399 Class Loss: 0.00 Loss: 0.03



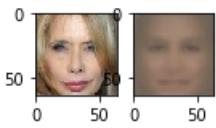
49%|# |ETA: 0:01:06 Epoch: 4 Iter: 1498 Class Loss: 0.00 Loss: 0.04



52%|## |ETA: 0:01:02 Epoch: 4 Iter: 1598 Class Loss: 0.00 Loss: 0.03



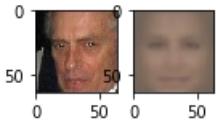
55%|## |ETA: 0:00:58 Epoch: 4 Iter: 1699 Class Loss: 0.00 Loss: 0.03



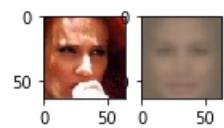
58%|## |ETA: 0:00:53 Epoch: 4 Iter: 1798 Class Loss: 0.00 Loss: 0.04



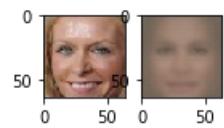
62%|## |ETA: 0:00:57 Epoch: 4 Iter: 1898 Class Loss: 0.00 Loss: 0.03



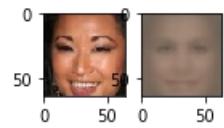
65%|## |ETA: 0:00:52 Epoch: 4 Iter: 1999 Class Loss: 0.01 Loss: 0.04



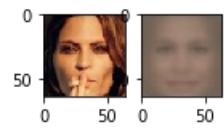
68%|##| ETA: 0:00:40 Epoch: 4 Iter: 2098 Class Loss: 0.00 Loss: 0.03



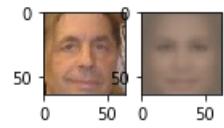
72%|##| ETA: 0:00:36 Epoch: 4 Iter: 2199 Class Loss: 0.00 Loss: 0.03



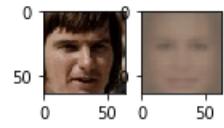
75%|###| ETA: 0:00:32 Epoch: 4 Iter: 2298 Class Loss: 0.00 Loss: 0.03



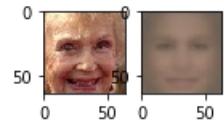
78%|###| ETA: 0:00:28 Epoch: 4 Iter: 2399 Class Loss: 0.00 Loss: 0.03



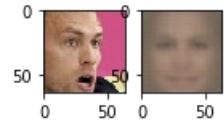
81%|###| ETA: 0:00:28 Epoch: 4 Iter: 2499 Class Loss: 0.00 Loss: 0.03



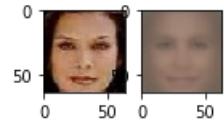
85%|###| ETA: 0:00:19 Epoch: 4 Iter: 2598 Class Loss: 0.02 Loss: 0.05



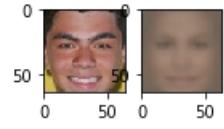
88%|###| ETA: 0:00:15 Epoch: 4 Iter: 2699 Class Loss: 0.00 Loss: 0.03



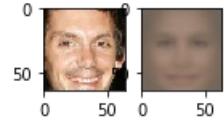
91%|###| ETA: 0:00:10 Epoch: 4 Iter: 2797 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 4 Iter: 2898 Class Loss: 0.00 Loss: 0.04



98%|###| ETA: 0:00:02 Epoch: 4 Iter: 2999 Class Loss: 0.00 Loss: 0.03



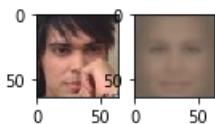
100%|####| Time: 0:02:21 Epoch: 4 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

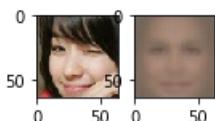
N/A%| |ETA: --:--- Epoch: nan Iter: nan Class Loss: nan Loss: nan



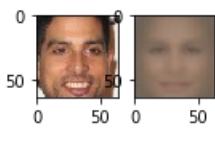
3%| |ETA: 0:02:07 Epoch: 5 Iter: 97 Class Loss: 0.00 Loss: 0.04



6%| |ETA: 0:02:03 Epoch: 5 Iter: 198 Class Loss: 0.00 Loss: 0.03



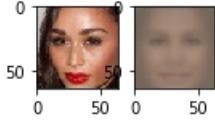
9%| |ETA: 0:01:58 Epoch: 5 Iter: 299 Class Loss: 0.00 Loss: 0.03



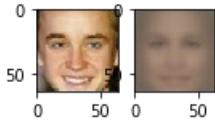
13%| |ETA: 0:01:53 Epoch: 5 Iter: 398 Class Loss: 0.00 Loss: 0.03



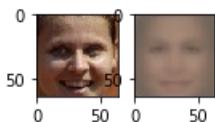
16%| |ETA: 0:01:51 Epoch: 5 Iter: 498 Class Loss: 0.00 Loss: 0.04



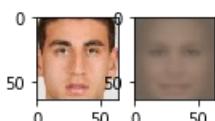
19%| |ETA: 0:02:02 Epoch: 5 Iter: 599 Class Loss: 0.00 Loss: 0.04



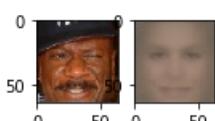
22%|# |ETA: 0:01:52 Epoch: 5 Iter: 697 Class Loss: 0.00 Loss: 0.03



26%|# |ETA: 0:01:37 Epoch: 5 Iter: 799 Class Loss: 0.01 Loss: 0.04



29%|# |ETA: 0:01:33 Epoch: 5 Iter: 897 Class Loss: 0.00 Loss: 0.03



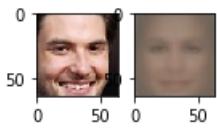
32%|# |ETA: 0:01:28 Epoch: 5 Iter: 998 Class Loss: 0.00 Loss: 0.03



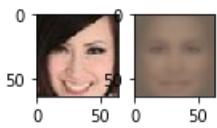
36%|# |ETA: 0:01:24 Epoch: 5 Iter: 1099 Class Loss: 0.00 Loss: 0.03



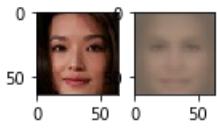
39%|# |ETA: 0:01:20 Epoch: 5 Iter: 1197 Class Loss: 0.00 Loss: 0.03



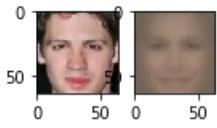
42%|# |ETA: 0:01:15 Epoch: 5 Iter: 1298 Class Loss: 0.00 Loss: 0.04



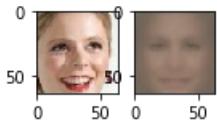
45%|# |ETA: 0:01:11 Epoch: 5 Iter: 1399 Class Loss: 0.01 Loss: 0.04



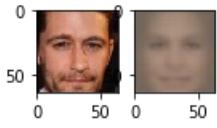
49%|# |ETA: 0:01:06 Epoch: 5 Iter: 1498 Class Loss: 0.00 Loss: 0.03



52%|## |ETA: 0:01:02 Epoch: 5 Iter: 1599 Class Loss: 0.00 Loss: 0.03



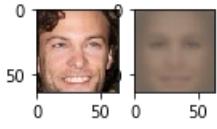
55%|## |ETA: 0:00:58 Epoch: 5 Iter: 1697 Class Loss: 0.00 Loss: 0.03



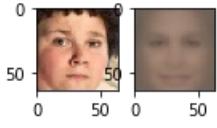
58%|## |ETA: 0:00:54 Epoch: 5 Iter: 1798 Class Loss: 0.00 Loss: 0.04



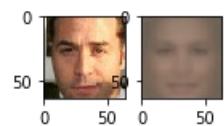
62%|## |ETA: 0:00:49 Epoch: 5 Iter: 1898 Class Loss: 0.00 Loss: 0.04



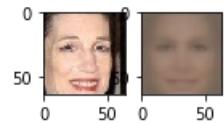
65%|## |ETA: 0:00:45 Epoch: 5 Iter: 1997 Class Loss: 0.01 Loss: 0.04



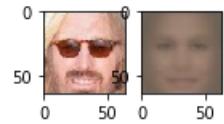
68%|## |ETA: 0:00:41 Epoch: 5 Iter: 2098 Class Loss: 0.00 Loss: 0.03



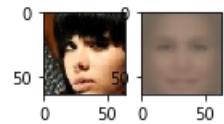
71%|##| ETA: 0:00:37 Epoch: 5 Iter: 2197 Class Loss: 0.00 Loss: 0.04



75%|###| ETA: 0:00:36 Epoch: 5 Iter: 2299 Class Loss: 0.00 Loss: 0.03



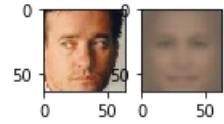
78%|###| ETA: 0:00:32 Epoch: 5 Iter: 2398 Class Loss: 0.00 Loss: 0.04



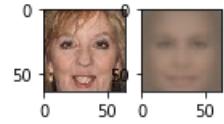
81%|###| ETA: 0:00:24 Epoch: 5 Iter: 2497 Class Loss: 0.00 Loss: 0.03



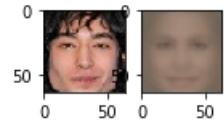
85%|###| ETA: 0:00:19 Epoch: 5 Iter: 2598 Class Loss: 0.00 Loss: 0.04



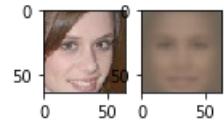
88%|###| ETA: 0:00:15 Epoch: 5 Iter: 2699 Class Loss: 0.00 Loss: 0.04



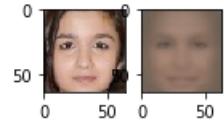
91%|###| ETA: 0:00:10 Epoch: 5 Iter: 2799 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 5 Iter: 2898 Class Loss: 0.00 Loss: 0.03



98%|###| ETA: 0:00:02 Epoch: 5 Iter: 2999 Class Loss: 0.00 Loss: 0.04



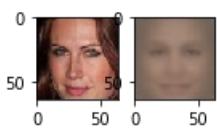
100%|####| Time: 0:02:18 Epoch: 5 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

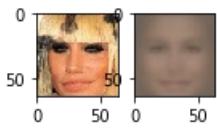
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



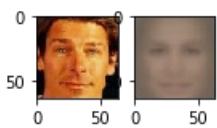
3% | ETA: 0:02:06 Epoch: 6 Iter: 98 Class Loss: 0.02 Loss: 0.06



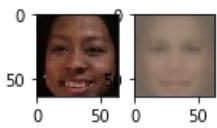
6% | ETA: 0:02:02 Epoch: 6 Iter: 199 Class Loss: 0.00 Loss: 0.03



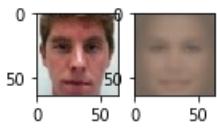
9% | ETA: 0:01:58 Epoch: 6 Iter: 297 Class Loss: 0.00 Loss: 0.03



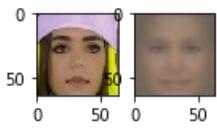
13% | ETA: 0:01:55 Epoch: 6 Iter: 398 Class Loss: 0.00 Loss: 0.03



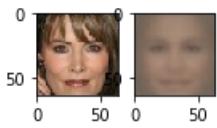
16% | ETA: 0:01:51 Epoch: 6 Iter: 499 Class Loss: 0.00 Loss: 0.03



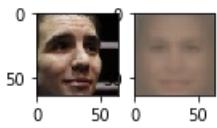
19% | ETA: 0:01:45 Epoch: 6 Iter: 599 Class Loss: 0.00 Loss: 0.03



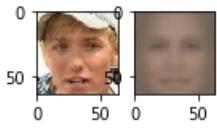
22% | # ETA: 0:01:42 Epoch: 6 Iter: 697 Class Loss: 0.00 Loss: 0.03



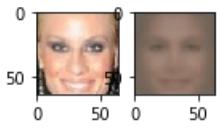
26% | # ETA: 0:01:37 Epoch: 6 Iter: 799 Class Loss: 0.00 Loss: 0.04



29% | # ETA: 0:01:32 Epoch: 6 Iter: 897 Class Loss: 0.02 Loss: 0.06



32% | # ETA: 0:01:48 Epoch: 6 Iter: 998 Class Loss: 0.00 Loss: 0.04



36% | # ETA: 0:01:43 Epoch: 6 Iter: 1099 Class Loss: 0.00 Loss: 0.04



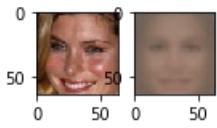
39%|# |ETA: 0:01:19 Epoch: 6 Iter: 1197 Class Loss: 0.00 Loss: 0.04



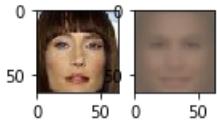
42%|# |ETA: 0:01:15 Epoch: 6 Iter: 1299 Class Loss: 0.00 Loss: 0.03



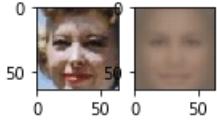
45%|# |ETA: 0:01:10 Epoch: 6 Iter: 1399 Class Loss: 0.14 Loss: 0.18



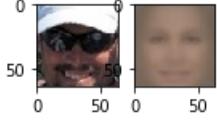
49%|# |ETA: 0:01:07 Epoch: 6 Iter: 1498 Class Loss: 0.00 Loss: 0.03



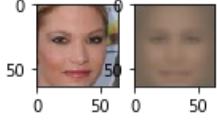
52%|## |ETA: 0:01:02 Epoch: 6 Iter: 1599 Class Loss: 0.00 Loss: 0.03



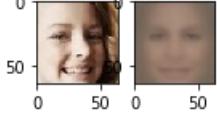
55%|## |ETA: 0:00:58 Epoch: 6 Iter: 1697 Class Loss: 0.05 Loss: 0.08



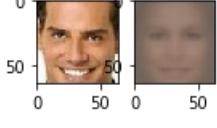
58%|## |ETA: 0:00:54 Epoch: 6 Iter: 1798 Class Loss: 0.00 Loss: 0.03



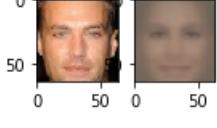
62%|## |ETA: 0:00:49 Epoch: 6 Iter: 1899 Class Loss: 0.00 Loss: 0.04



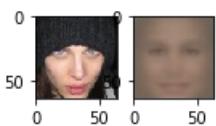
65%|## |ETA: 0:00:45 Epoch: 6 Iter: 1997 Class Loss: 0.00 Loss: 0.03



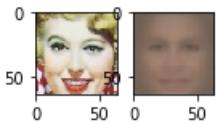
68%|## |ETA: 0:00:41 Epoch: 6 Iter: 2098 Class Loss: 0.00 Loss: 0.04



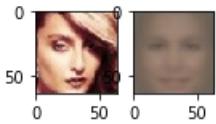
72%|## |ETA: 0:00:37 Epoch: 6 Iter: 2198 Class Loss: 0.00 Loss: 0.03



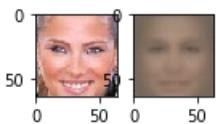
75%|###| ETA: 0:00:32 Epoch: 6 Iter: 2298 Class Loss: 0.00 Loss: 0.03



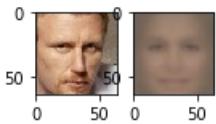
78%|###| ETA: 0:00:28 Epoch: 6 Iter: 2399 Class Loss: 0.00 Loss: 0.03



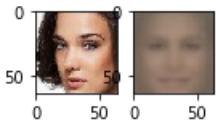
81%|###| ETA: 0:00:23 Epoch: 6 Iter: 2499 Class Loss: 0.12 Loss: 0.15



85%|###| ETA: 0:00:19 Epoch: 6 Iter: 2598 Class Loss: 0.00 Loss: 0.03



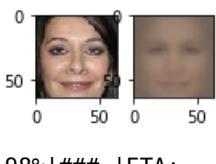
88%|###| ETA: 0:00:15 Epoch: 6 Iter: 2699 Class Loss: 0.00 Loss: 0.03



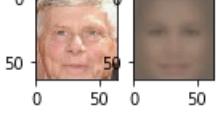
91%|###| ETA: 0:00:12 Epoch: 6 Iter: 2798 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:07 Epoch: 6 Iter: 2898 Class Loss: 0.00 Loss: 0.03



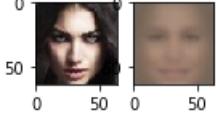
98%|###| ETA: 0:00:02 Epoch: 6 Iter: 2997 Class Loss: 0.00 Loss: 0.04



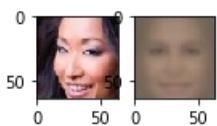
100%|####| Time: 0:02:20 Epoch: 6 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

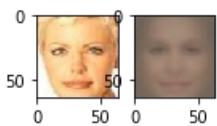
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



3%| | ETA: 0:02:08 Epoch: 7 Iter: 98 Class Loss: 0.01 Loss: 0.03



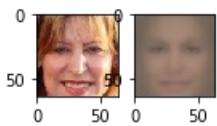
6% | ETA: 0:02:03 Epoch: 7 Iter: 198 Class Loss: 0.00 Loss: 0.03



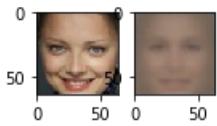
9% | ETA: 0:01:58 Epoch: 7 Iter: 299 Class Loss: 0.00 Loss: 0.04



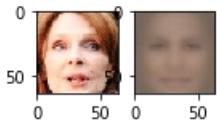
13% | ETA: 0:01:55 Epoch: 7 Iter: 397 Class Loss: 0.00 Loss: 0.03



16% | ETA: 0:01:49 Epoch: 7 Iter: 499 Class Loss: 0.00 Loss: 0.03



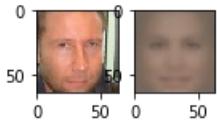
19% | ETA: 0:01:46 Epoch: 7 Iter: 597 Class Loss: 0.00 Loss: 0.03



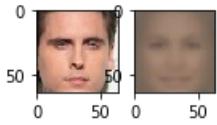
22% | # ETA: 0:01:41 Epoch: 7 Iter: 698 Class Loss: 0.00 Loss: 0.03



26% | # ETA: 0:01:36 Epoch: 7 Iter: 799 Class Loss: 0.00 Loss: 0.03



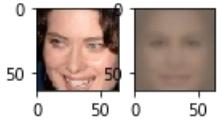
29% | # ETA: 0:01:32 Epoch: 7 Iter: 897 Class Loss: 0.00 Loss: 0.03



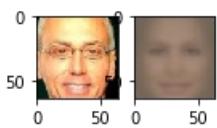
32% | # ETA: 0:01:27 Epoch: 7 Iter: 998 Class Loss: 0.00 Loss: 0.04



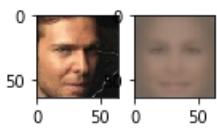
36% | # ETA: 0:01:24 Epoch: 7 Iter: 1099 Class Loss: 0.00 Loss: 0.04



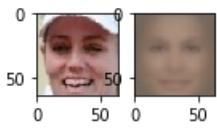
39% | # ETA: 0:01:19 Epoch: 7 Iter: 1198 Class Loss: 0.00 Loss: 0.02



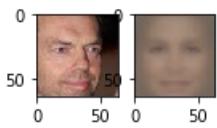
42%|# |ETA: 0:01:15 Epoch: 7 Iter: 1299 Class Loss: 0.00 Loss: 0.03



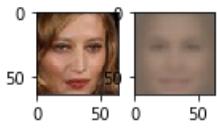
45%|# |ETA: 0:01:21 Epoch: 7 Iter: 1397 Class Loss: 0.00 Loss: 0.03



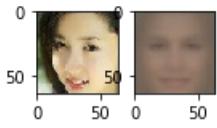
49%|# |ETA: 0:01:16 Epoch: 7 Iter: 1498 Class Loss: 0.00 Loss: 0.03



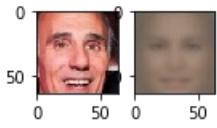
52%|## |ETA: 0:01:02 Epoch: 7 Iter: 1599 Class Loss: 0.00 Loss: 0.03



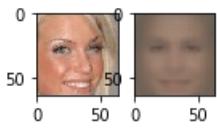
55%|## |ETA: 0:00:58 Epoch: 7 Iter: 1697 Class Loss: 0.00 Loss: 0.04



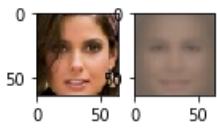
58%|## |ETA: 0:00:53 Epoch: 7 Iter: 1798 Class Loss: 0.00 Loss: 0.03



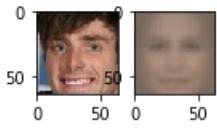
62%|## |ETA: 0:00:49 Epoch: 7 Iter: 1899 Class Loss: 0.00 Loss: 0.04



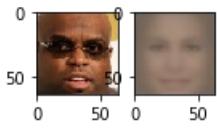
65%|## |ETA: 0:00:45 Epoch: 7 Iter: 1999 Class Loss: 0.00 Loss: 0.03



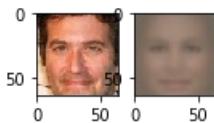
68%|## |ETA: 0:00:40 Epoch: 7 Iter: 2098 Class Loss: 0.00 Loss: 0.04



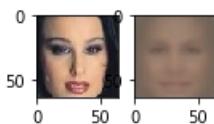
71%|## |ETA: 0:00:36 Epoch: 7 Iter: 2197 Class Loss: 0.00 Loss: 0.03



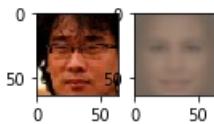
75%|### |ETA: 0:00:32 Epoch: 7 Iter: 2298 Class Loss: 0.00 Loss: 0.03



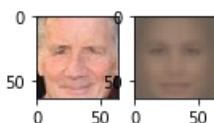
78%|###| ETA: 0:00:28 Epoch: 7 Iter: 2398 Class Loss: 0.00 Loss: 0.03



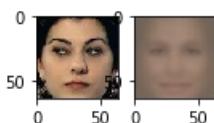
81%|###| ETA: 0:00:23 Epoch: 7 Iter: 2497 Class Loss: 0.00 Loss: 0.03



85%|###| ETA: 0:00:19 Epoch: 7 Iter: 2597 Class Loss: 0.00 Loss: 0.03



88%|###| ETA: 0:00:15 Epoch: 7 Iter: 2698 Class Loss: 0.00 Loss: 0.03



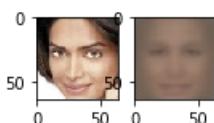
91%|###| ETA: 0:00:10 Epoch: 7 Iter: 2797 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 7 Iter: 2898 Class Loss: 0.00 Loss: 0.03



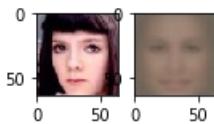
98%|###| ETA: 0:00:02 Epoch: 7 Iter: 2999 Class Loss: 0.00 Loss: 0.03



100%|####| Time: 0:02:16 Epoch: 7 Iter: 3052 Class Loss: 0.00 Loss: 0.03

Recomputing the sampling probabilities

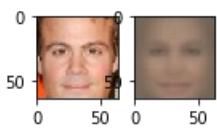
N/A%| | ETA: --:--:-- Epoch: nan Iter: nan Class Loss: nan Loss: nan



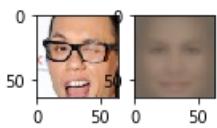
3%| | ETA: 0:02:28 Epoch: 8 Iter: 99 Class Loss: 0.00 Loss: 0.04



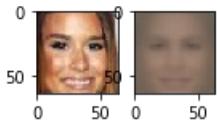
6%| | ETA: 0:02:20 Epoch: 8 Iter: 199 Class Loss: 0.01 Loss: 0.04



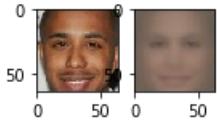
9% | ETA: 0:01:57 Epoch: 8 Iter: 297 Class Loss: 0.00 Loss: 0.03



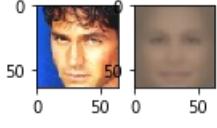
13% | ETA: 0:01:53 Epoch: 8 Iter: 398 Class Loss: 0.00 Loss: 0.03



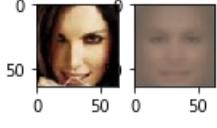
16% | ETA: 0:01:49 Epoch: 8 Iter: 499 Class Loss: 0.00 Loss: 0.03



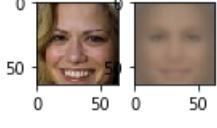
19% | ETA: 0:01:45 Epoch: 8 Iter: 597 Class Loss: 0.00 Loss: 0.03



22% | # ETA: 0:01:41 Epoch: 8 Iter: 698 Class Loss: 0.00 Loss: 0.03



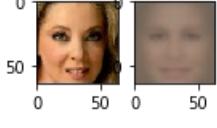
26% | # ETA: 0:01:37 Epoch: 8 Iter: 799 Class Loss: 0.00 Loss: 0.03



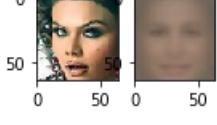
29% | # ETA: 0:01:32 Epoch: 8 Iter: 897 Class Loss: 0.00 Loss: 0.04



32% | # ETA: 0:01:28 Epoch: 8 Iter: 998 Class Loss: 0.00 Loss: 0.04



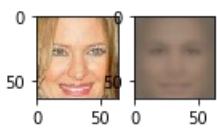
35% | # ETA: 0:01:23 Epoch: 8 Iter: 1097 Class Loss: 0.00 Loss: 0.03



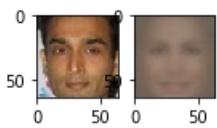
39% | # ETA: 0:01:18 Epoch: 8 Iter: 1198 Class Loss: 0.00 Loss: 0.03



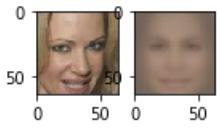
42% | # ETA: 0:01:15 Epoch: 8 Iter: 1299 Class Loss: 0.00 Loss: 0.03



45%|# |ETA: 0:01:10 Epoch: 8 Iter: 1397 Class Loss: 0.00 Loss: 0.03



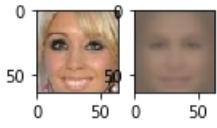
49%|# |ETA: 0:01:06 Epoch: 8 Iter: 1498 Class Loss: 0.04 Loss: 0.07



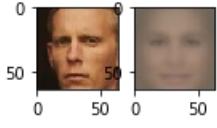
52%|## |ETA: 0:01:16 Epoch: 8 Iter: 1598 Class Loss: 0.00 Loss: 0.03



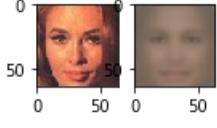
55%|## |ETA: 0:01:11 Epoch: 8 Iter: 1699 Class Loss: 0.00 Loss: 0.03



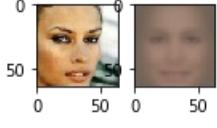
58%|## |ETA: 0:00:58 Epoch: 8 Iter: 1799 Class Loss: 0.00 Loss: 0.03



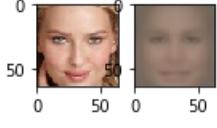
62%|## |ETA: 0:00:56 Epoch: 8 Iter: 1899 Class Loss: 0.00 Loss: 0.03



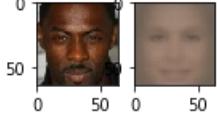
65%|## |ETA: 0:00:47 Epoch: 8 Iter: 1997 Class Loss: 0.00 Loss: 0.03



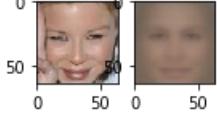
68%|## |ETA: 0:00:41 Epoch: 8 Iter: 2098 Class Loss: 0.00 Loss: 0.03



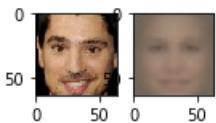
72%|## |ETA: 0:00:36 Epoch: 8 Iter: 2199 Class Loss: 0.00 Loss: 0.03



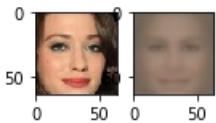
75%|### |ETA: 0:00:32 Epoch: 8 Iter: 2298 Class Loss: 0.00 Loss: 0.03



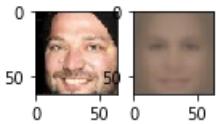
78%|### |ETA: 0:00:27 Epoch: 8 Iter: 2399 Class Loss: 0.00 Loss: 0.04



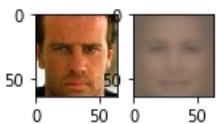
81%|###|ETA: 0:00:23 Epoch: 8 Iter: 2497 Class Loss: 0.00 Loss: 0.03



85%|###|ETA: 0:00:19 Epoch: 8 Iter: 2597 Class Loss: 0.00 Loss: 0.04



88%|###|ETA: 0:00:15 Epoch: 8 Iter: 2698 Class Loss: 0.00 Loss: 0.03



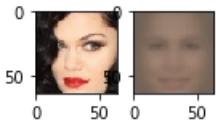
91%|###|ETA: 0:00:10 Epoch: 8 Iter: 2797 Class Loss: 0.00 Loss: 0.03



94%|###|ETA: 0:00:06 Epoch: 8 Iter: 2898 Class Loss: 0.05 Loss: 0.08



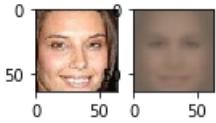
98%|###|ETA: 0:00:02 Epoch: 8 Iter: 2999 Class Loss: 0.00 Loss: 0.03



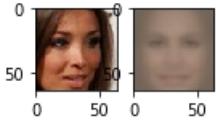
100%|####|Time: 0:02:20 Epoch: 8 Iter: 3052 Class Loss: 0.00 Loss: 0.04

Recomputing the sampling probabilities

N/A%| |ETA: ---- Epoch: nan Iter: nan Class Loss: nan Loss: nan



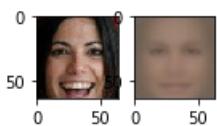
3%| |ETA: 0:02:05 Epoch: 9 Iter: 98 Class Loss: 0.00 Loss: 0.03



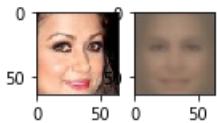
6%| |ETA: 0:02:01 Epoch: 9 Iter: 198 Class Loss: 0.00 Loss: 0.03



9%| |ETA: 0:01:58 Epoch: 9 Iter: 299 Class Loss: 0.00 Loss: 0.03



13% | ETA: 0:02:17 Epoch: 9 Iter: 398 Class Loss: 0.00 Loss: 0.04



16% | ETA: 0:02:06 Epoch: 9 Iter: 497 Class Loss: 0.00 Loss: 0.03



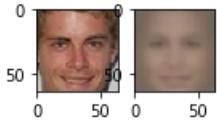
19% | ETA: 0:02:01 Epoch: 9 Iter: 598 Class Loss: 0.00 Loss: 0.03



22% | # ETA: 0:01:40 Epoch: 9 Iter: 697 Class Loss: 0.00 Loss: 0.04



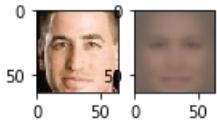
26% | # ETA: 0:01:37 Epoch: 9 Iter: 799 Class Loss: 0.00 Loss: 0.03



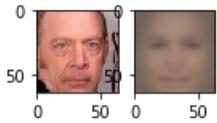
29% | # ETA: 0:01:32 Epoch: 9 Iter: 897 Class Loss: 0.01 Loss: 0.04



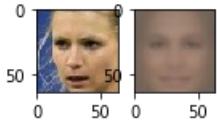
32% | # ETA: 0:01:27 Epoch: 9 Iter: 998 Class Loss: 0.00 Loss: 0.04



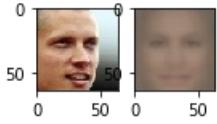
36% | # ETA: 0:01:23 Epoch: 9 Iter: 1099 Class Loss: 0.00 Loss: 0.03



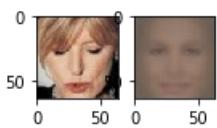
39% | # ETA: 0:01:19 Epoch: 9 Iter: 1197 Class Loss: 0.00 Loss: 0.03



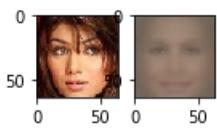
42% | # ETA: 0:01:15 Epoch: 9 Iter: 1298 Class Loss: 0.00 Loss: 0.04



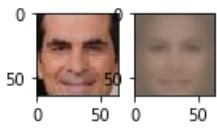
45% | # ETA: 0:01:10 Epoch: 9 Iter: 1397 Class Loss: 0.00 Loss: 0.03



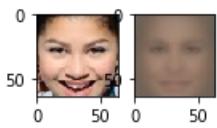
49%|#| ETA: 0:01:06 Epoch: 9 Iter: 1498 Class Loss: 0.00 Loss: 0.04



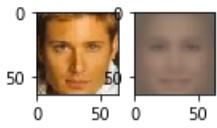
52%|##| ETA: 0:01:02 Epoch: 9 Iter: 1599 Class Loss: 0.00 Loss: 0.03



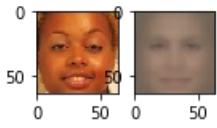
55%|##| ETA: 0:00:58 Epoch: 9 Iter: 1697 Class Loss: 0.00 Loss: 0.03



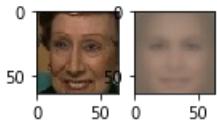
58%|##| ETA: 0:00:53 Epoch: 9 Iter: 1798 Class Loss: 0.07 Loss: 0.11



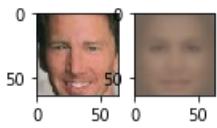
62%|##| ETA: 0:00:49 Epoch: 9 Iter: 1899 Class Loss: 0.00 Loss: 0.03



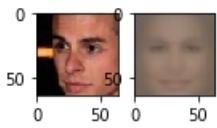
65%|##| ETA: 0:00:44 Epoch: 9 Iter: 1997 Class Loss: 0.00 Loss: 0.03



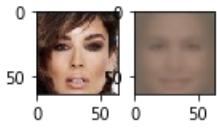
68%|##| ETA: 0:00:41 Epoch: 9 Iter: 2098 Class Loss: 0.00 Loss: 0.03



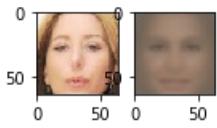
72%|##| ETA: 0:00:36 Epoch: 9 Iter: 2199 Class Loss: 0.00 Loss: 0.04



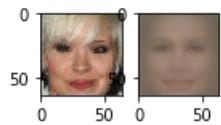
75%|###| ETA: 0:00:37 Epoch: 9 Iter: 2298 Class Loss: 0.00 Loss: 0.03



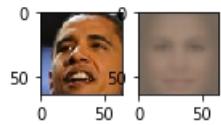
78%|###| ETA: 0:00:33 Epoch: 9 Iter: 2399 Class Loss: 0.00 Loss: 0.04



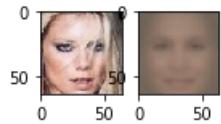
81%|###| ETA: 0:00:23 Epoch: 9 Iter: 2497 Class Loss: 0.00 Loss: 0.03



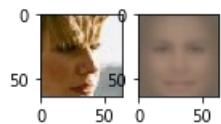
85%|###| ETA: 0:00:19 Epoch: 9 Iter: 2598 Class Loss: 0.00 Loss: 0.03



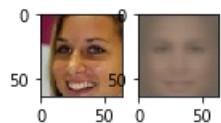
88%|###| ETA: 0:00:15 Epoch: 9 Iter: 2698 Class Loss: 0.00 Loss: 0.03



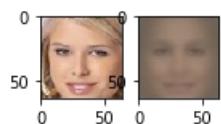
91%|###| ETA: 0:00:10 Epoch: 9 Iter: 2797 Class Loss: 0.00 Loss: 0.03



94%|###| ETA: 0:00:06 Epoch: 9 Iter: 2898 Class Loss: 0.00 Loss: 0.03



98%|###| ETA: 0:00:02 Epoch: 9 Iter: 2999 Class Loss: 0.00 Loss: 0.03



100%|####| Time: 0:02:18 Epoch: 9 Iter: 3052 Class Loss: 0.00 Loss: 0.03
100% (97 of 97) |#####| Elapsed Time: 0:01:06 Time: 0:01:06
N/A% (0 of 72) | Elapsed Time: 0:00:00 ETA: -----

male lighter: 1.0

100% (72 of 72) |#####| Elapsed Time: 0:00:49 Time: 0:00:49
N/A% (0 of 78) | Elapsed Time: 0:00:00 ETA: -----

female lighter: 0.986111111111

100% (78 of 78) |#####| Elapsed Time: 0:00:51 Time: 0:00:51
N/A% (0 of 71) | Elapsed Time: 0:00:00 ETA: -----

male darker: 0.897435897436

100% (71 of 71) |#####| Elapsed Time: 0:00:48 Time: 0:00:48

female darker: 0.985915492958

